

Reinforcement Learning in Control

Dr. Saeed Shamaghdari

**Electrical Engineering Department
Control Group**

Fall 2025 | 4041

Temporal Difference Learning

Temporal Difference Learning

Q Box:

Similarity to dynamic programming (DP)?

Similarity to Monte Carlo methods (MCs)?

Reminder Box

Value function approximation in DP and MC

$$\underline{V(S_t)} \leftarrow \underline{V(S_t)} + \alpha[\underline{G_t} - \underline{V(S_t)}]$$

New value of state t

Former estimation
of value of state t
(= Expected return
starting at that state)

Learning
Rate

Return at
timestep
t

Former estimation
of value of state t
(= Expected return
starting at that state)

The Constant-alpha method in MC (requires episode completion)
Using Return as an estimate of the expected value (Sampling)

Temporal Difference Prediction

Value function approximation TD

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

New value
of state t

Former
estimation of
value of state
t

Learning Rate

Reward

Discounted value of next
state

Q Box:

Target in MC and TD?

Temporal Difference Prediction

Value function approximation TD

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

Diagram illustrating the components of the TD update rule:

- New value of state t (green bar)
- Former estimation of value of state t (blue bar)
- Learning Rate (red bar)
- Reward (orange bar)
- Discounted value of next state (purple bar)
- TD Target (green bar, labeled below)

A blue curved arrow points from the "TD Target" bar to a smiley face icon.

TD(0): one step TD

Q Box:
Target in MC and TD?

To Recap ...

TD Learning Approach:

Temporal Difference Learning: learning at each time step.

$$\underline{V(S_t)} \leftarrow \underline{V(S_t)} + \alpha [R_{t+1} + \gamma \underline{V(S_{t+1})} - \underline{V(S_t)}]$$

New value of state t Former estimation of value of state t Learning Rate Reward Discounted value of next state

TD Target

To Recap ...

TD Approach:



At the end of one step (State, Action, Reward, Next State):

- We have R_{t+1} and S_{t+1}
- We update $V(S_t)$:
 - **We estimate G_t** by adding R_{t+1} and the discounted value of next state.

TD target : $R_{t+1} + \gamma V(S_{t+1})$

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

Now we **continue to interact with this environment** with our updated value function. By running more and more steps, **the agent will learn to play better and better.**

To Recap ...

Two Learning Approaches:

Monte Carlo: learning at the end of the episode

A 4x4 grid-based diagram illustrating Monte Carlo learning. At the top left is a cat icon. Three cheese icons are arranged in a row across the top. In the bottom-left corner, a mouse icon is positioned next to a cheese icon. Arrows indicate a path from the mouse to the cheese, then up to another cheese, right to a third cheese, up to the row of three, right to the final cheese, down to the mouse, left to the first cheese, up to the row of three, right to the final cheese, and finally up to the cat.

Temporal Difference Learning: learning at each step.

A 4x4 grid-based diagram illustrating Temporal Difference learning. The layout is identical to the Monte Carlo diagram: a cat at the top left, three cheese icons at the top, a mouse icon and one cheese icon in the bottom-left. However, the arrows show a continuous loop: from the mouse to the cheese, up to the first cheese, right to the second cheese, up to the third cheese, right to the fourth cheese, up to the mouse, left to the first cheese, up to the second cheese, right to the third cheese, up to the fourth cheese, and finally right to the cat.

Deep RL Course

Temporal Difference Prediction

Value function approximation TD

$$\begin{aligned}v_{\pi}(s) &\doteq \mathbb{E}_{\pi}[G_t \mid S_t = s] \\&= \mathbb{E}_{\pi}[R_{t+1} + \gamma G_{t+1} \mid S_t = s] \\&= \mathbb{E}_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_t = s]\end{aligned}$$

- **DP:** Computing the expected value using the model, utilizing previous value function estimates in new estimates (**bootstrapping**).
- **TD:** Sampling, using previous value function estimates in new estimates (**bootstrapping**).

Algorithm

Tabular TD(0) for estimating v_π

Input: the policy π to be evaluated

Algorithm parameter: step size $\alpha \in (0, 1]$

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

$A \leftarrow$ action given by π for S

 Take action A , observe R, S'

$V(S) \leftarrow V(S) + \alpha [R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

 until S is terminal

TD Error

$$\delta_t \doteq R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$

Monte Carlo Error

$$\begin{aligned}
G_t - V(S_t) &= R_{t+1} + \gamma G_{t+1} - V(S_t) + \gamma V(S_{t+1}) - \gamma V(S_{t+1}) \\
&= \delta_t + \gamma(G_{t+1} - V(S_{t+1})) \\
&= \delta_t + \gamma\delta_{t+1} + \gamma^2(G_{t+2} - V(S_{t+2})) \\
&= \delta_t + \gamma\delta_{t+1} + \gamma^2\delta_{t+2} + \cdots + \gamma^{T-t-1}\delta_{T-1} + \gamma^{T-t}(G_T - V(S_T)) \\
&= \delta_t + \gamma\delta_{t+1} + \gamma^2\delta_{t+2} + \cdots + \gamma^{T-t-1}\delta_{T-1} + \gamma^{T-t}(0 - 0) \\
&= \sum_{k=t}^{T-1} \gamma^{k-t} \delta_k.
\end{aligned}$$

Example: Driving Home

Value function approximation TD

<i>State</i>	<i>Elapsed Time</i> (minutes)	<i>Predicted</i> <i>Time to Go</i>	<i>Predicted</i> <i>Total Time</i>
leaving office, friday at 6	0	30	30
reach car, raining	5	35	40
exiting highway	20	15	35
2ndary road, behind truck	30	10	40
entering home street	40	3	43
arrive home	43	0	43

Reward?

Q Box

Return? T_Go

Value? E(T_Go)

Example: Driving Home

Value function approximation TD

<i>State</i>	<i>Elapsed Time</i> (minutes)	<i>Predicted</i> <i>Time to Go</i>	<i>Predicted</i> <i>Total Time</i>
leaving office, friday at 6	0	30	30
reach car, raining	5	35	40
exiting highway	20	15	35
2ndary road, behind truck	30	10	40
entering home street	40	3	43
arrive home	43	0	43



Estimation error at highway exit (based on Monte Carlo)?

Q Box

Estimation correction at highway exit (MC-based) for alpha=0.5?

$$\alpha(G_t - V(s_t))?$$

TD: Faster Learning



Temporal Difference Learning

TD Convergence

- Very small α
- α with decay conditions

$$\sum_{n=1}^{\infty} \alpha_n(a) = \infty$$



and

$$\sum_{n=1}^{\infty} \alpha_n^2(a) < \infty$$

Convergence speed of TD and MC

Markov Reward Process: Random Walk Example

In this example we empirically compare the prediction abilities of TD(0) and constant- α MC when applied to the following Markov reward process:



Random Walk Under Batch Updating

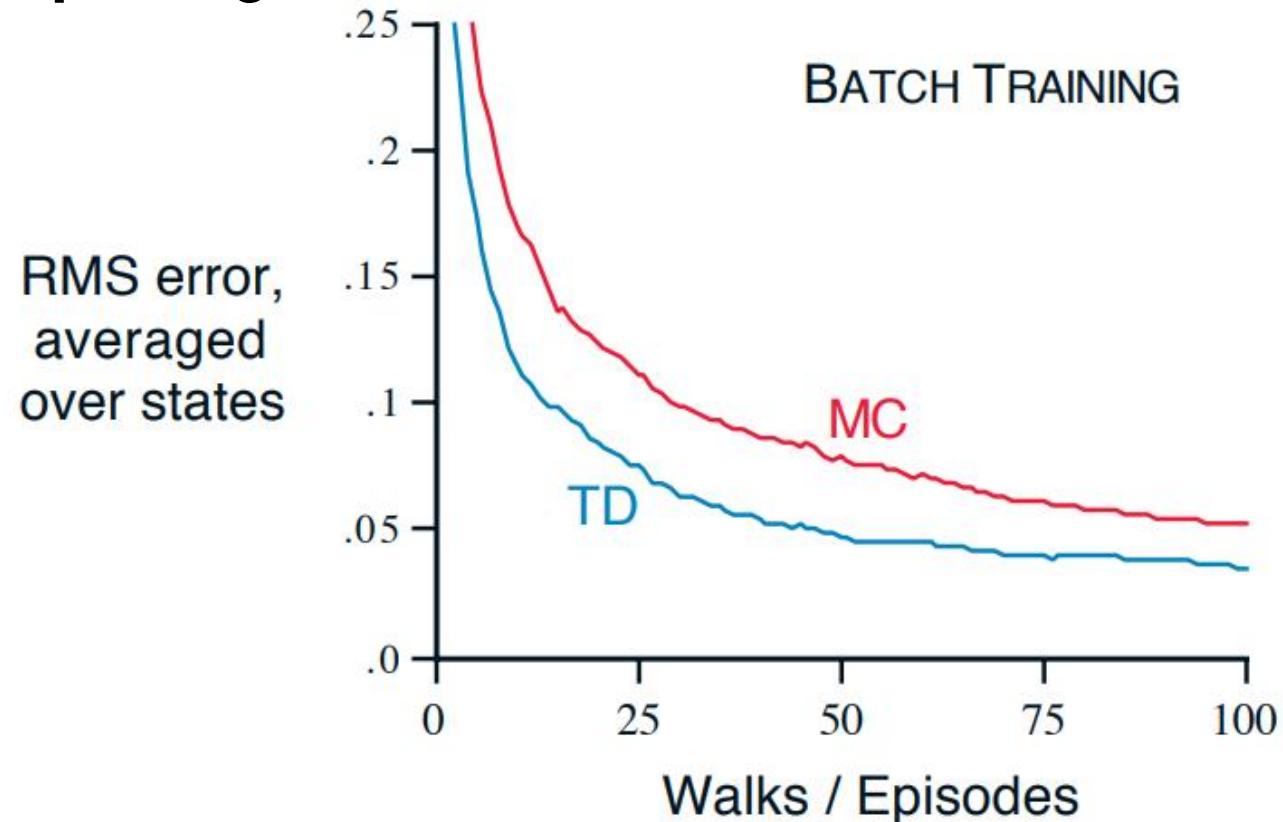


Figure 6.2: Performance of TD(0) and constant- α MC under batch training on the random walk task.

Example: You are the Predictor

A, 0, B, 0

B, 1

B, 1


B, 1

B, 1

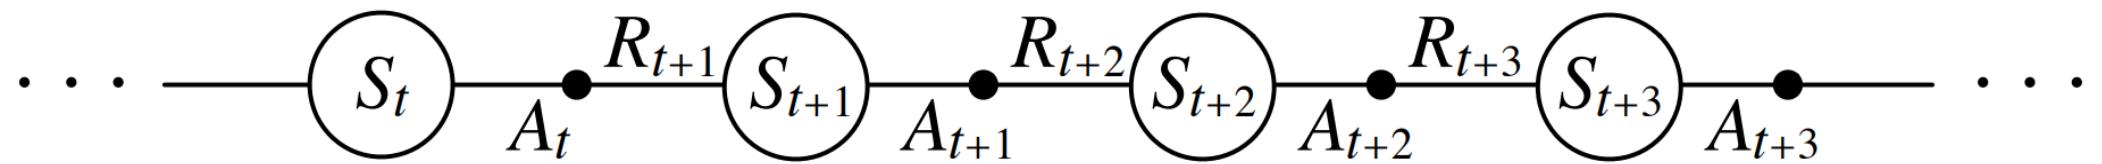
B, 1

B, 1

B, 0




SARSA: On-Policy TD Control



$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t) \right]$$

SARSA: On-Policy TD Control



Algorithm

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R, S'


 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$$


$S \leftarrow S'; A \leftarrow A'$;

 until S is terminal


Better Intuition



Example: Windy Grid



Example: Windy Grid



Q-Learning: Off-Policy TD Control


$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

The diagram illustrates the Q-Learning update rule with color-coded components:

- New Q-value estimation** (green bar)
- Former Q-value estimation** (blue bar)
- Learning Rate** (red bar)
- Immediate Reward** (orange bar)
- Discounted Estimate optimal Q-value of next state** (purple bar)
- TD Target** (dark blue bar)
- TD Error** (yellow bar)

The update rule is shown as: $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$

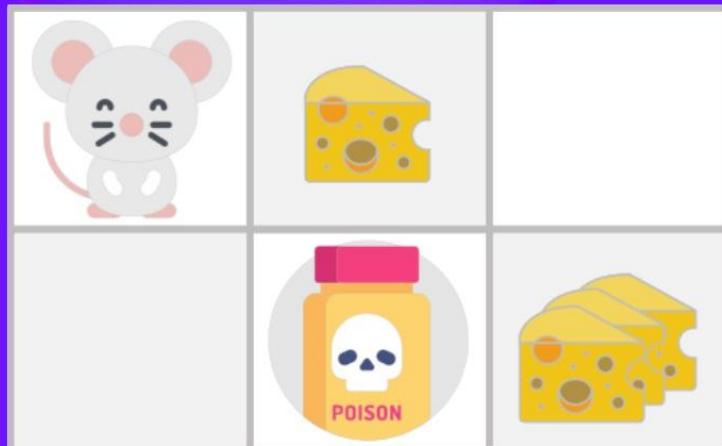
Example



- You always start at the **same starting point**.
- The goal: eat the **big pile of cheese** (at the bottom right-hand corner) and **avoid the poison**.
- The episode ends if we eat the poison, eat the big pile of cheese or if we spent more than 5 steps.
- Learning rate = 0.1
- Gamma = 0.99

Example

- The reward function:
 - 0: Going to a state **with no cheese in it.**
 - +1: Going to a state with a **small cheese in it.**
 - +10: Going to the state with **the big pile of cheese.**
 - -10: Going to the state **with the poison and thus die.**



Example, Step 1


Initialize Q arbitrarily (e.g., $Q(s, a) = 0$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$, and $Q(\text{terminal-state}, \cdot) = 0$)

	\leftarrow	\rightarrow	\uparrow	\downarrow
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0

We initialize the Q-Table

Example, Step 2

Choose action A_t using policy derived from Q (e.g., ϵ -greedy)



We took a random action (exploration)

Example, Step 3

Take action A_t and observe R_{t+1}, S_{t+1}



Example, Step 4

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

New
Q-value
estimation

Former
Q-value
estimation

Learning
Rate

Immediate
Reward

Discounted Estimate
optimal Q-value
of next state

Former
Q-value
estimation

TD Target

TD Error

Update our Q-value estimation

Example, Step 4

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t))$$

$$Q(\text{Initial state, Right}) = 0 + 0.1 * [1 + 0.99 * 0 - 0]$$

$$Q(\text{Initial state, Right}) = 0.1$$

	\leftarrow	\rightarrow	\uparrow	\downarrow
	0	0.1	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0

Q-Learning Recap: Algorithm

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)


 Take action A , observe R, S'

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$$S \leftarrow S'$$

 until S is terminal

Example: Cliff Walking



Example: Cliff Walking

