

# Reinforcement Learning in Control

**Dr. Saeed Shamaghdari**

**Electrical Engineering Department  
Control Group**

Fall 2025 | 4041

# Deep Reinforcement Learning

---

# Playing Atari with Deep Reinforcement Learning

---

**Volodymyr Mnih    Koray Kavukcuoglu    David Silver    Alex Graves    Ioannis Antonoglou**

**Daan Wierstra    Martin Riedmiller**

DeepMind Technologies

{vlad, koray, david, alex.graves, ioannis, daan, martin.riedmiller} @ deepmind.com

[☆ Save](#) [✉ Cite](#) [Cited by 19338](#) [Related articles](#) [All 37 versions](#) [»](#)

## Human-level control through deep reinforcement learning

Volodymyr Mnih<sup>1\*</sup>, Koray Kavukcuoglu<sup>1\*</sup>, David Silver<sup>1\*</sup>, Andrei A. Rusu<sup>1</sup>, Joel Veness<sup>1</sup>, Marc G. Bellemare<sup>1</sup>, Alex Graves<sup>1</sup>, Martin Riedmiller<sup>1</sup>, Andreas K. Fidjeland<sup>1</sup>, Georg Ostrovski<sup>1</sup>, Stig Petersen<sup>1</sup>, Charles Beattie<sup>1</sup>, Amir Sadik<sup>1</sup>, Ioannis Antonoglou<sup>1</sup>, Helen King<sup>1</sup>, Dharshan Kumaran<sup>1</sup>, Daan Wierstra<sup>1</sup>, Shane Legg<sup>1</sup> & Demis Hassabis<sup>1</sup>

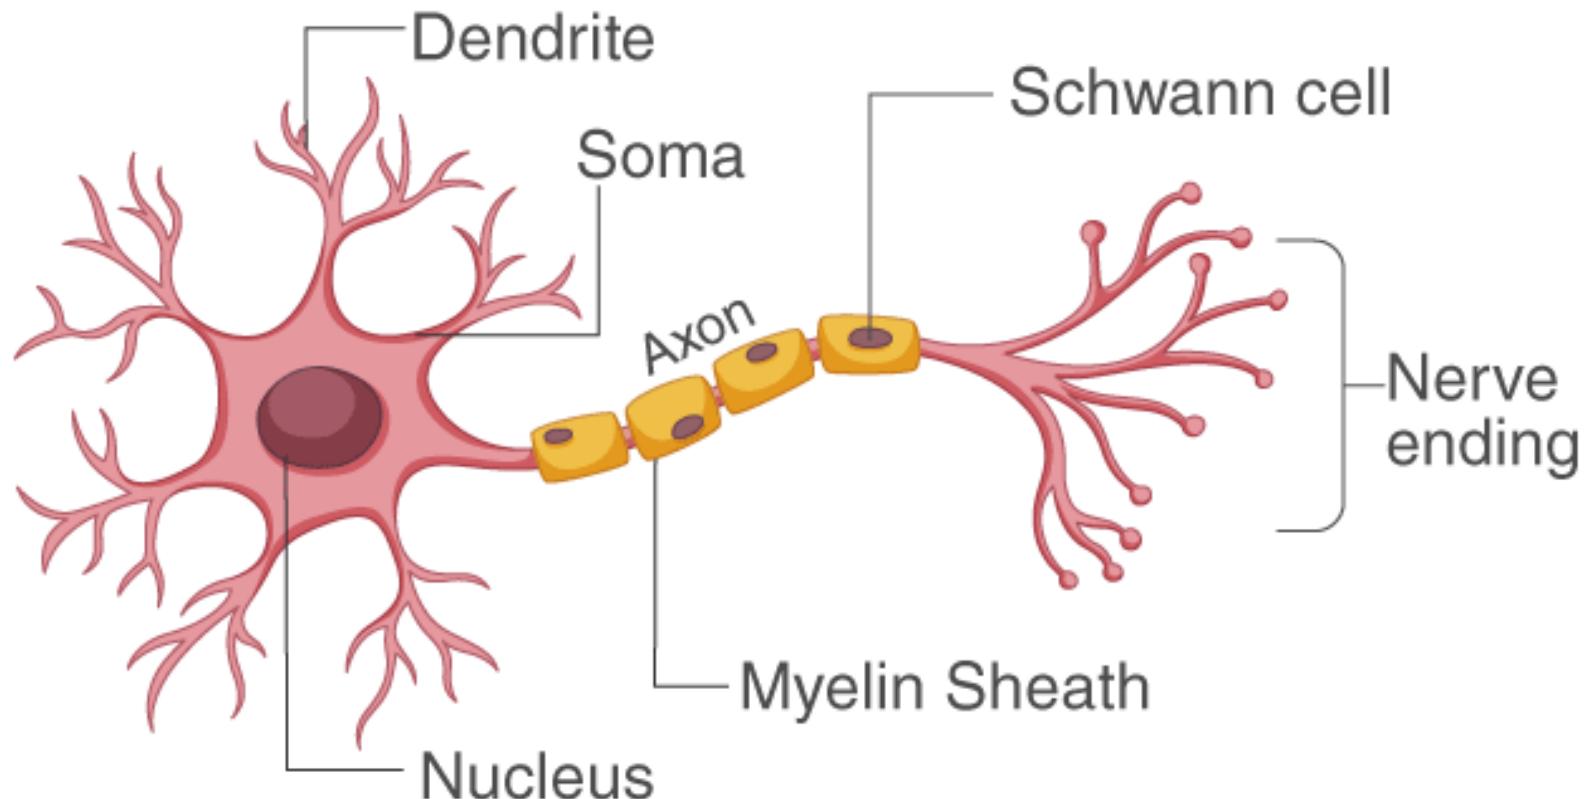
---

<sup>1</sup>Google DeepMind, 5 New Street Square, London EC4A 3TW, UK.

\*These authors contributed equally to this work.

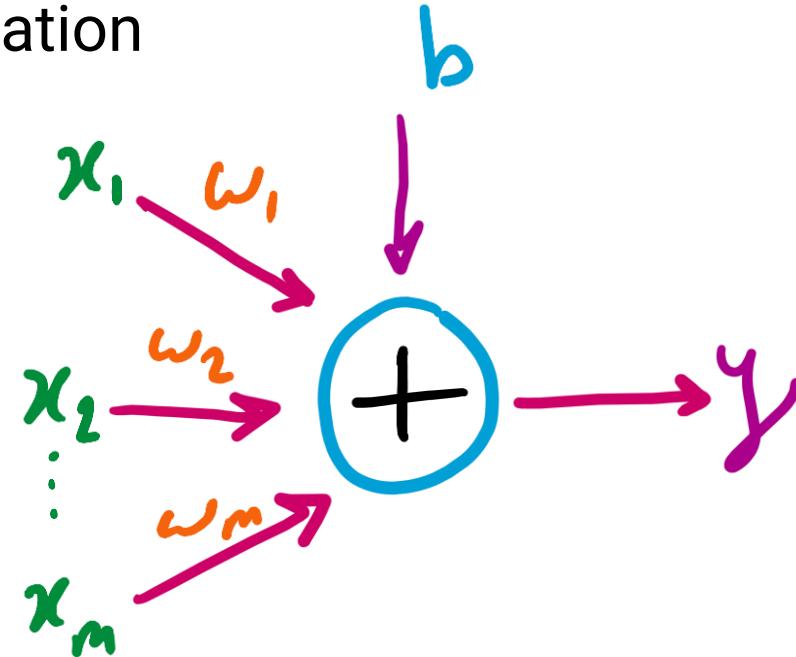
# Neural Networks and Deep Learning: A Simple Review

## Neuron



# Neural Networks and Deep Learning: A Simple Review

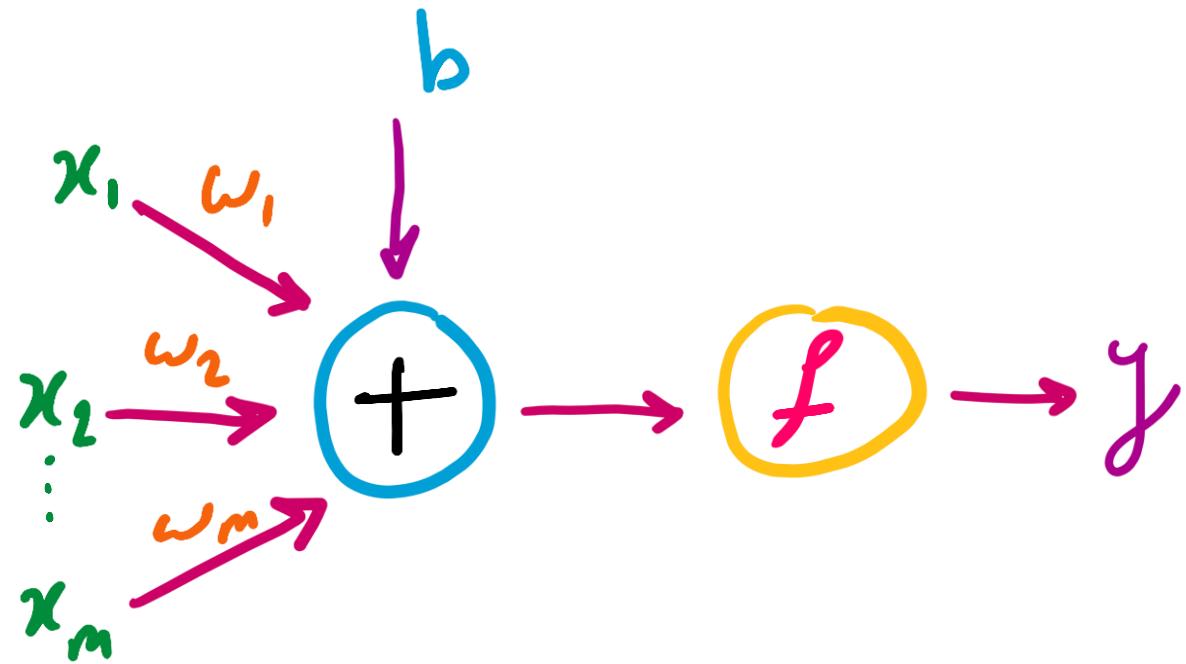
## Artificial Neuron Formulation



$$y = w_1x_1 + w_2x_2 + \dots + w_mx_m + b$$

# Neural Networks and Deep Learning: A Simple Review

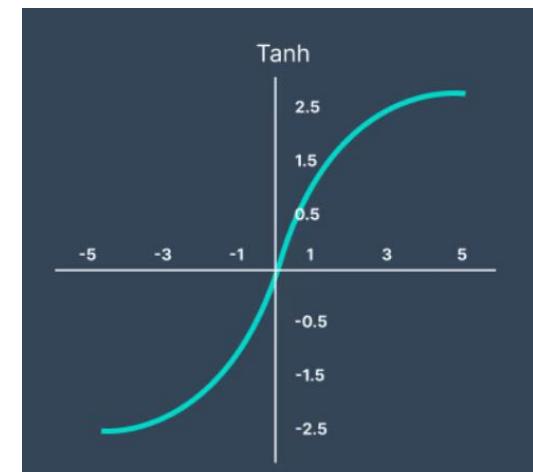
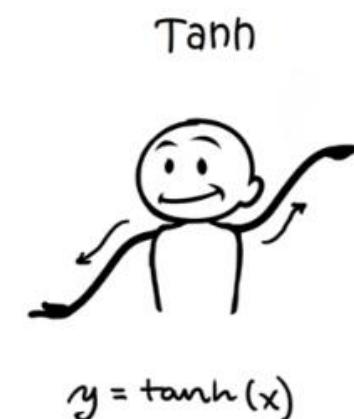
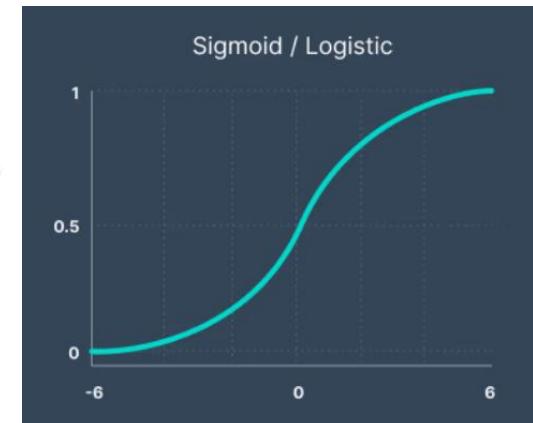
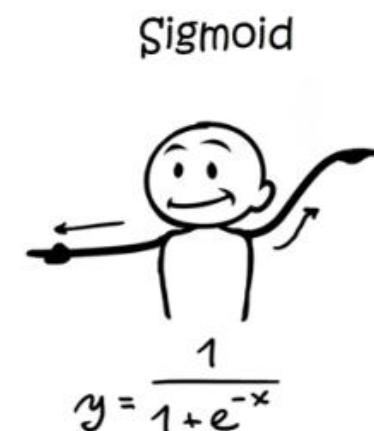
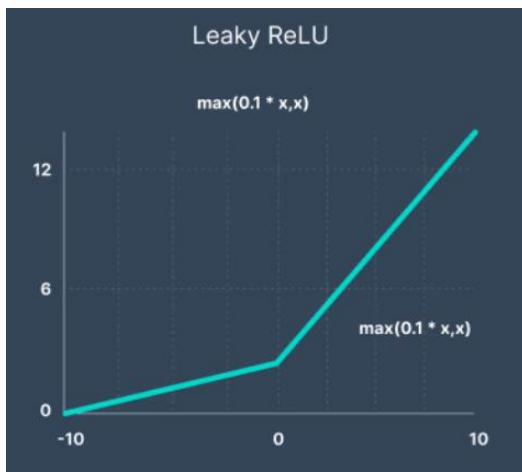
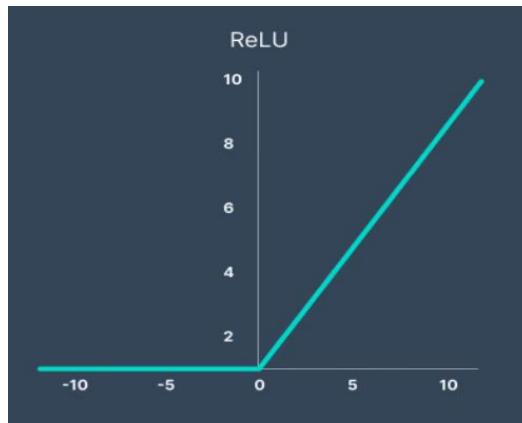
Q: Problems?



$$y = f(w_1x_1 + w_2x_2 + \dots + w_mx_m + b) = f(\mathbf{w}^\top \mathbf{x} + b)$$

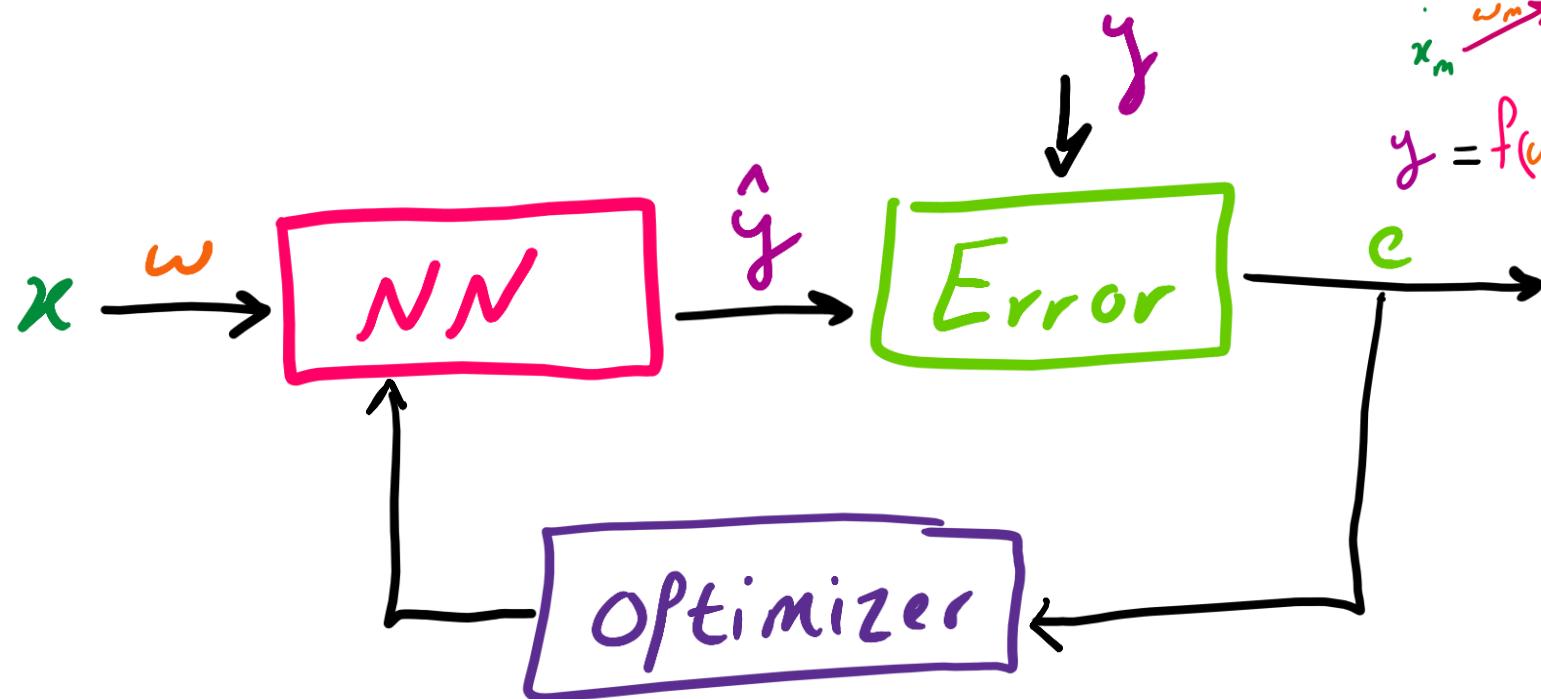
# Neural Networks and Deep Learning: A Simple Review

## Activation Functions



# Neural Networks and Deep Learning: A Simple Review

## Learning Block Diagram



A detailed diagram of a single neuron model. Inputs  $x_1, x_2, \dots, x_m$  (green) are multiplied by weights  $w_1, w_2, \dots, w_m$  (orange), and a bias  $b$  (blue) is added. The result is summed up in a blue circle labeled  $\Sigma$ . The output is passed through an activation function  $f$  (yellow) to produce the final output  $y$  (purple).

$$y = f(w_1x_1 + w_2x_2 + \dots + w_mx_m + b)$$

## Neural Networks and Deep Learning: A Simple Review

Learning Block Diagram

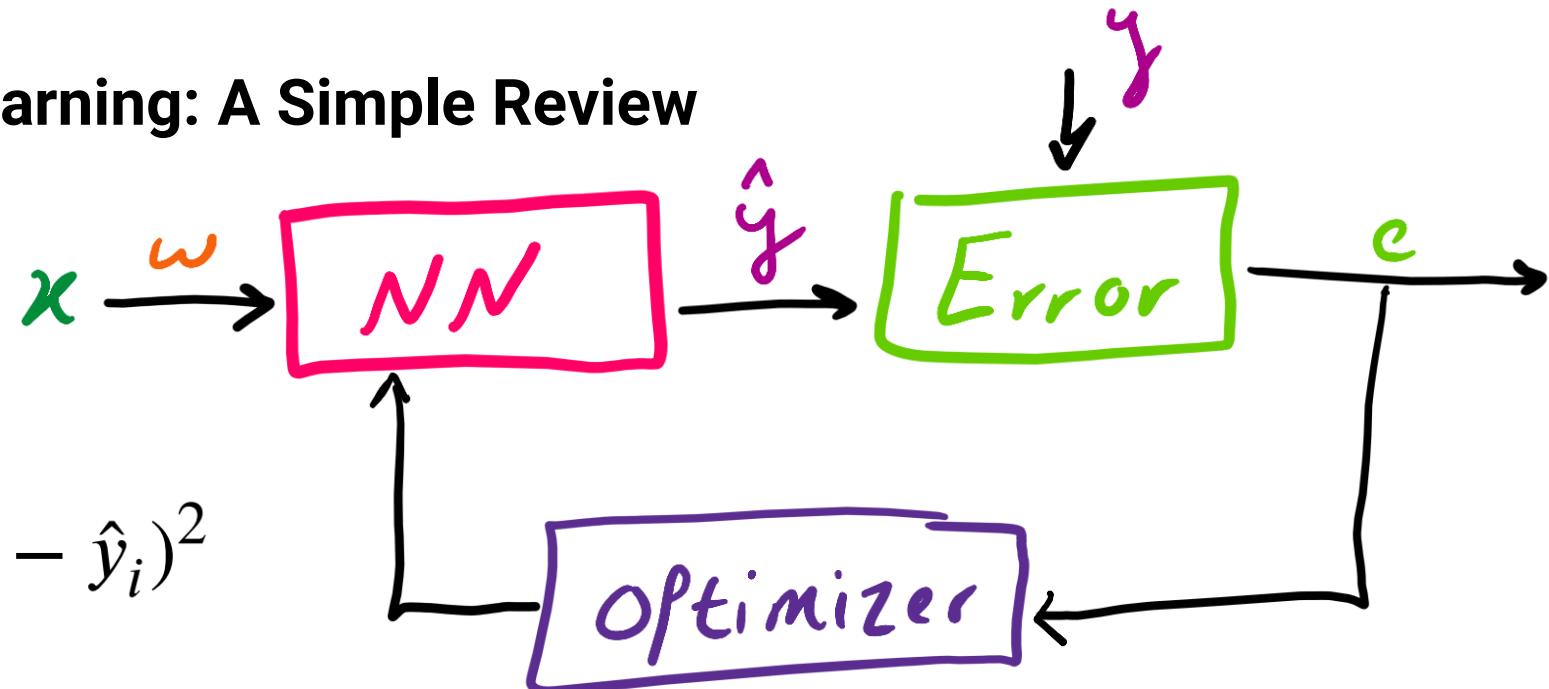
Loss Function:

MSE:

$$L(y, \hat{y}) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

MAE:

$$L(y, \hat{y}) = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$



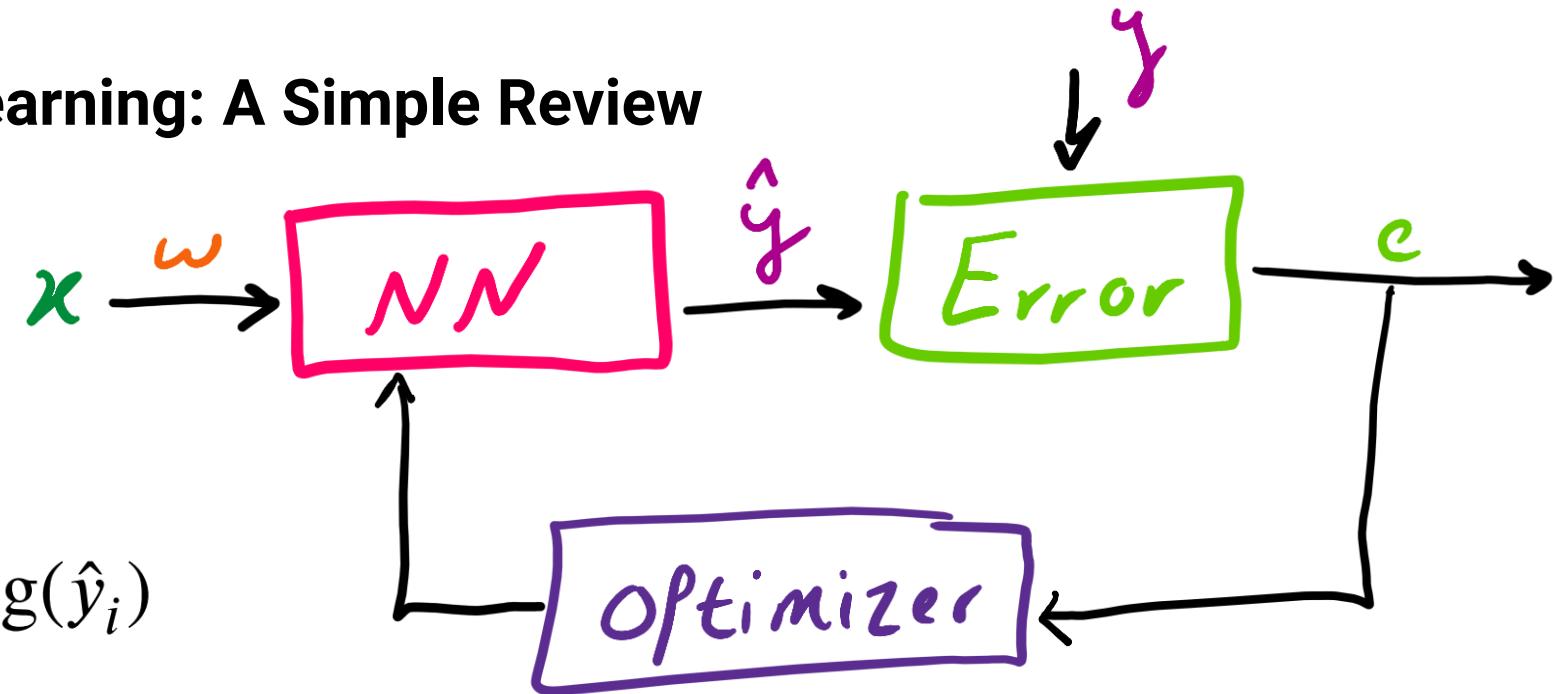
## Neural Networks and Deep Learning: A Simple Review

Learning Block Diagram

Loss Function:

*Cross Entropy*:

$$L(y, \hat{y}) = - \sum_{i=1}^N y_i \log(\hat{y}_i)$$

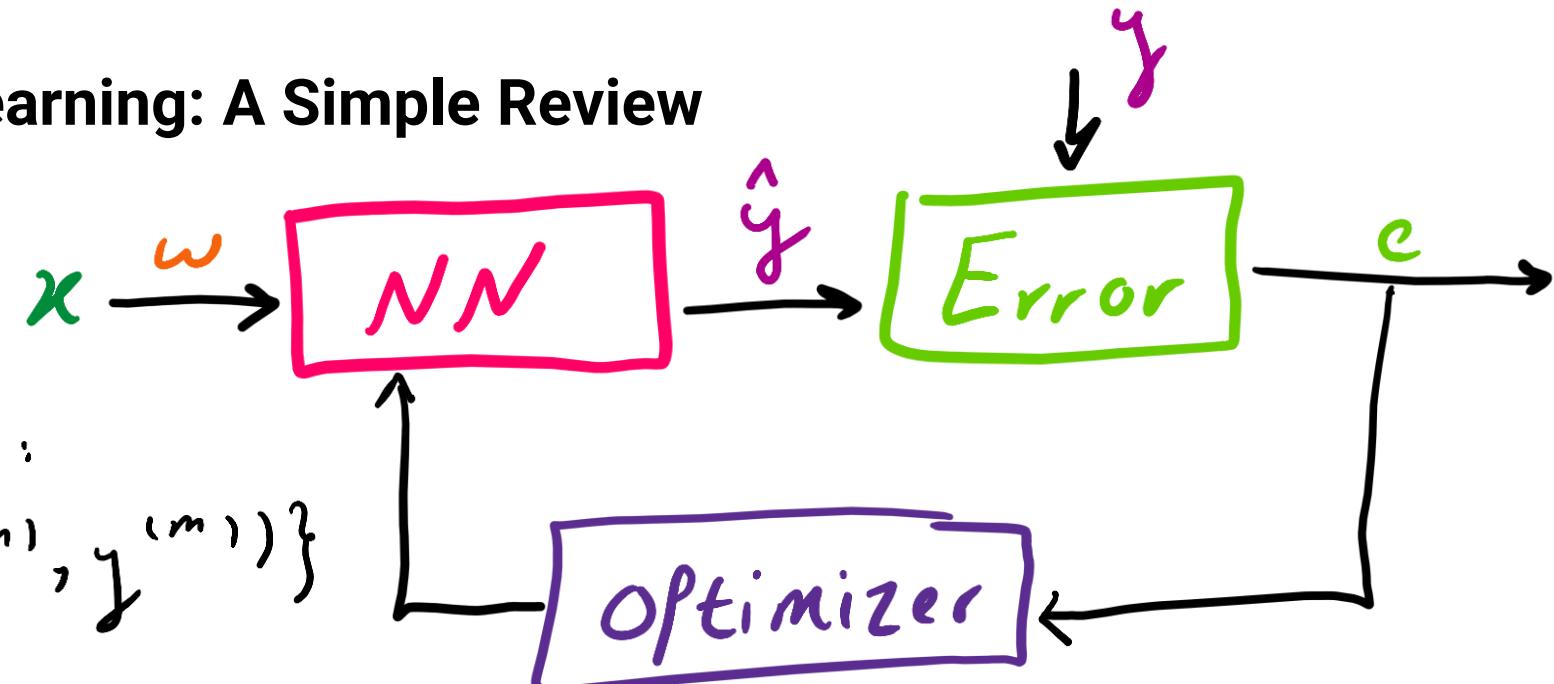


*Binary Cross Entropy*:

$$L(y, \hat{y}) = -(y \log(\hat{y}) + (1 - y) \log(1 - \hat{y}))$$

## Neural Networks and Deep Learning: A Simple Review

Learning Block Diagram



Given  $m$  train examples:

$$\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$$

want  $\hat{y}^{(i)} \approx y^{(i)}$

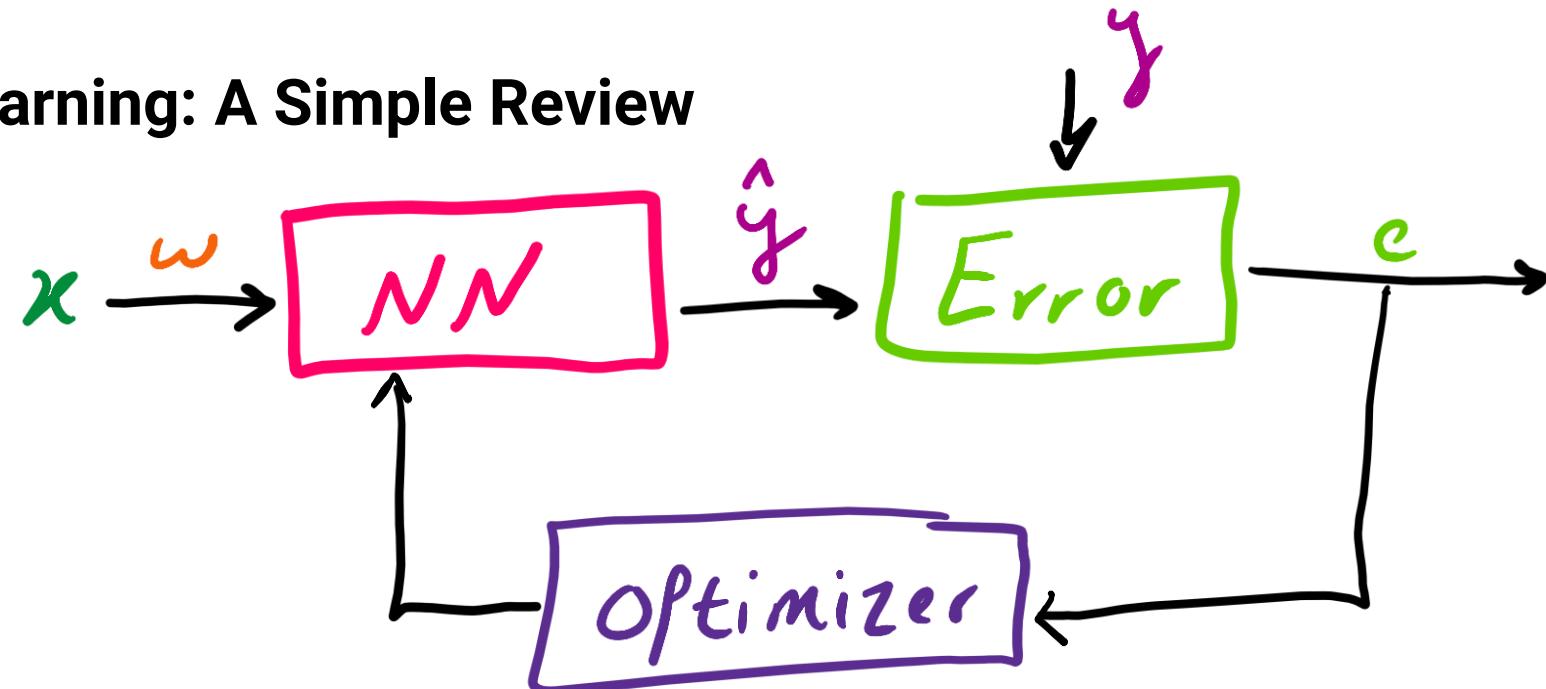
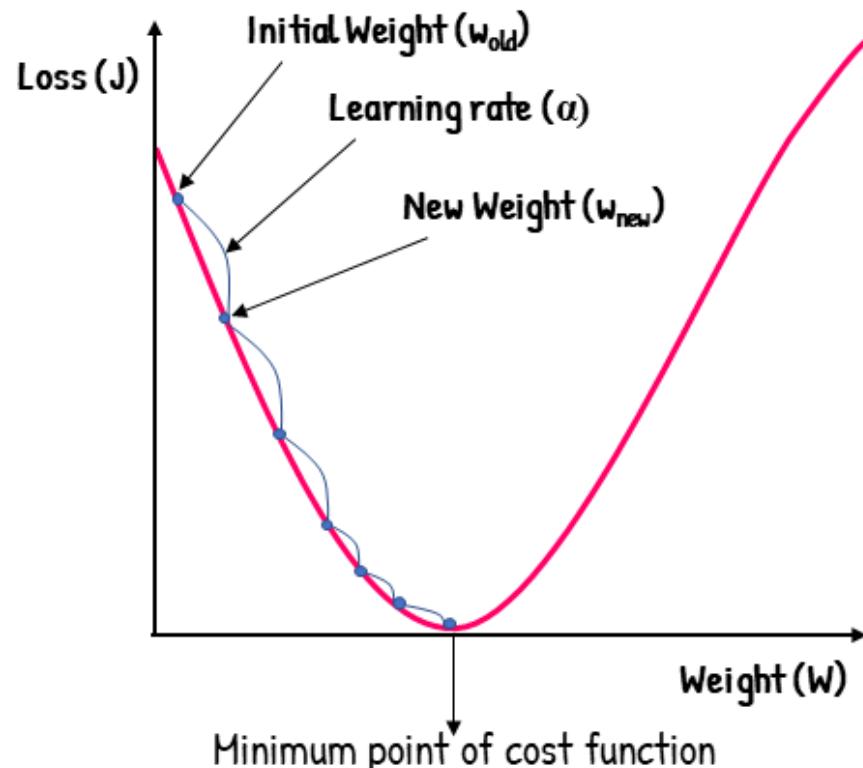
So the cost function is:

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m L(y, \hat{y})$$

# Neural Networks and Deep Learning: A Simple Review

Learning Block Diagram

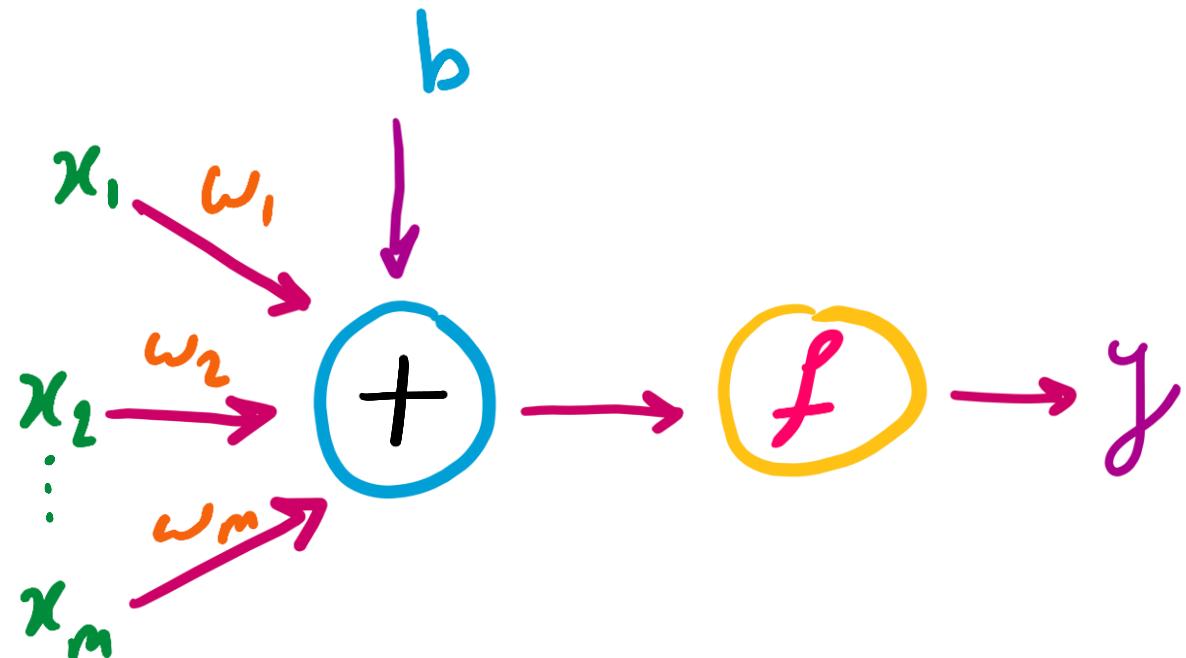
**Gradient Descent:**



$$w_{new} = w_{old} - \alpha \frac{\delta J}{\delta w}$$

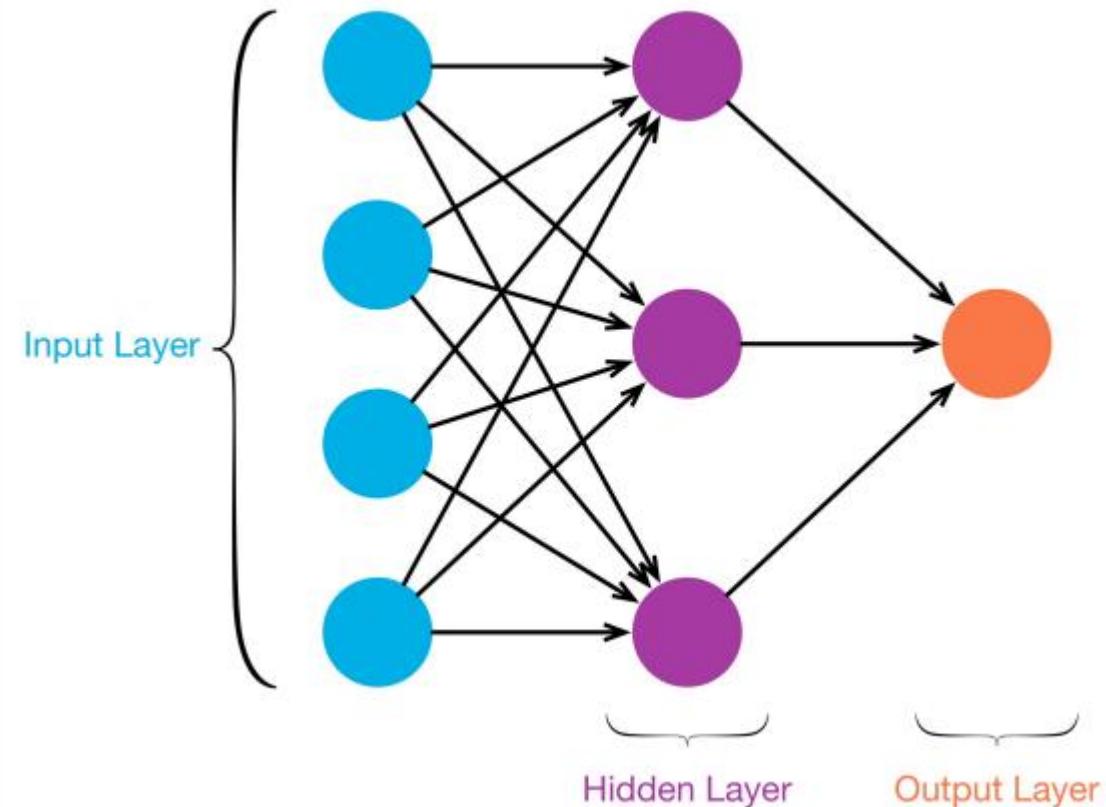
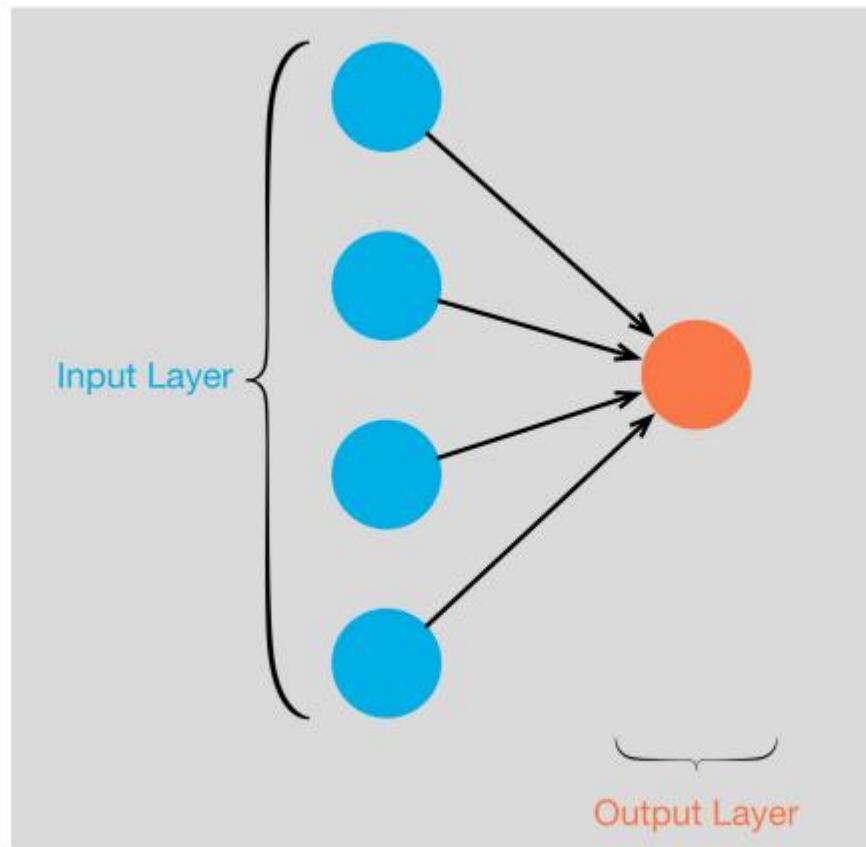
## Neural Networks and Deep Learning: A Simple Review

### Single-Layer Perceptron

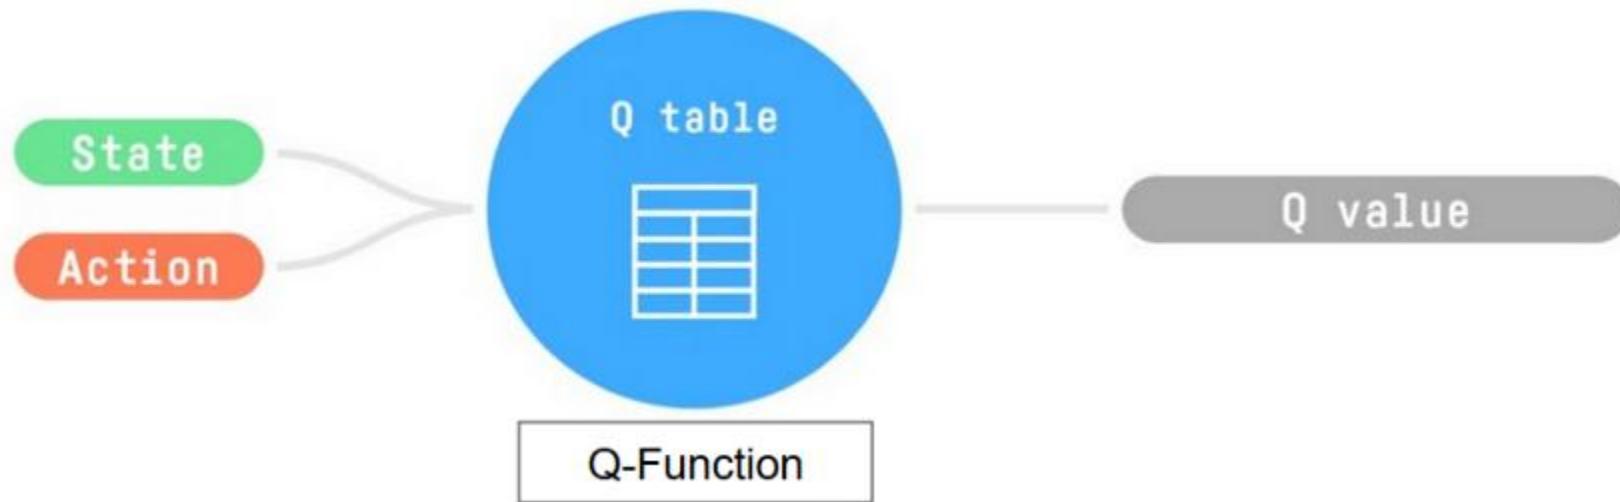


# Neural Networks and Deep Learning: A Simple Review

Multi-Layer Perceptron ([MLP](#)) ([Play!](#))



## Q-Learning Recap ...



## Pseudocode

### Q-Learning

---

**Algorithm 14:** Sarsamax (Q-Learning)
 

---

**Input:** policy  $\pi$ , positive integer  $num\_episodes$ , small positive fraction  $\alpha$ , GLIE  $\{\epsilon_i\}$

**Output:** value function  $Q$  ( $\approx q_\pi$  if  $num\_episodes$  is large enough)

Initialize  $Q$  arbitrarily (e.g.,  $Q(s, a) = 0$  for all  $s \in \mathcal{S}$  and  $a \in \mathcal{A}(s)$ , and  $Q(\text{terminal-state}, \cdot) = 0$ )

**for**  $i \leftarrow 1$  **to**  $num\_episodes$  **do**

→ Step 1

$\epsilon \leftarrow \epsilon_i$

    Observe  $S_0$

$t \leftarrow 0$

**repeat**

        Choose action  $A_t$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)   Step 2

        Take action  $A_t$  and observe  $R_{t+1}, S_{t+1}$    Step 3

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t))$    Step 4

$t \leftarrow t + 1$

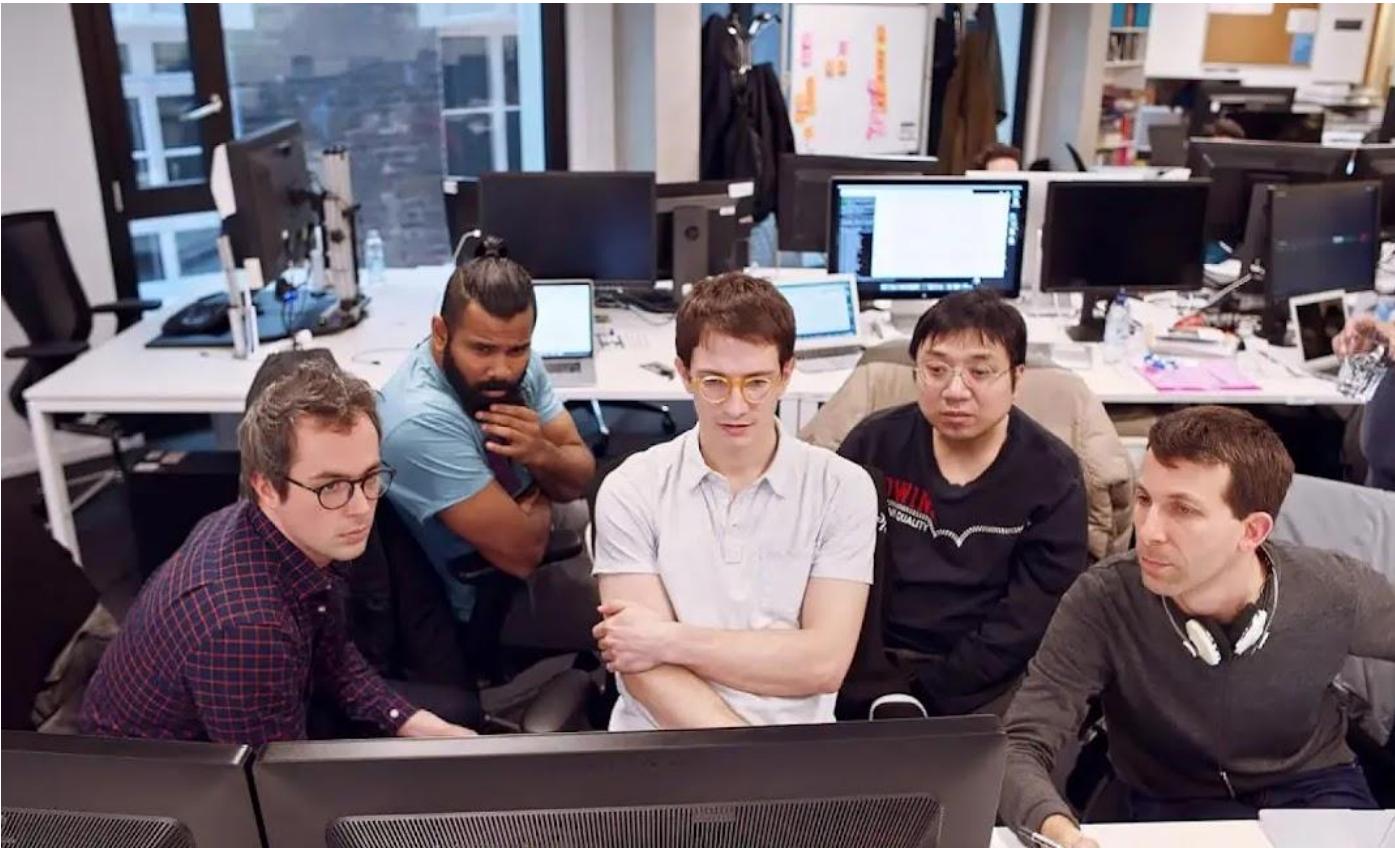
**until**  $S_t$  is terminal;

**end**

**return**  $Q$

---

## Why Deep Reinforcement Learning?



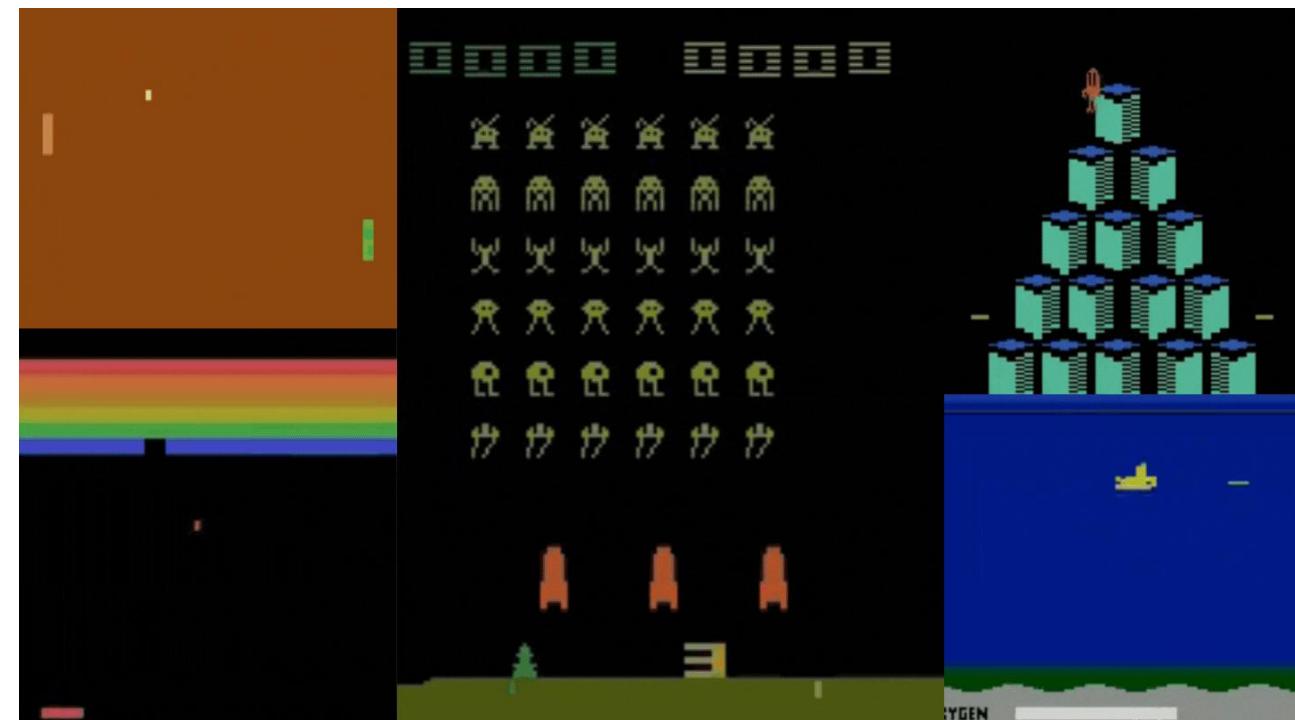
Playing Atari with Deep Reinforcement Learning

([Paper](#))

RL in Control | IUST

## Why Deep Reinforcement Learning?

چالش جدی RL: کنترل مبتنی بر یادگیری مستقیماً از دیتای با ابعاد بزرگ (تصویر یا صوت یا سنسورها)



Q: Number of states in an 8\*8 Gridworld?  
What about an Atari game?

Playing Atari with Deep Reinforcement Learning ([Paper](#))

## Why Deep Reinforcement Learning?



(Left-to-right) Pong, Breakout, Space Invaders, Seaquest, Beam Rider

“Atari 2600 :  
*visual input ( $210 \times 160$  RGB video at 60Hz)*

Playing Atari with Deep Reinforcement Learning ([Paper](#))

## Why Deep Reinforcement Learning?



(Left-to-right) Pong, Breakout, Space Invaders, Seaquest, Beam Rider

Each Frame: (210, 160, 3) containing values ranging from 0 to 255

**Q:** Number of states?

**A:**  $256^{210 \times 160 \times 3} = 256^{100800}$

Idea: approximate Q-values

Using parametrized Q-function  $Q_\theta(s, a)$

Playing Atari with Deep Reinforcement Learning ([Paper](#))

## DeepRL: Why Is It Still Hard?

- روش های موفق Deep Learning دیتاهای عظیم لیبل دار
- یادگیری RL از پاداش اسکالر به صورت sparse و نویزی و تاخیر دار (بر عکس DL)

Playing Atari with Deep Reinforcement Learning ([Paper](#))

## DeepRL: Why Is It Still Hard?

DL data samples: **IID**

**RL:** Highly correlated states.

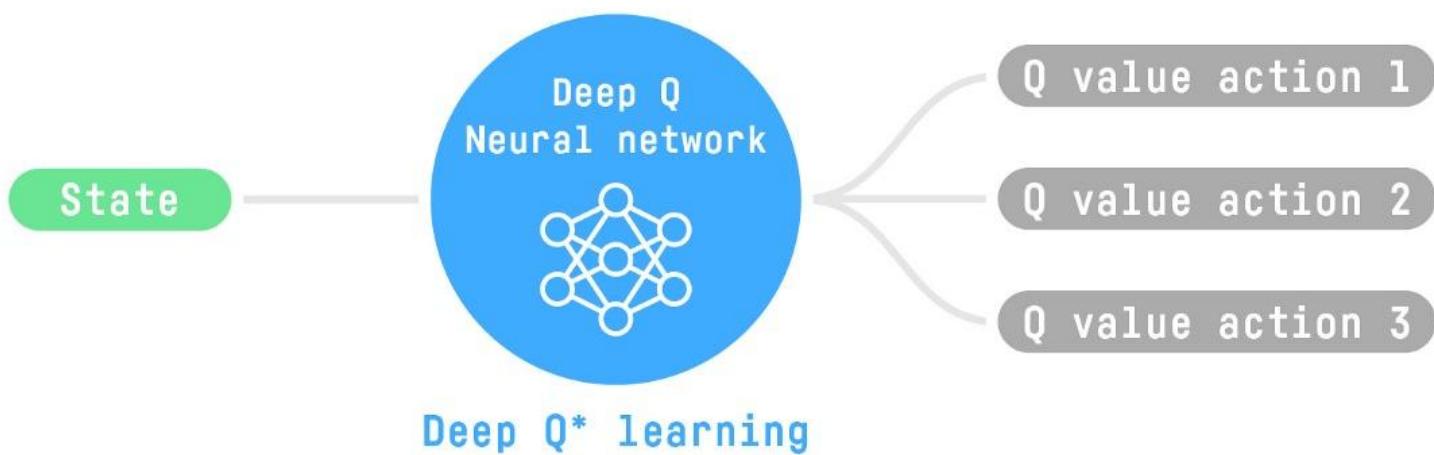
تغییر توزیع دیتا در RL در فرآیند یادگیری

Playing Atari with Deep Reinforcement Learning ([Paper](#))

## Deep Q-Learning (DQN)

The best idea is to approximate the Q-values using a parametrized Q-function  $Q_\theta(s, a)$ .

:DL مشابه



Playing Atari with Deep Reinforcement Learning ([Paper](#))

## Deep Q-Learning (DQN)

*Emulator: agent interacts with an environment  $\mathcal{E}$*

**Goal:** maximizes future rewards.

The future discounted *return* at time  $t$ :

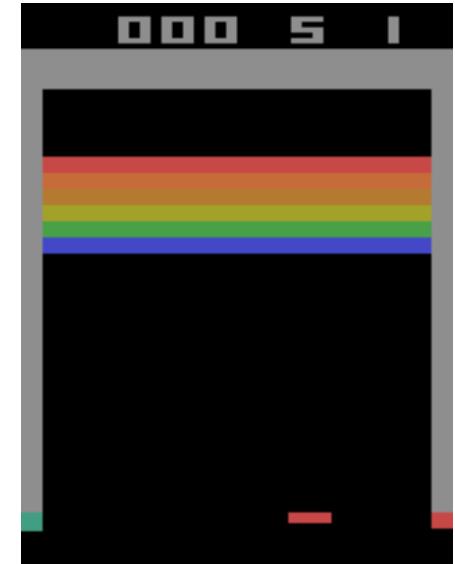
$$R_t = \sum_{t'=t}^T \gamma^{t'-t} r_{t'}$$

$r_t$ : reward

The optimal action-value function:

$$Q^*(s, a) = \max_{\pi} \mathbb{E}[R_t | s_t = s, a_t = a, \pi]$$

Playing Atari with Deep Reinforcement Learning ([Paper](#))



## Deep Q-Learning (DQN)

optimal strategy :

select  $a'$  :

maximizing the expected value of  $r + \gamma Q^*(s', a')$

$$Q^*(s, a) = \mathbb{E}_{s' \sim \mathcal{E}} \left[ r + \gamma \max_{a'} Q^*(s', a') \middle| s, a \right]$$

### Reminder Box: Value Iteration

$$Q_{i+1}(s, a) = \mathbb{E} [r + \gamma \max_{a'} Q_i(s', a') | s, a]$$
$$Q_i \rightarrow Q^* \text{ as } i \rightarrow \infty$$

Playing Atari with Deep Reinforcement Learning ([Paper](#))

## Deep Q-Learning (DQN)

A function approximator to estimate the action-value function:

$$Q(s, a) \approx \text{A Neural Network!}$$

Playing Atari with Deep Reinforcement Learning ([Paper](#))

## Deep Q-Learning (DQN)

*Neural network function approximator with weights  $\theta$  :  **$Q$ -network***

*Training:*

*minimizing*      
$$L_i(\theta_i) = \mathbb{E}_{s,a \sim \rho(\cdot)} \left[ (y_i - Q(s, a; \theta_i))^2 \right]$$

*“Where*

*target:*

$$y_i = \mathbb{E}_{s' \sim \mathcal{E}} [r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) | s, a]$$

$\rho(s, a)$ : behaviour

تارگت وابسته به وزنهای شبکه است بر عکس Suprvised DL که ثابت است

Playing Atari with Deep Reinforcement Learning ([Paper](#))

## Deep Q-Learning (DQN)

مشتق تابع Loss نسبت به وزنها:

$$\nabla_{\theta_i} L_i(\theta_i) = \mathbb{E}_{s,a \sim \rho(\cdot); s' \sim \mathcal{E}} \left[ \left( r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i) \right) \nabla_{\theta_i} Q(s, a; \theta_i) \right]$$

### Reminder Box

$$L_i(\theta_i) = \mathbb{E}_{s,a \sim \rho(\cdot)} \left[ (y_i - Q(s, a; \theta_i))^2 \right]$$

*stochastic gradient descent*

model-free and off-policy.

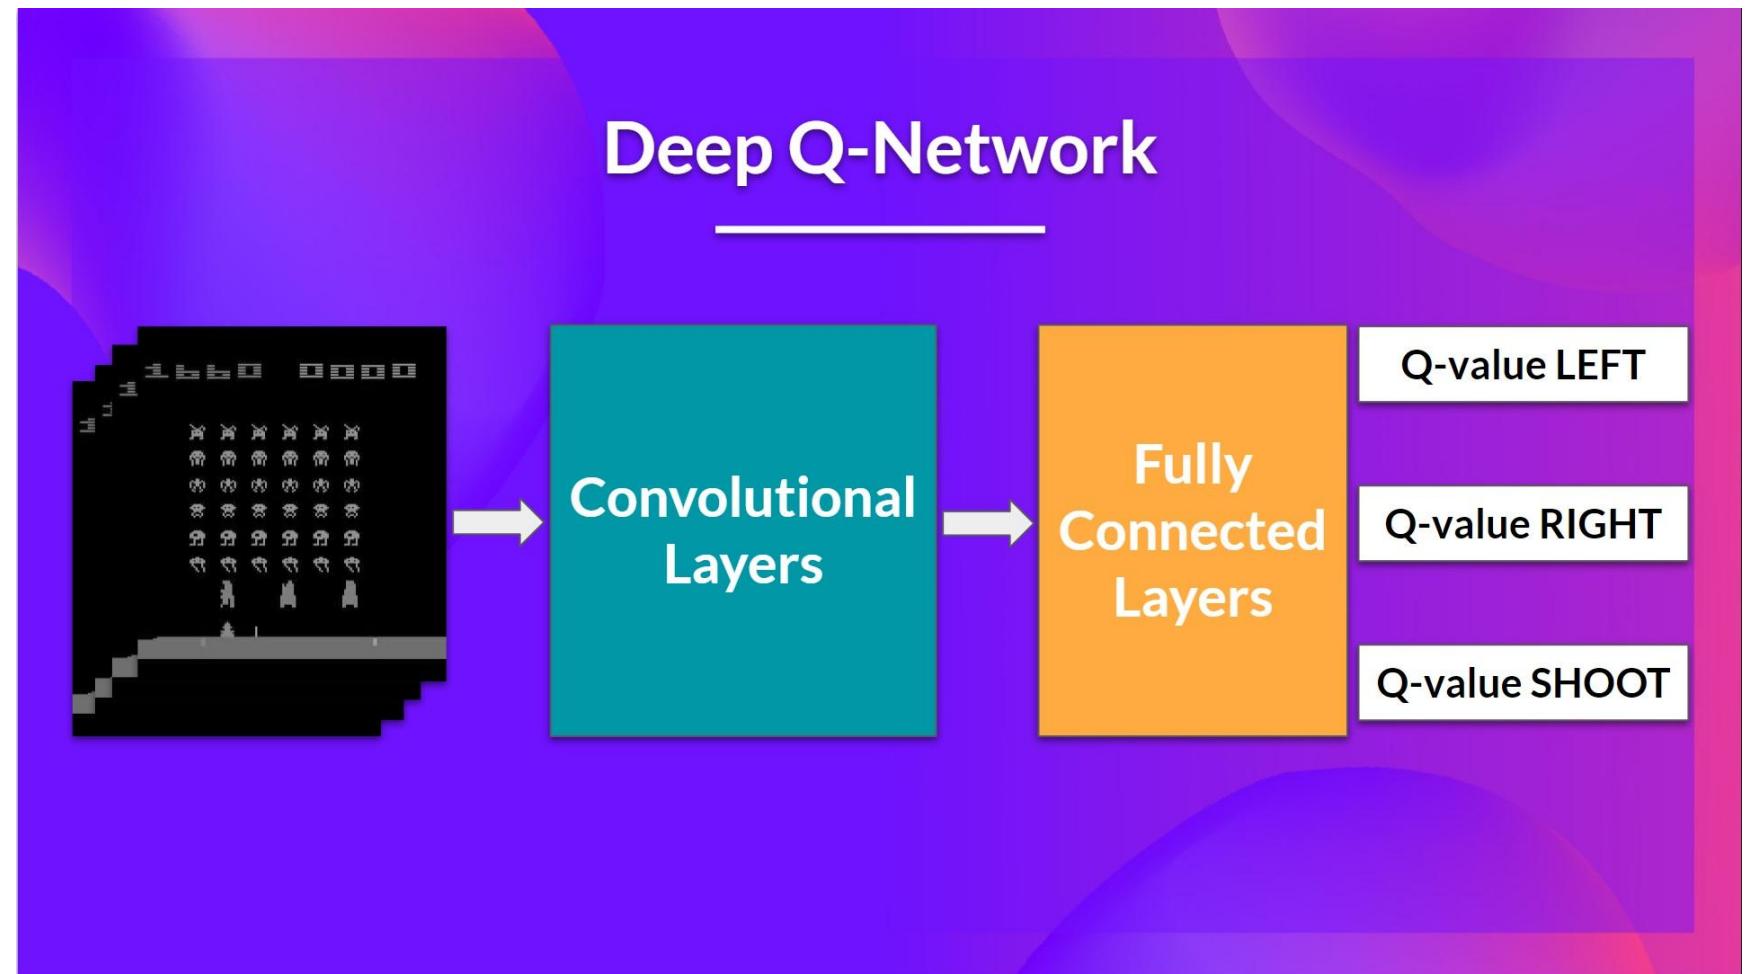
Playing Atari with Deep Reinforcement Learning ([Paper](#))

## Architecture

**Input:** Stack of 4 gray-scale frames

**Q:** Why?

**Output:** A vector of Q-values for each possible action at that state

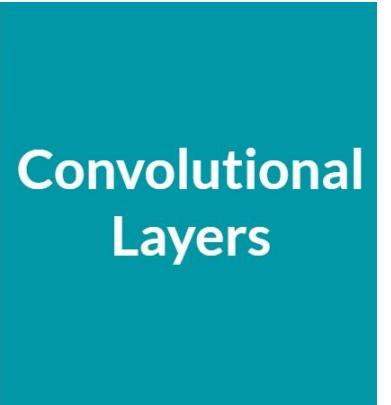


Playing Atari with Deep Reinforcement Learning ([Paper](#))

## Architecture: What Is a Convolutional Layer?

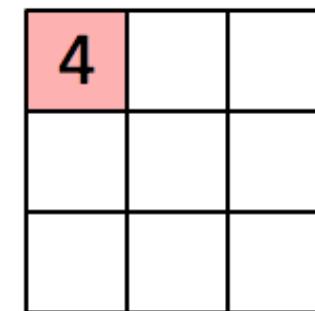
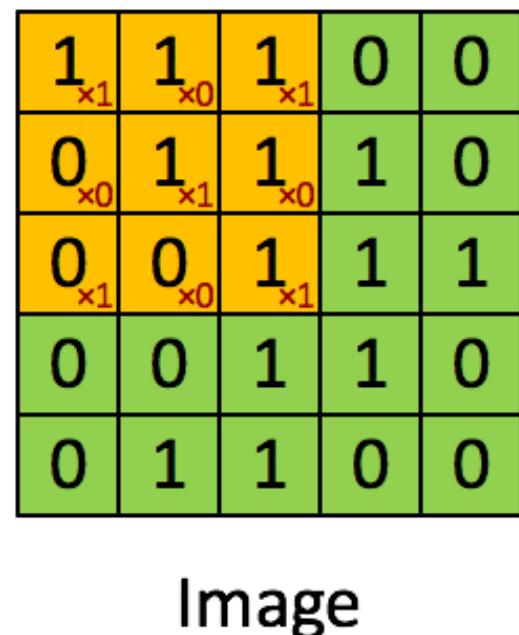
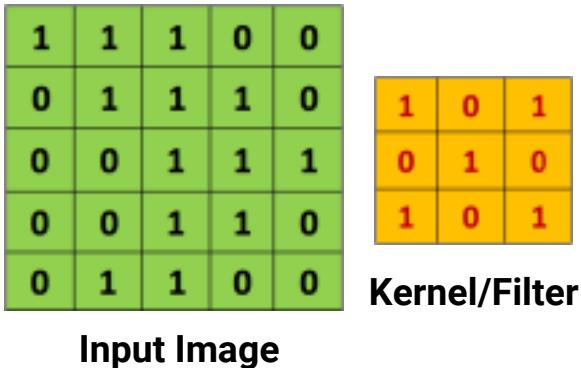


Input

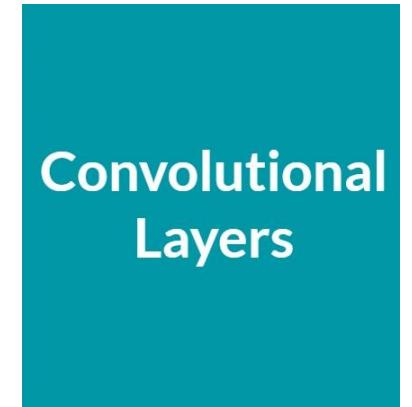


Playing Atari with Deep Reinforcement Learning ([Paper](#))

## Architecture: What Is a Convolutional Layer?

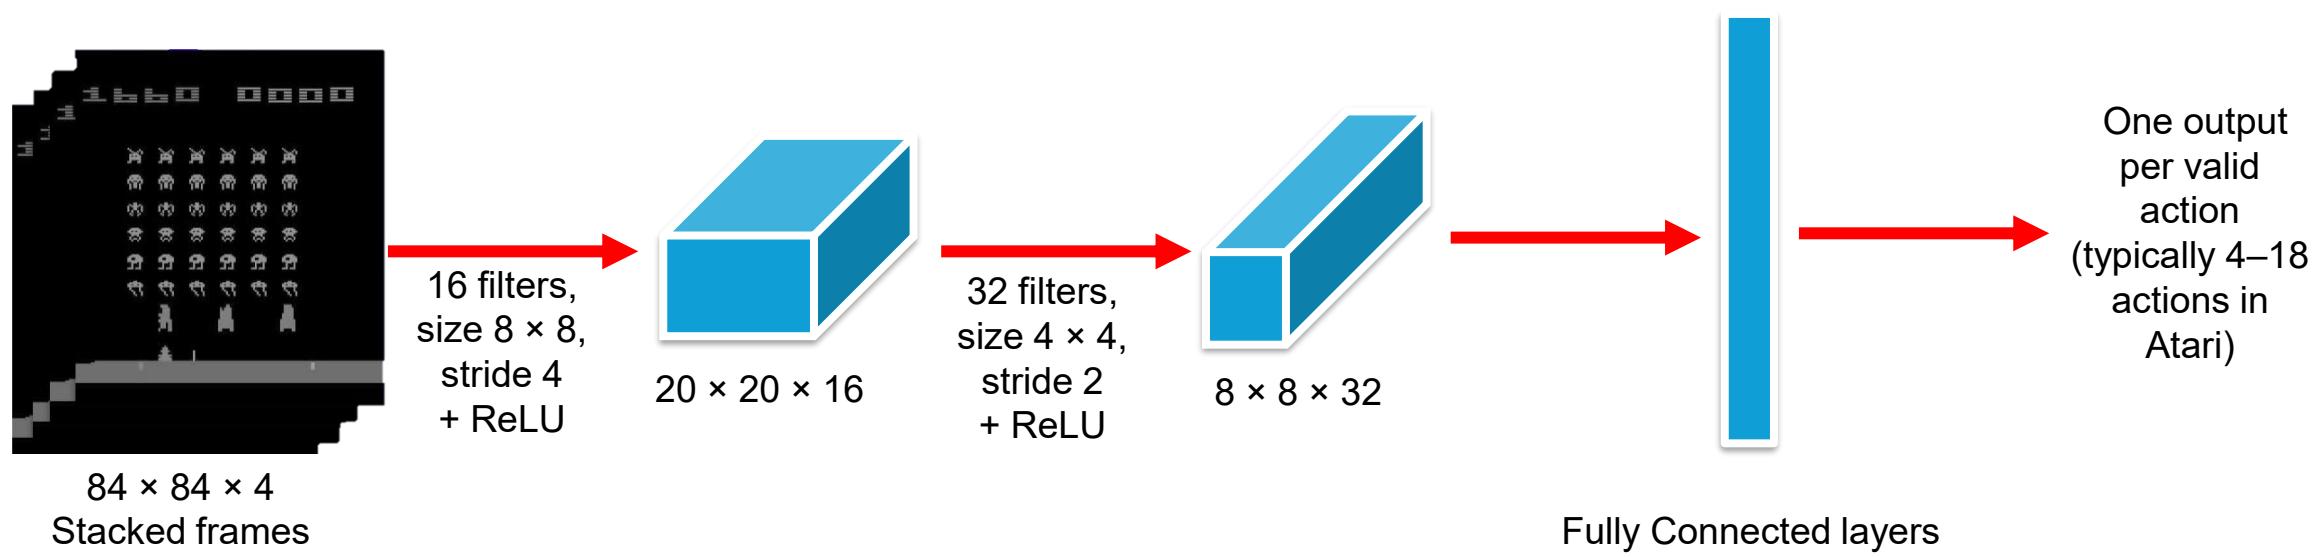


Convolved  
Feature



Playing Atari with Deep Reinforcement Learning ([Paper](#))

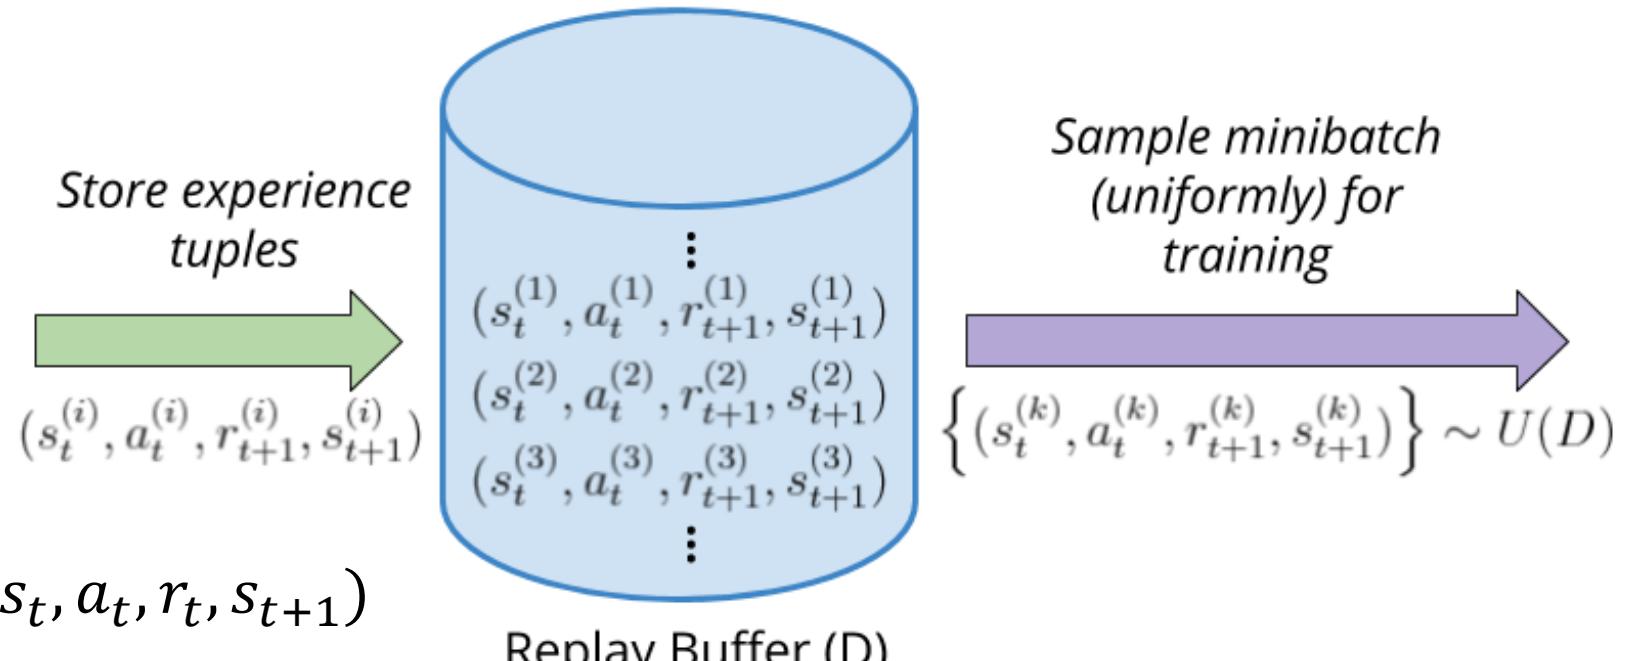
## Architecture



This architecture is known as a **Deep Q-Network (DQN)**

Playing Atari with Deep Reinforcement Learning ([Paper](#))

## Experience Replay Replay Buffer



Agent's experiences :  $e_t = (s_t, a_t, r_t, s_{t+1})$

Data-set  $D = e_1, \dots, e_N$

$Q$ -learning updates, or minibatch updates :

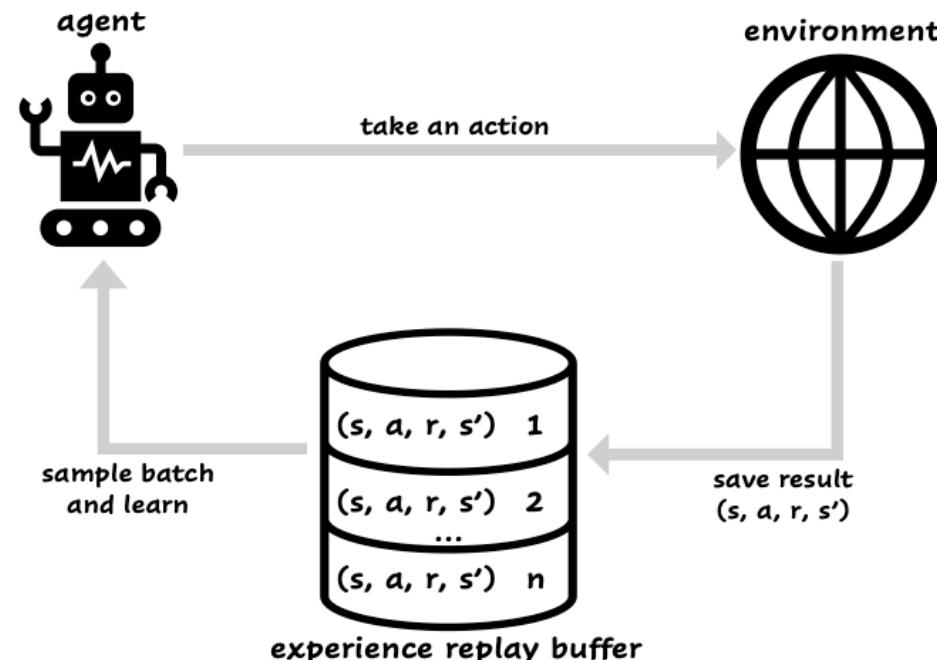
*Random* :  $e \sim D$

Playing Atari with Deep Reinforcement Learning ([Paper](#))

## Algorithm: Replay Buffer

*Stores the last  $N$  experience*

*Samples uniformly at random from  $D$*



**Q:** Why Replay Buffer?

Playing Atari with Deep Reinforcement Learning ([Paper](#))

# Algorithm

## Reminder Box

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

New  
Q-value  
estimation

Former  
Q-value  
estimation

Learning  
Rate

Immediate  
Reward

Discounted Estimate  
optimal Q-value  
of next state

Former  
Q-value  
estimation

TD Target

TD Error

Playing Atari with Deep Reinforcement Learning ([Paper](#))

# Algorithm

## Intuition

### Q-Target

$$y_j = r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-)$$

$$R_{t+1} + \gamma \max_a Q(S_{t+1}, a)$$

Immediate Reward      Discounted Estimate optimal Q-value of next state

TD Target

### Q-Loss

$$y_j - Q(\phi_j, a_j; \theta)$$

$$[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

Immediate Reward      Discounted Estimate optimal Q-value of next state

TD Target

TD Error

Former Q-value estimation

Playing Atari with Deep Reinforcement Learning ([Paper](#))

# Algorithm

---

**Algorithm 1** Deep Q-learning with Experience Replay
 

---

Initialize replay memory  $\mathcal{D}$  to capacity  $N$

Initialize action-value function  $Q$  with random weights

**for** episode = 1,  $M$  **do**

  Initialise sequence  $s_1 = \{x_1\}$  and preprocessed sequenced  $\phi_1 = \phi(s_1)$

**for**  $t = 1, T$  **do**

    With probability  $\epsilon$  select a random action  $a_t$

    otherwise select  $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

    Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$

    Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$

    Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $\mathcal{D}$

    Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $\mathcal{D}$

    Set  $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

    Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  according to equation 3

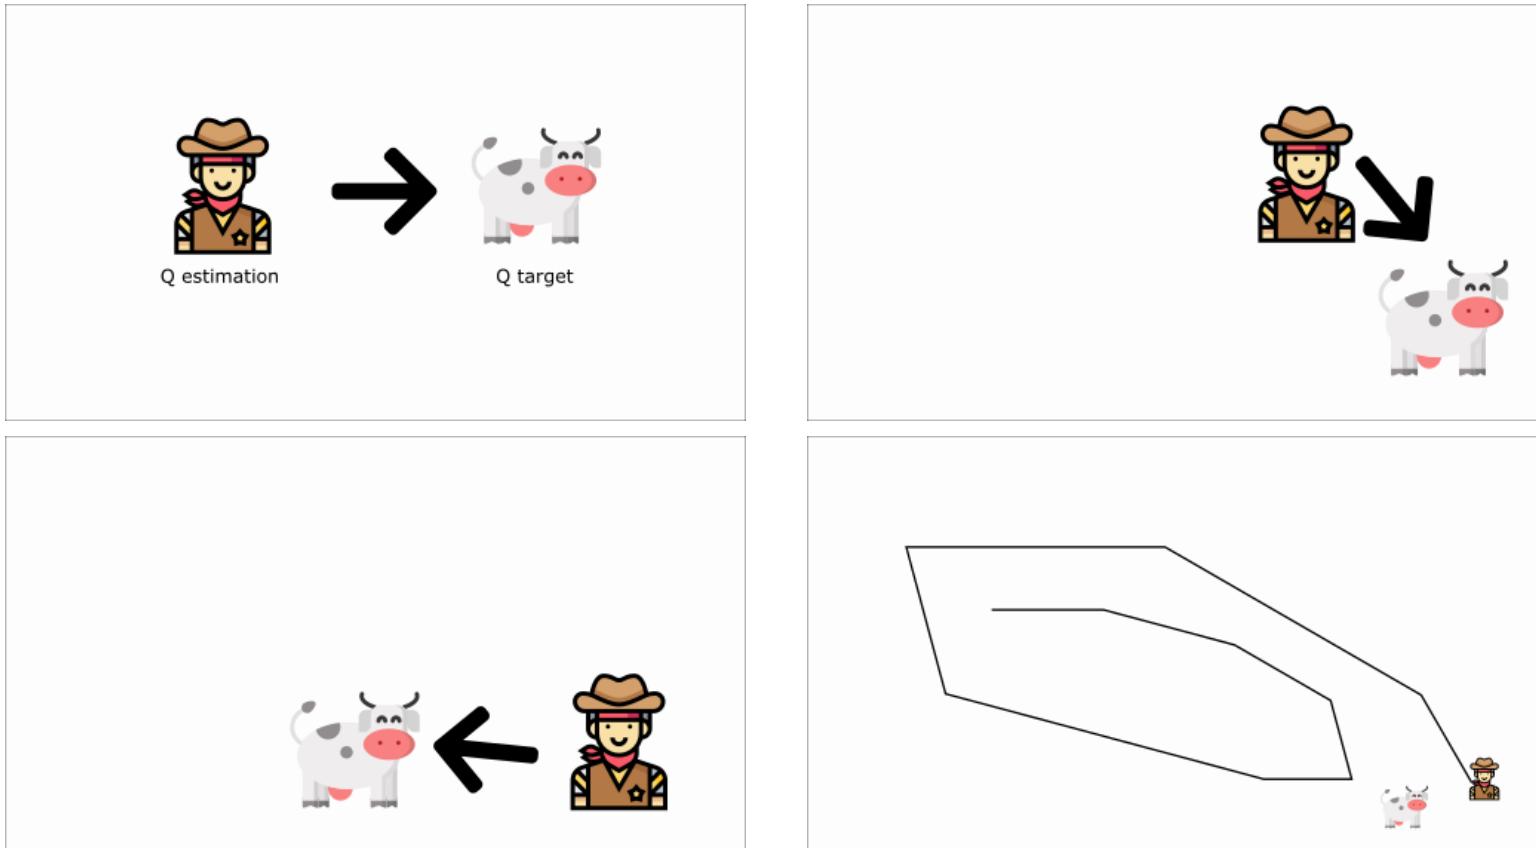
**end for**

**end for**

---

Playing Atari with Deep Reinforcement Learning ([Paper](#))

## Algorithm: A Challenge!



Playing Atari with Deep Reinforcement Learning ([Paper](#))

## Algorithm: A Challenge!

### Solution:

- Use a **separate network** with fixed parameters for estimating the TD Target
- Copy the parameters from our Deep Q-Network **every C steps** to update the target network.

Playing Atari with Deep Reinforcement Learning ([Paper](#), [Paper](#))

# Algorithm

## Algorithm 1: deep Q-learning with experience replay.

```

Initialize replay memory  $D$  to capacity  $N$ 
Initialize action-value function  $Q$  with random weights  $\theta$ 
Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- = \theta$ 
For episode = 1,  $M$  do
    Initialize sequence  $s_1 = \{x_1\}$  and preprocessed sequence  $\phi_1 = \phi(s_1)$ 
    For  $t = 1, T$  do
        With probability  $\varepsilon$  select a random action  $a_t$ 
        otherwise select  $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$ 
        Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$ 
        Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$ 
        Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $D$ 
        Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $D$ 
        Set  $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$ 
        Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  with respect to the
        network parameters  $\theta$ 
        Every  $C$  steps reset  $\hat{Q} = Q$ 
    End For
End For

```

Playing Atari with Deep Reinforcement Learning ([Paper](#), [Paper](#))

## Conclusion

“We apply our method to seven Atari 2600 games from the Arcade Learning Environment, with no adjustment of the architecture or learning algorithm. We find that it outperforms all previous approaches on six of the games and surpasses a human expert on three of them, with no adjustment of the architecture or hyperparameters.”

Playing Atari with Deep Reinforcement Learning ([Paper](#), [Paper](#))