

Reinforcement Learning in Control

Dr. Saeed Shamaghdari

**Electrical Engineering Department
Control Group**

Fall 2025 | 4041

Deep Reinforcement Learning

Playing Atari with Deep Reinforcement Learning

Volodymyr Mnih Koray Kavukcuoglu David Silver Alex Graves Ioannis Antonoglou

Daan Wierstra Martin Riedmiller

DeepMind Technologies

{vlad, koray, david, alex.graves, ioannis, daan, martin.riedmiller} @ deepmind.com

[☆ Save](#) [✉ Cite](#) [Cited by 19338](#) [Related articles](#) [All 37 versions](#) [»](#)

Human-level control through deep reinforcement learning

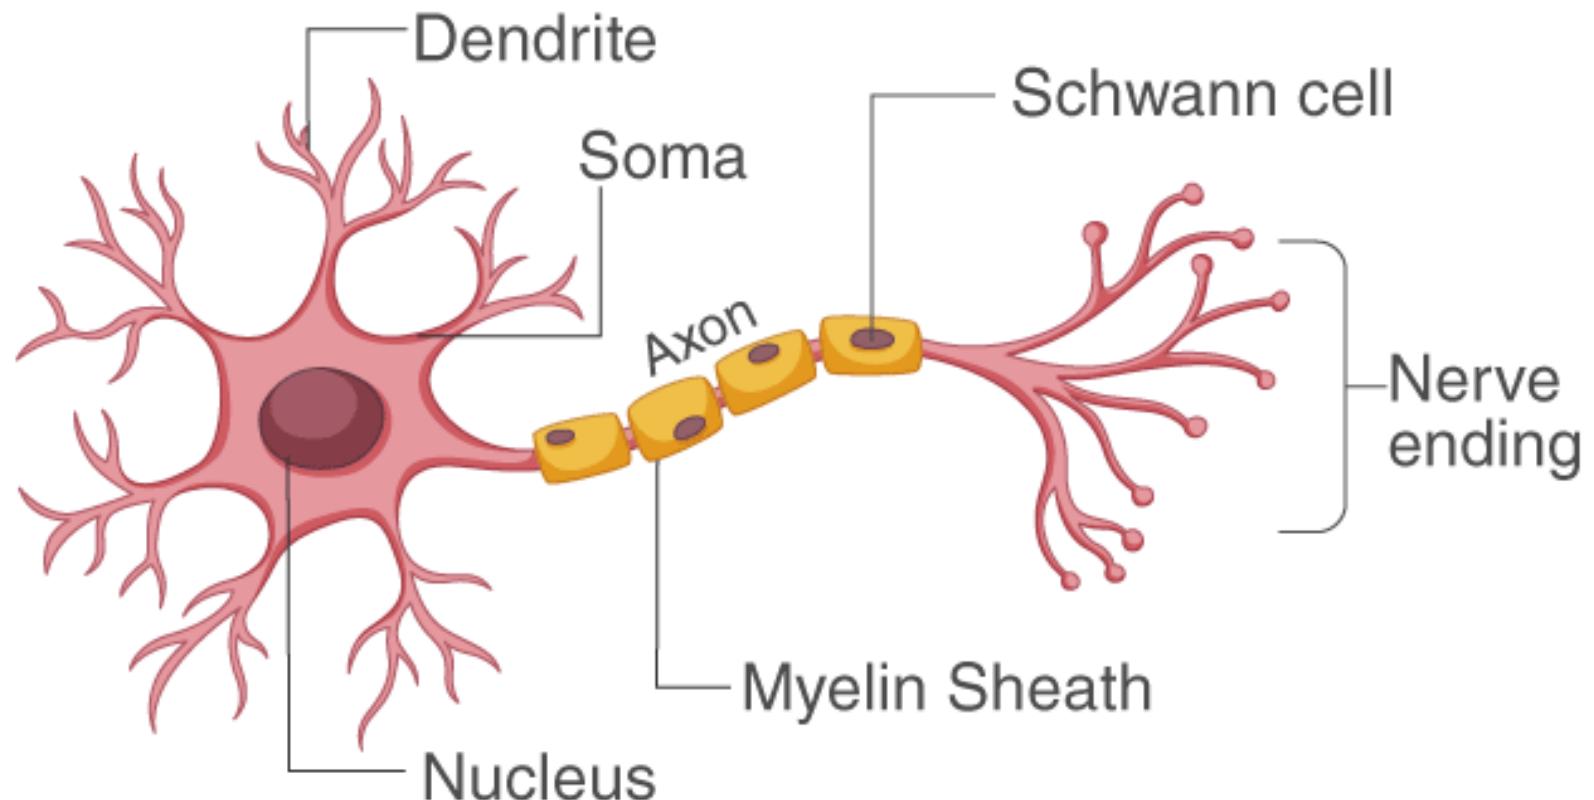
Volodymyr Mnih^{1*}, Koray Kavukcuoglu^{1*}, David Silver^{1*}, Andrei A. Rusu¹, Joel Veness¹, Marc G. Bellemare¹, Alex Graves¹, Martin Riedmiller¹, Andreas K. Fidjeland¹, Georg Ostrovski¹, Stig Petersen¹, Charles Beattie¹, Amir Sadik¹, Ioannis Antonoglou¹, Helen King¹, Dharshan Kumaran¹, Daan Wierstra¹, Shane Legg¹ & Demis Hassabis¹

¹Google DeepMind, 5 New Street Square, London EC4A 3TW, UK.

*These authors contributed equally to this work.

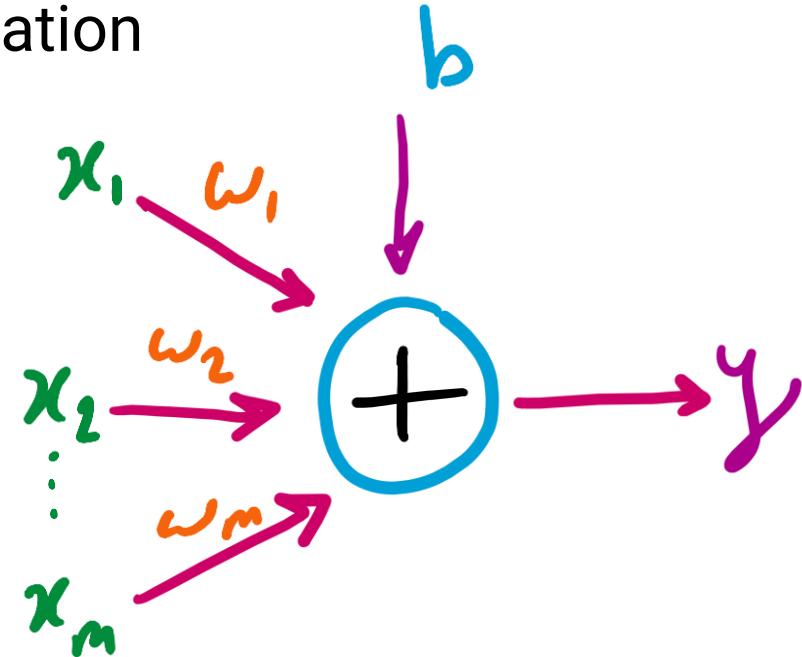
Neural Networks and Deep Learning: A Simple Review

Neuron



Neural Networks and Deep Learning: A Simple Review

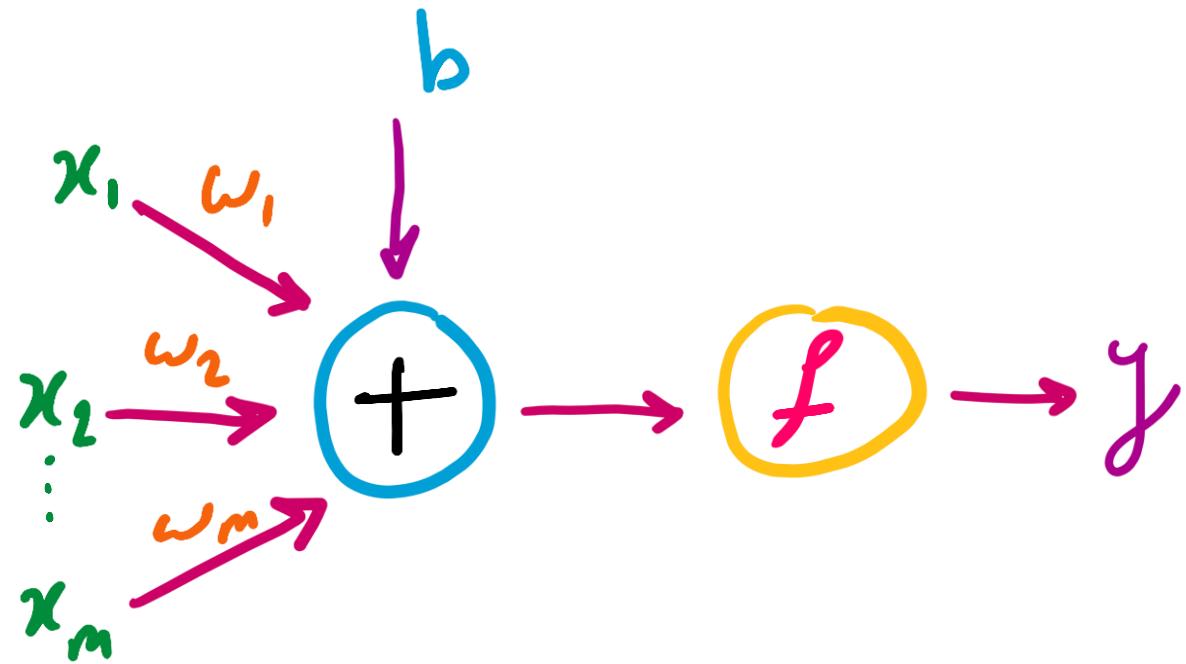
Artificial Neuron Formulation



$$y = w_1x_1 + w_2x_2 + \dots + w_mx_m + b$$

Neural Networks and Deep Learning: A Simple Review

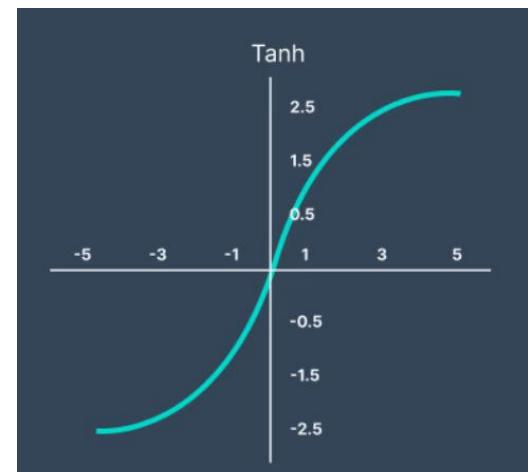
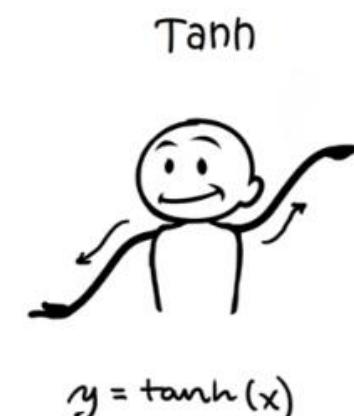
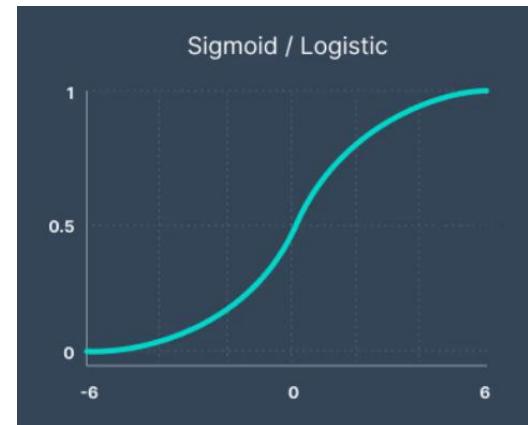
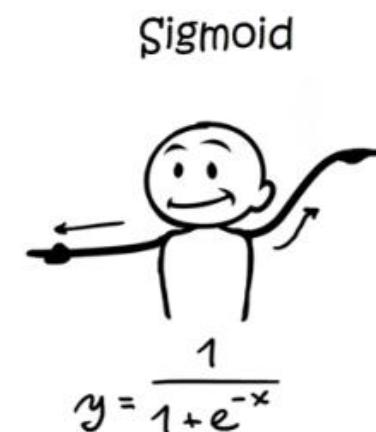
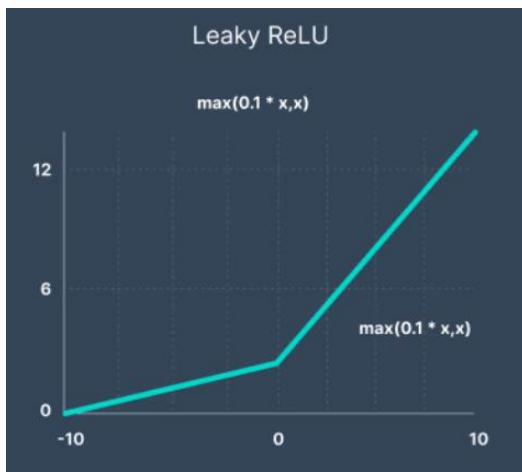
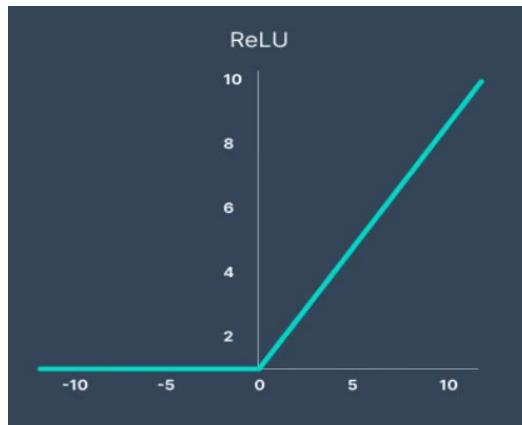
Q: Problems?



$$y = f(w_1x_1 + w_2x_2 + \dots + w_mx_m + b) = f(\mathbf{w}^\top \mathbf{x} + b)$$

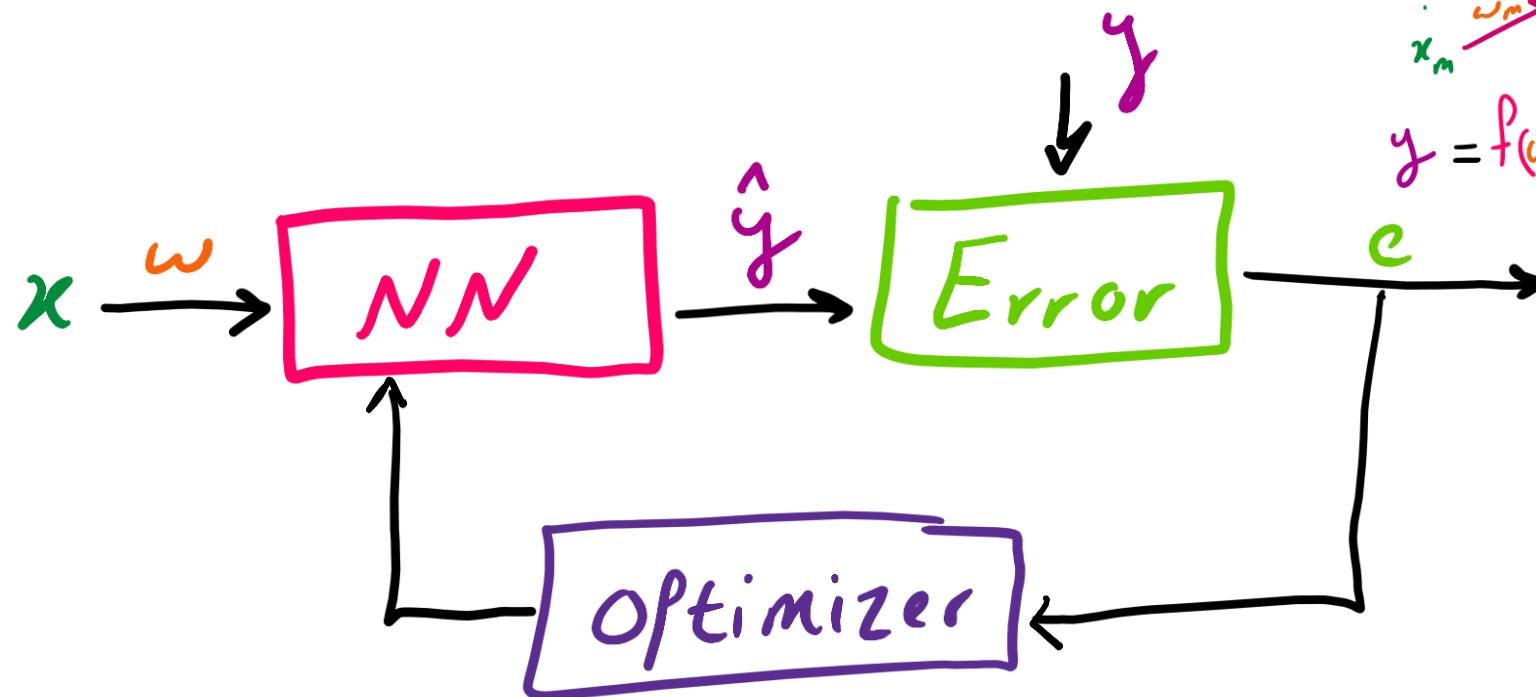
Neural Networks and Deep Learning: A Simple Review

Activation Functions



Neural Networks and Deep Learning: A Simple Review

Learning Block Diagram



A detailed diagram of a single neuron model. The input layer consists of nodes x_1, x_2, \dots, x_m (green). These inputs are multiplied by weights w_1, w_2, \dots, w_m (orange) and summed up by a blue circle containing a summation symbol (Σ). The sum is then passed through an activation function f (yellow) to produce the output y (purple). The bias term b (blue) is also shown being added to the weighted sum.

$$y = f(w_1x_1 + w_2x_2 + \dots + w_mx_m + b)$$

Neural Networks and Deep Learning: A Simple Review

Learning Block Diagram

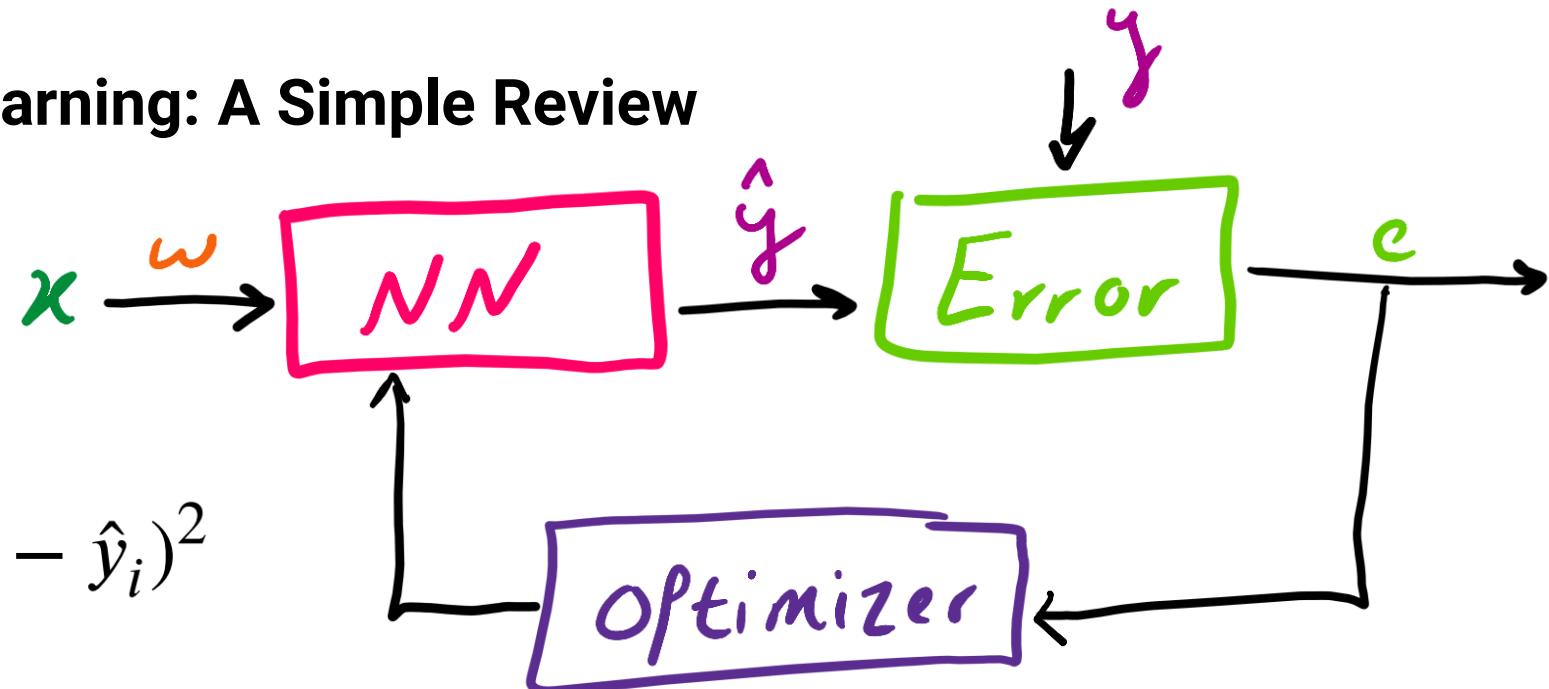
Loss Function:

MSE:

$$L(y, \hat{y}) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

MAE:

$$L(y, \hat{y}) = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$



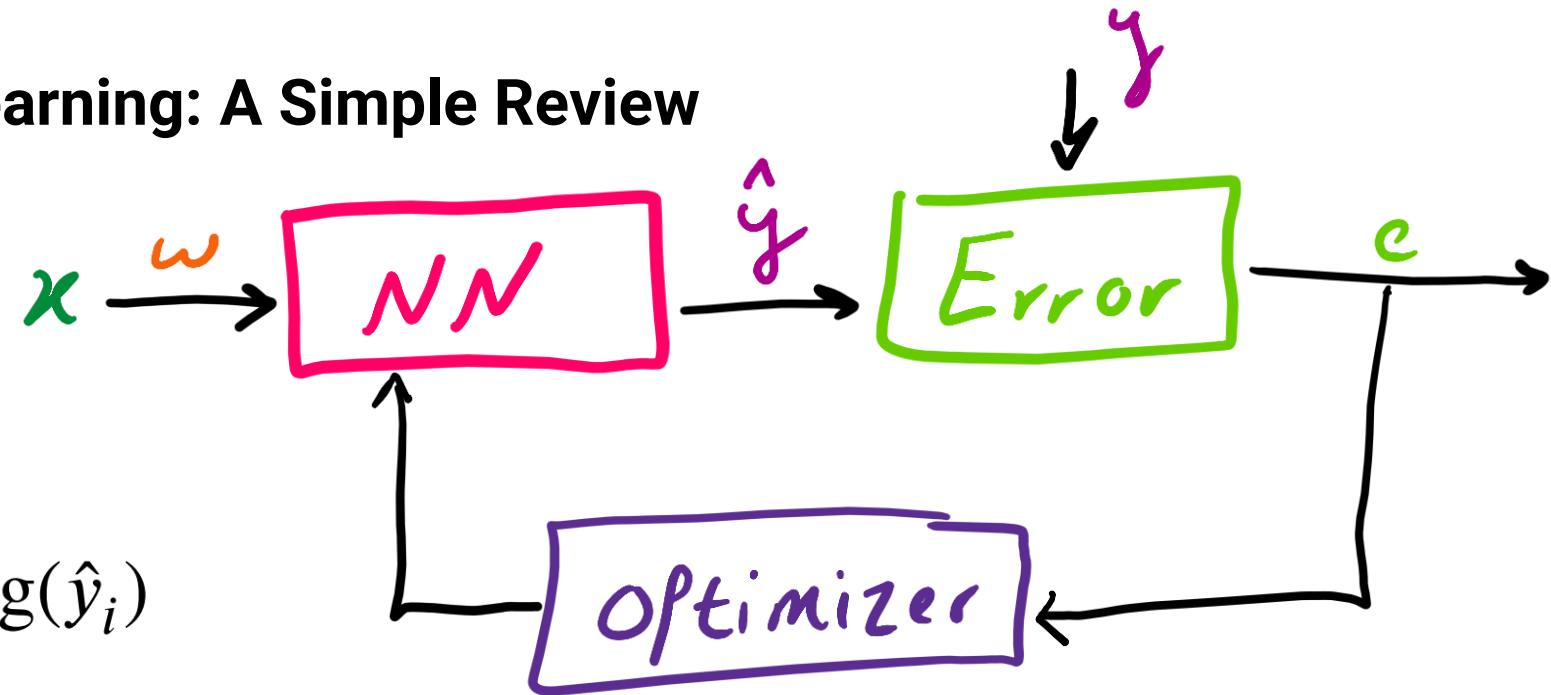
Neural Networks and Deep Learning: A Simple Review

Learning Block Diagram

Loss Function:

Cross Entropy:

$$L(y, \hat{y}) = - \sum_{i=1}^N y_i \log(\hat{y}_i)$$

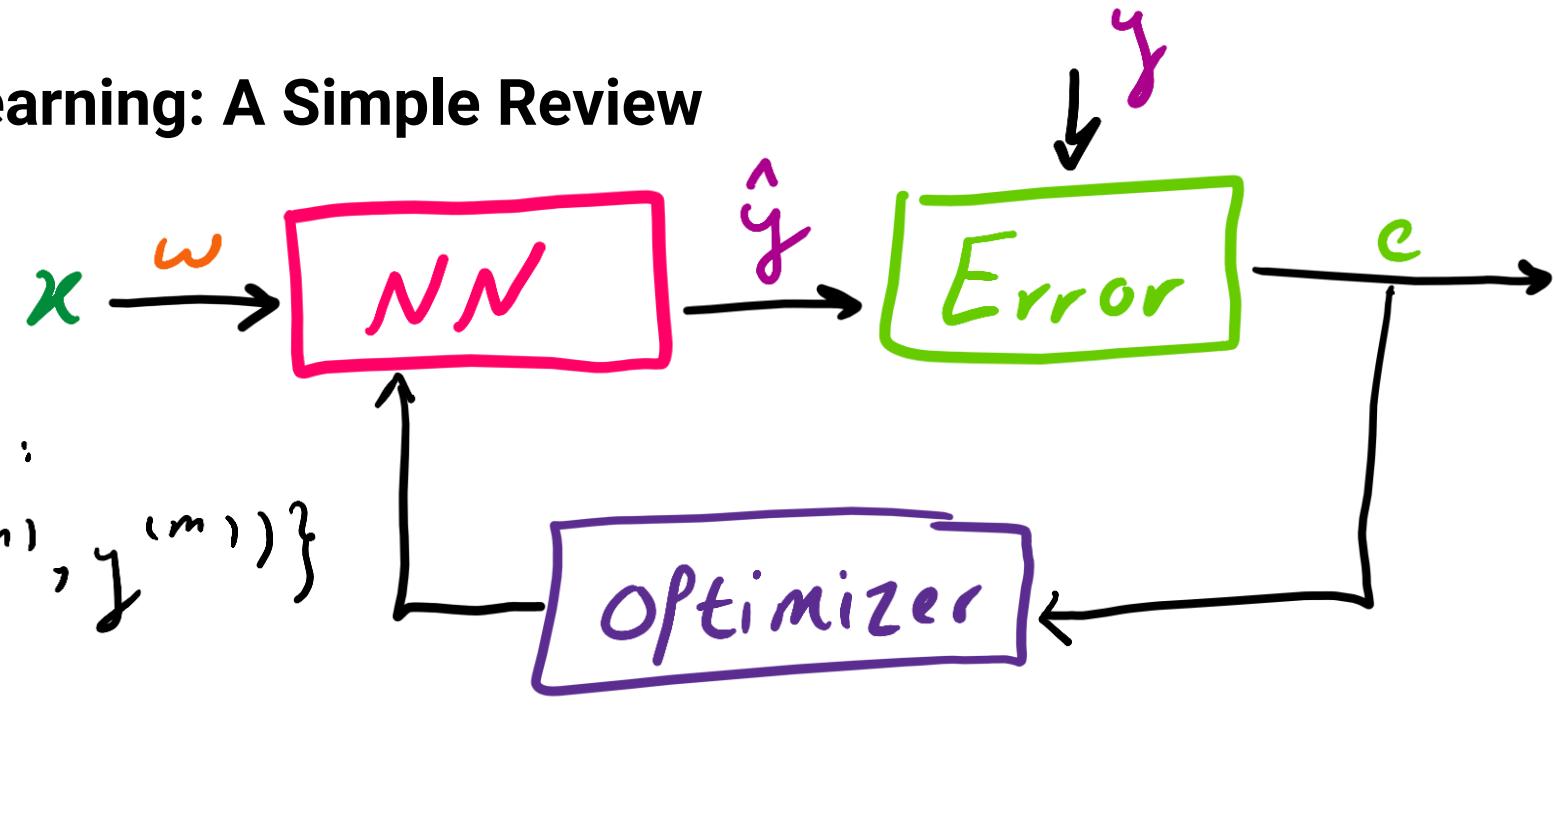


Binary Cross Entropy:

$$L(y, \hat{y}) = -(y \log(\hat{y}) + (1 - y) \log(1 - \hat{y}))$$

Neural Networks and Deep Learning: A Simple Review

Learning Block Diagram



Given m train examples:

$$\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$$

want $\hat{y}^{(i)} \approx y^{(i)}$

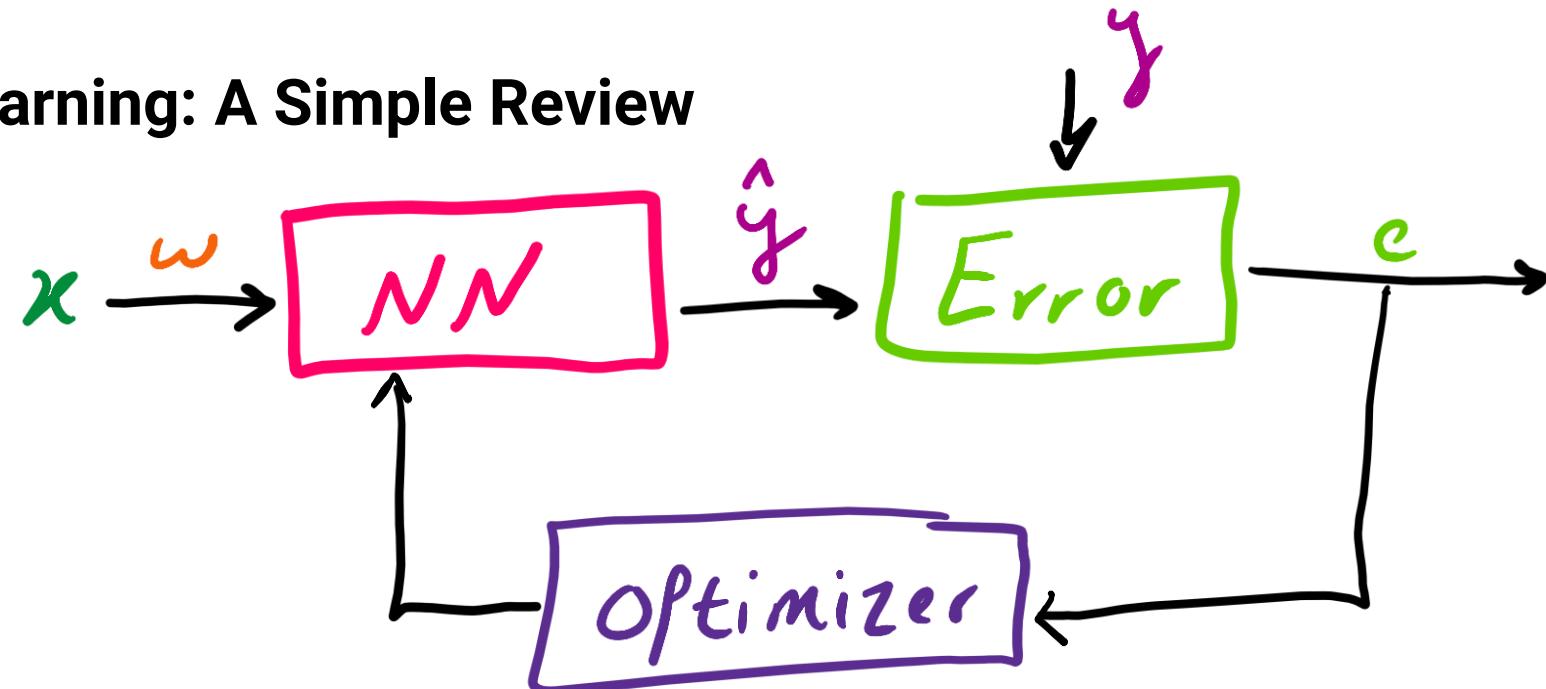
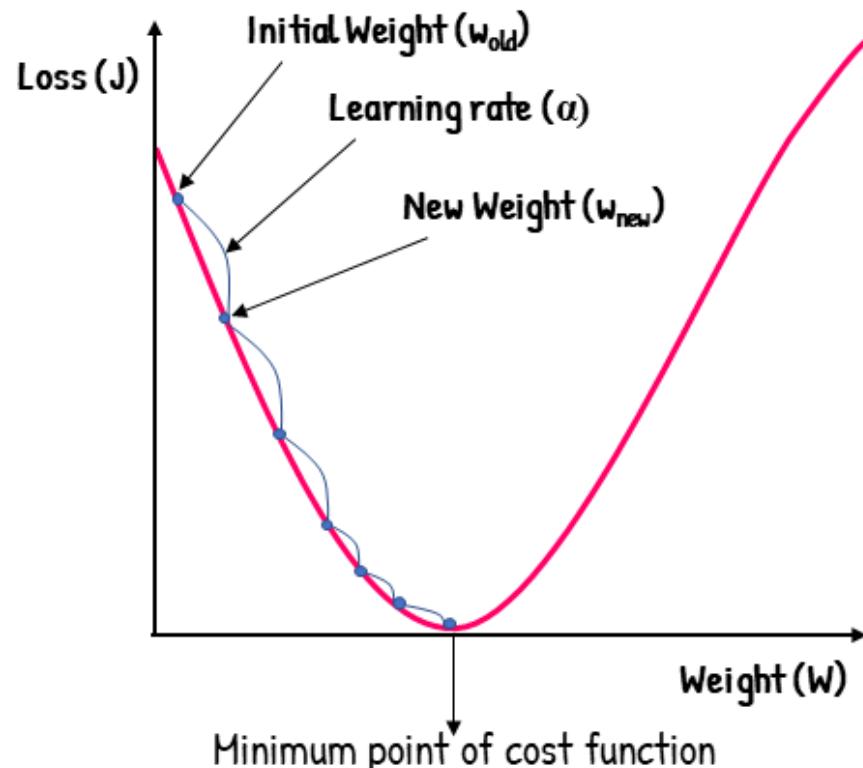
So the cost function is:

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m L(y, \hat{y})$$

Neural Networks and Deep Learning: A Simple Review

Learning Block Diagram

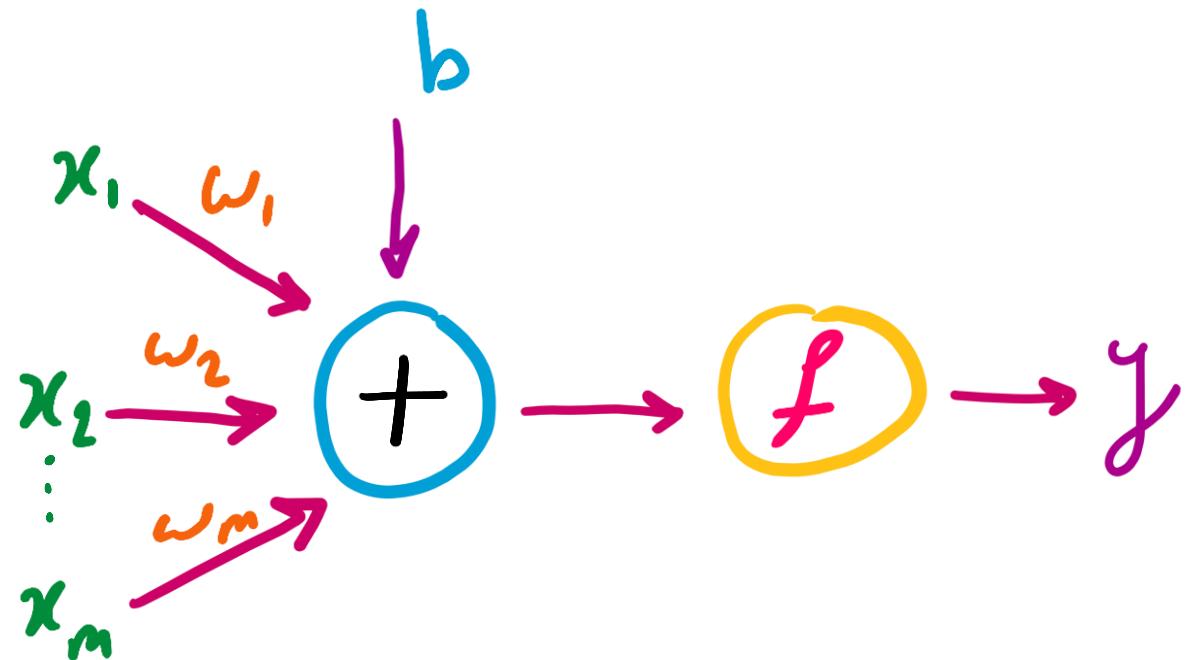
Gradient Descent:



$$w_{new} = w_{old} - \alpha \frac{\delta J}{\delta w}$$

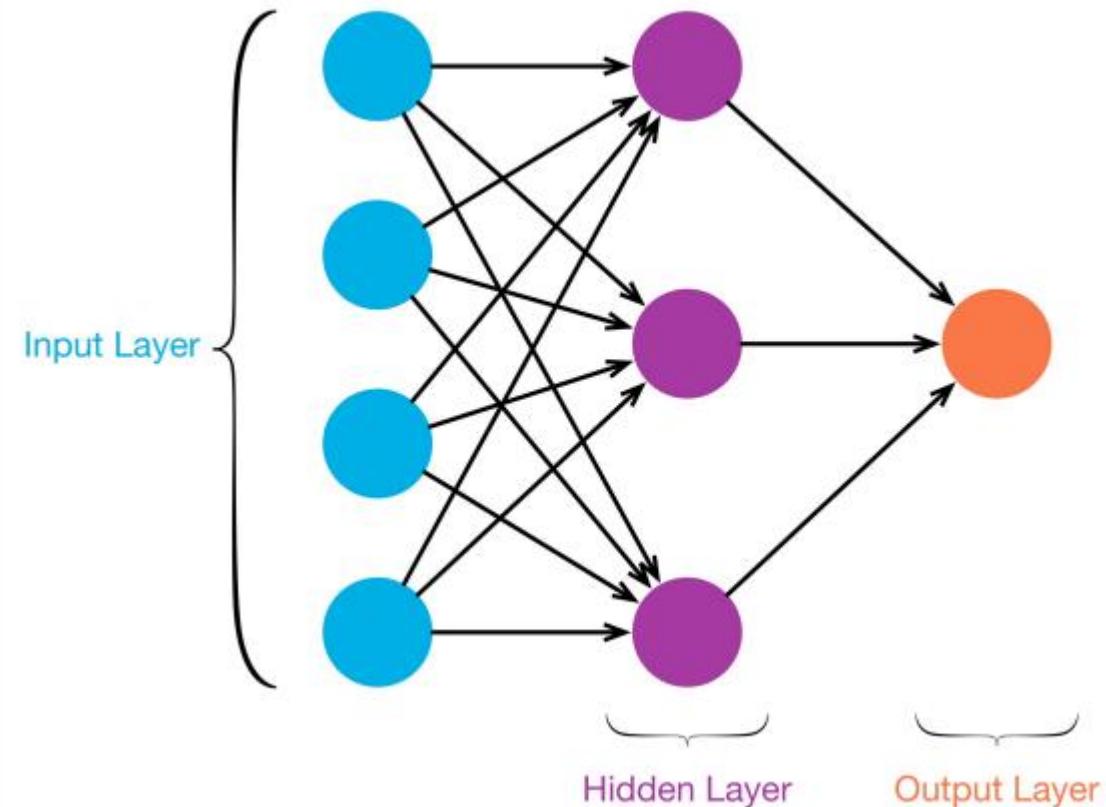
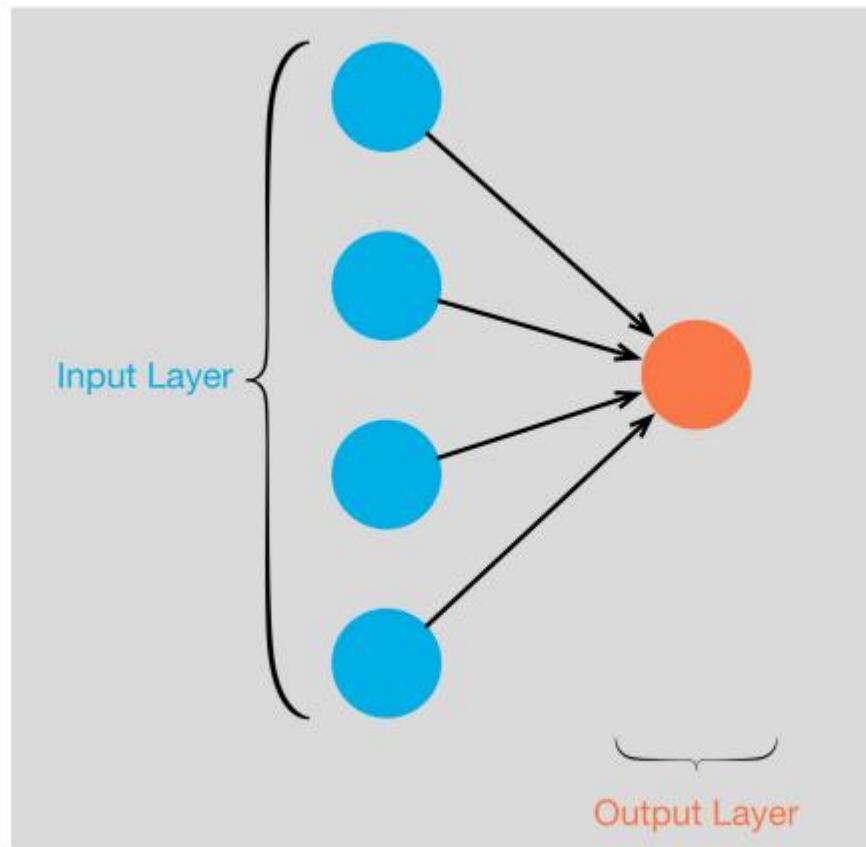
Neural Networks and Deep Learning: A Simple Review

Single-Layer Perceptron

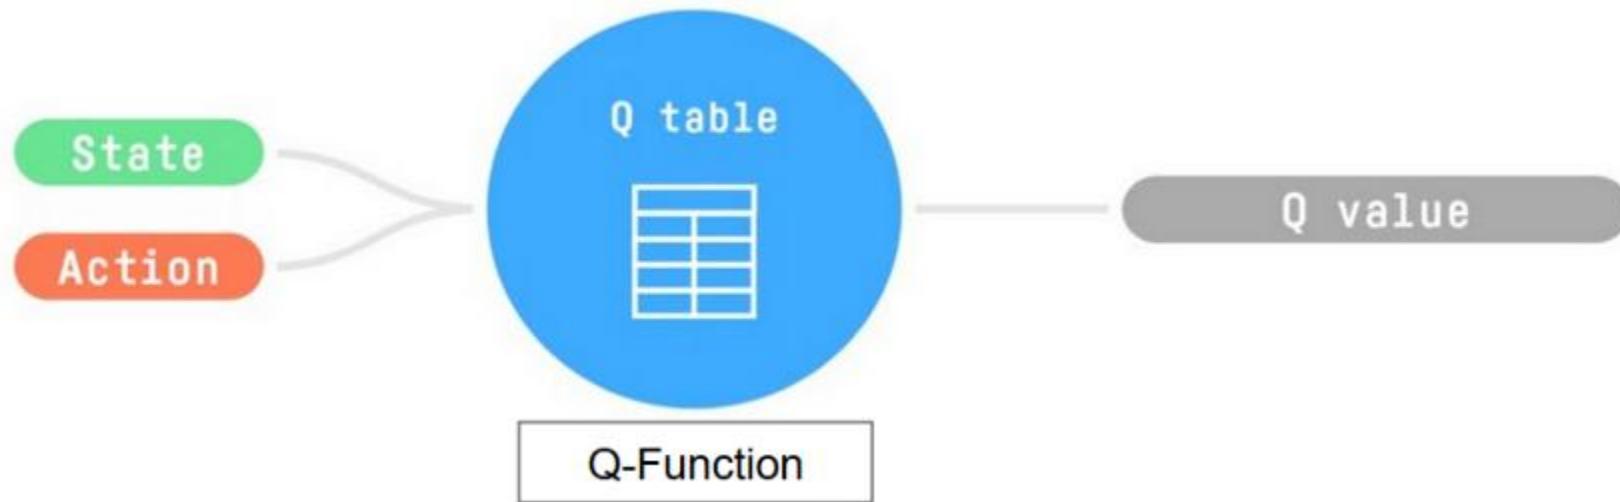


Neural Networks and Deep Learning: A Simple Review

Multi-Layer Perceptron ([MLP](#)) ([Play!](#))



Q-Learning Recap ...



Pseudocode

Q-Learning

Algorithm 14: Sarsamax (Q-Learning)

Input: policy π , positive integer $num_episodes$, small positive fraction α , GLIE $\{\epsilon_i\}$

Output: value function Q ($\approx q_\pi$ if $num_episodes$ is large enough)

Initialize Q arbitrarily (e.g., $Q(s, a) = 0$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$, and $Q(\text{terminal-state}, \cdot) = 0$)

for $i \leftarrow 1$ **to** $num_episodes$ **do** → Step 1

$\epsilon \leftarrow \epsilon_i$

 Observe S_0

$t \leftarrow 0$

repeat

 Choose action A_t using policy derived from Q (e.g., ϵ -greedy) Step 2

 Take action A_t and observe R_{t+1}, S_{t+1} Step 3

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t))$ Step 4

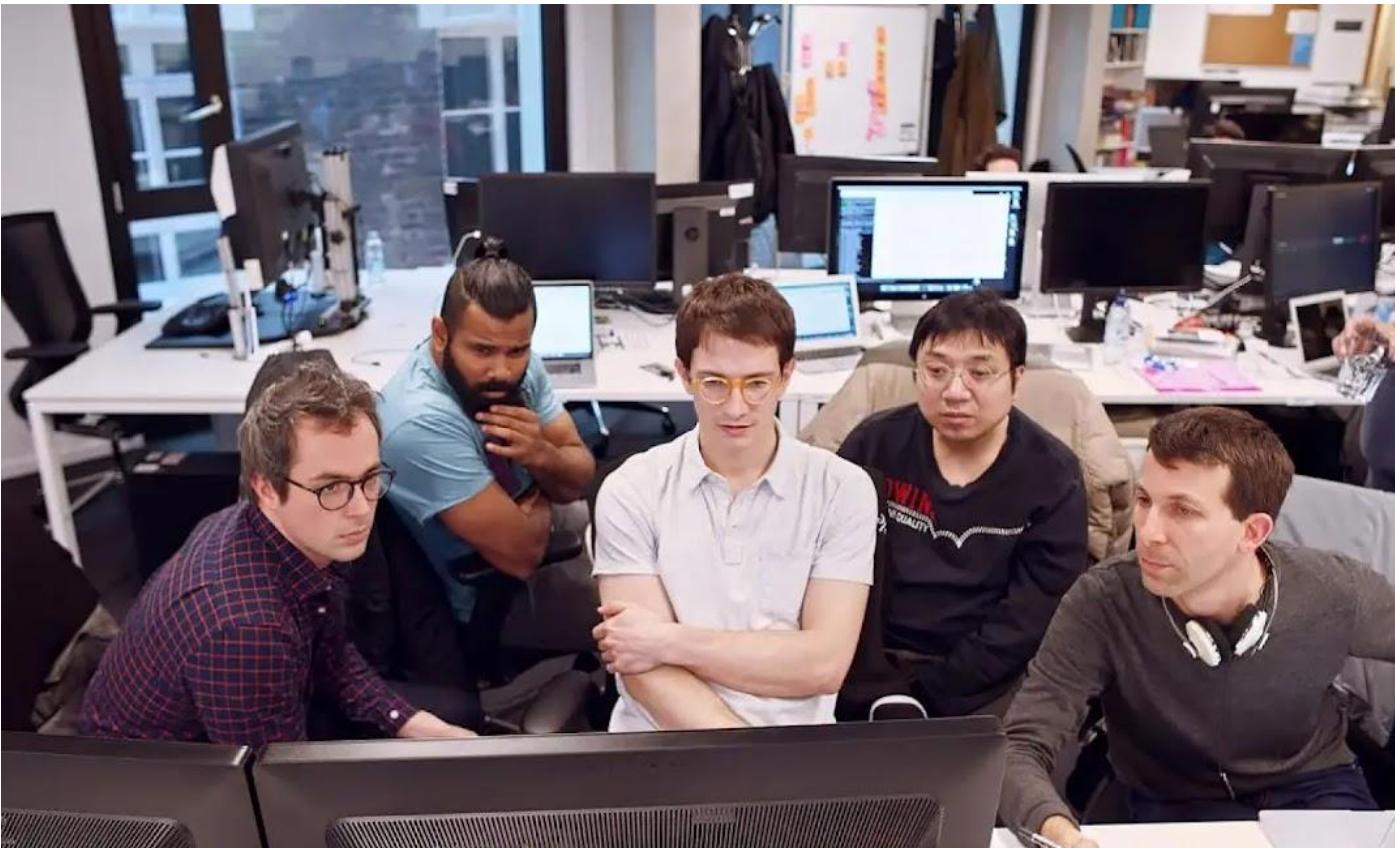
$t \leftarrow t + 1$

until S_t is terminal;

end

return Q

Why Deep Reinforcement Learning?



Playing Atari with Deep Reinforcement Learning

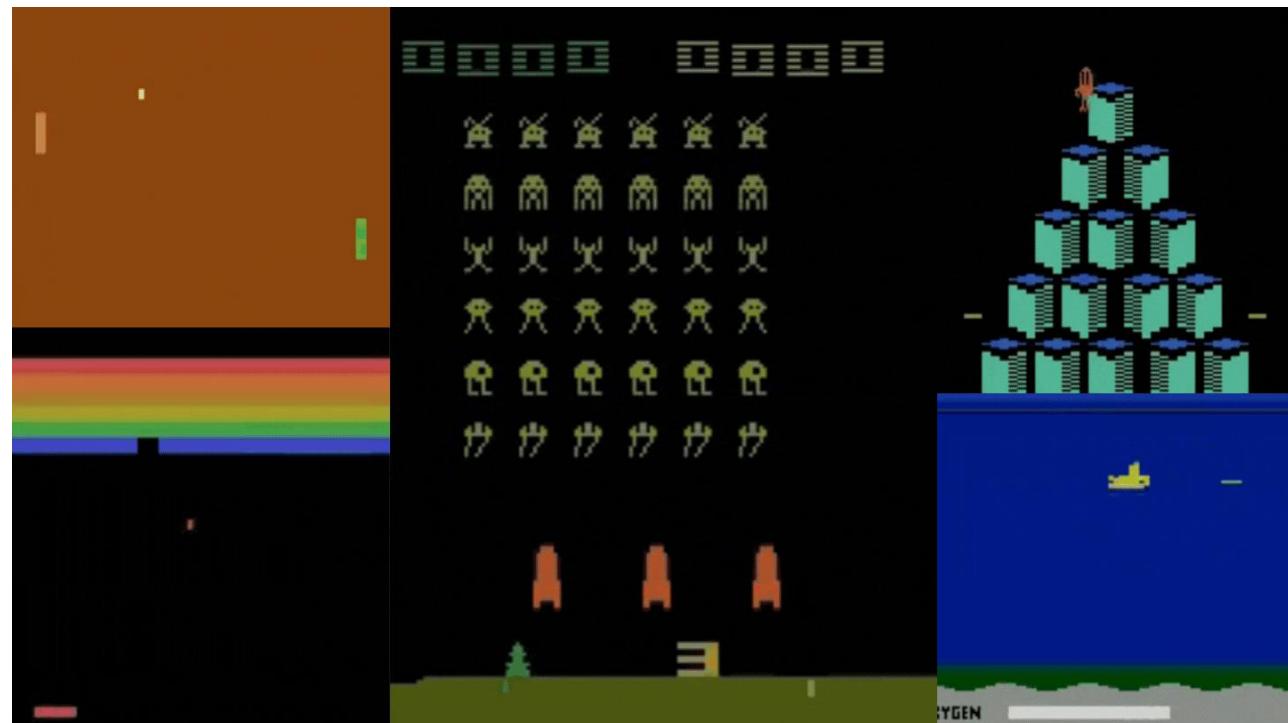
([Paper](#))

RL in Control | IUST

Why Deep Reinforcement Learning?

*“Learning to control agents directly from high-dimensional sensory inputs like **vision** and **speech** is one of the long-standing challenges of reinforcement learning (RL).”*

Q: Number of states in an 8*8 Gridworld?
What about an **Atari** game?



Playing Atari with Deep Reinforcement Learning ([Paper](#))

Why Deep Reinforcement Learning?



(Left-to-right) Pong, Breakout, Space Invaders, Seaquest, Beam Rider

*“Atari 2600 is a challenging RL testbed that presents agents with a high dimensional visual input (**210 × 160 RGB** video at 60Hz) and a diverse and interesting set of tasks that were designed to be difficult for humans players. “*

Playing Atari with Deep Reinforcement Learning ([Paper](#))

Why Deep Reinforcement Learning?



(Left-to-right) Pong, Breakout, Space Invaders, Seaquest, Beam Rider

Each Frame: (210, 160, 3) containing values ranging from 0 to 255

Q: Number of states?

A: $256^{210 \times 160 \times 3} = 256^{100800}$

In this case, the best idea is to approximate the Q-values using a parametrized Q-function $Q_\theta(s, a)$.

Playing Atari with Deep Reinforcement Learning ([Paper](#))

DeepRL: Why Is It Still Hard?

Firstly, most successful deep learning applications to date have required large amounts of hand-labelled training data.

- RL algorithms, on the other hand, must be able to learn from a scalar reward signal that is frequently sparse, noisy and delayed.

“The delay between actions and resulting rewards, which can be thousands of timesteps long, seems particularly daunting when compared to the direct association between inputs and targets found in supervised learning.”

DeepRL: Why Is It Still Hard?

Another issue is that most deep learning algorithms assume the data samples to be **independent (IID)**, while in reinforcement learning one typically encounters sequences of **highly correlated** states.

Furthermore, in RL the data distribution changes as the algorithm learns new behaviours, which can be problematic for deep learning methods that assume a fixed underlying distribution.

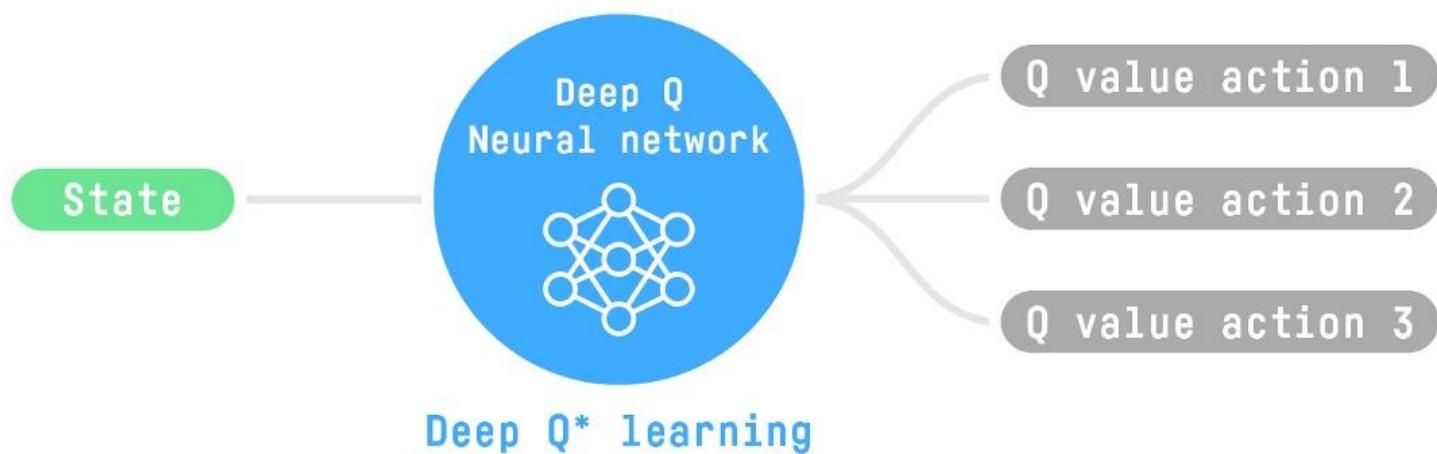
Playing Atari with Deep Reinforcement Learning ([Paper](#))

Deep Q-Learning (DQN)

This neural network will approximate, given a state, the different Q-values for each possible action at that state.

And that's exactly what Deep Q-Learning does.

The best idea is to approximate the Q-values using a parametrized Q-function $Q_\theta(s, a)$.



Playing Atari with Deep Reinforcement Learning ([Paper](#))

Deep Q-Learning (DQN)

“We consider tasks in which an agent interacts with an environment \mathcal{E} , in this case the Atari emulator.”

Goal: Interact with the emulator \mathcal{E} by selecting actions in a way that maximizes future rewards.

The future discounted *return* at time t :

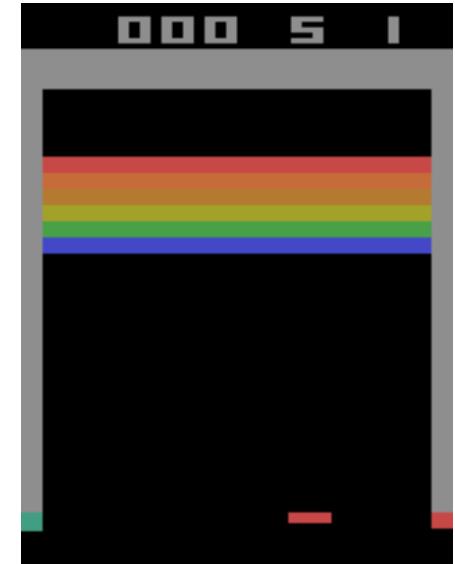
$$R_t = \sum_{t'=t}^T \gamma^{t'-t} r_{t'}$$

r_t : reward

The optimal action-value function:

$$Q^*(s, a) = \max_{\pi} \mathbb{E}[R_t | s_t = s, a_t = a, \pi]$$

Playing Atari with Deep Reinforcement Learning ([Paper](#))



Deep Q-Learning (DQN)

Intuition: If the optimal value $Q^*(s', a')$ of the sequence s' at the next time-step was known for all possible actions a' , then the optimal strategy is to select the action a' maximizing the expected value of $r + \gamma Q^*(s', a')$

$$Q^*(s, a) = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q^*(s', a') \middle| s, a \right]$$

Reminder Box: Value Iteration

$$Q_{i+1}(s, a) = \mathbb{E} [r + \gamma \max_{a'} Q_i(s', a') | s, a]$$
$$Q_i \rightarrow Q^* \text{ as } i \rightarrow \infty$$

Playing Atari with Deep Reinforcement Learning ([Paper](#))

Deep Q-Learning (DQN)

A function approximator to estimate the action-value function:

$$Q(s, a) \approx \text{A Neural Network!}$$

Playing Atari with Deep Reinforcement Learning ([Paper](#))

Deep Q-Learning (DQN)

“We refer to a neural network function approximator with weights θ as a Q -network. A Q -network can be trained by minimizing a sequence of loss functions $L_i(\theta_i)$ that changes at each iteration i :”

$$L_i(\theta_i) = \mathbb{E}_{s,a \sim \rho(\cdot)} \left[(y_i - Q(s, a; \theta_i))^2 \right]$$

“Where

$$y_i = \mathbb{E}_{s' \sim \mathcal{E}} [r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) | s, a]$$

*is the **target** for iteration i , and $\rho(s, a)$ is a probability distribution over sequences s and actions a that we refer to as the behaviour distribution.”*

Note: The targets depend on the network weights; this is in contrast with the targets used for supervised learning, which are fixed before learning begins.

Playing Atari with Deep Reinforcement Learning ([Paper](#))

Deep Q-Learning (DQN)

“Differentiating the loss function with respect to the weights we arrive at the following gradient:”

$$\nabla_{\theta_i} L_i(\theta_i) = \mathbb{E}_{s,a \sim \rho(\cdot); s' \sim \mathcal{E}} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i) \right) \nabla_{\theta_i} Q(s, a; \theta_i) \right]$$

“Rather than computing the full expectations in the above gradient, it is often computationally expedient to optimize the loss function by stochastic gradient descent.”

Note: This algorithm is model-free and off-policy.

Reminder Box

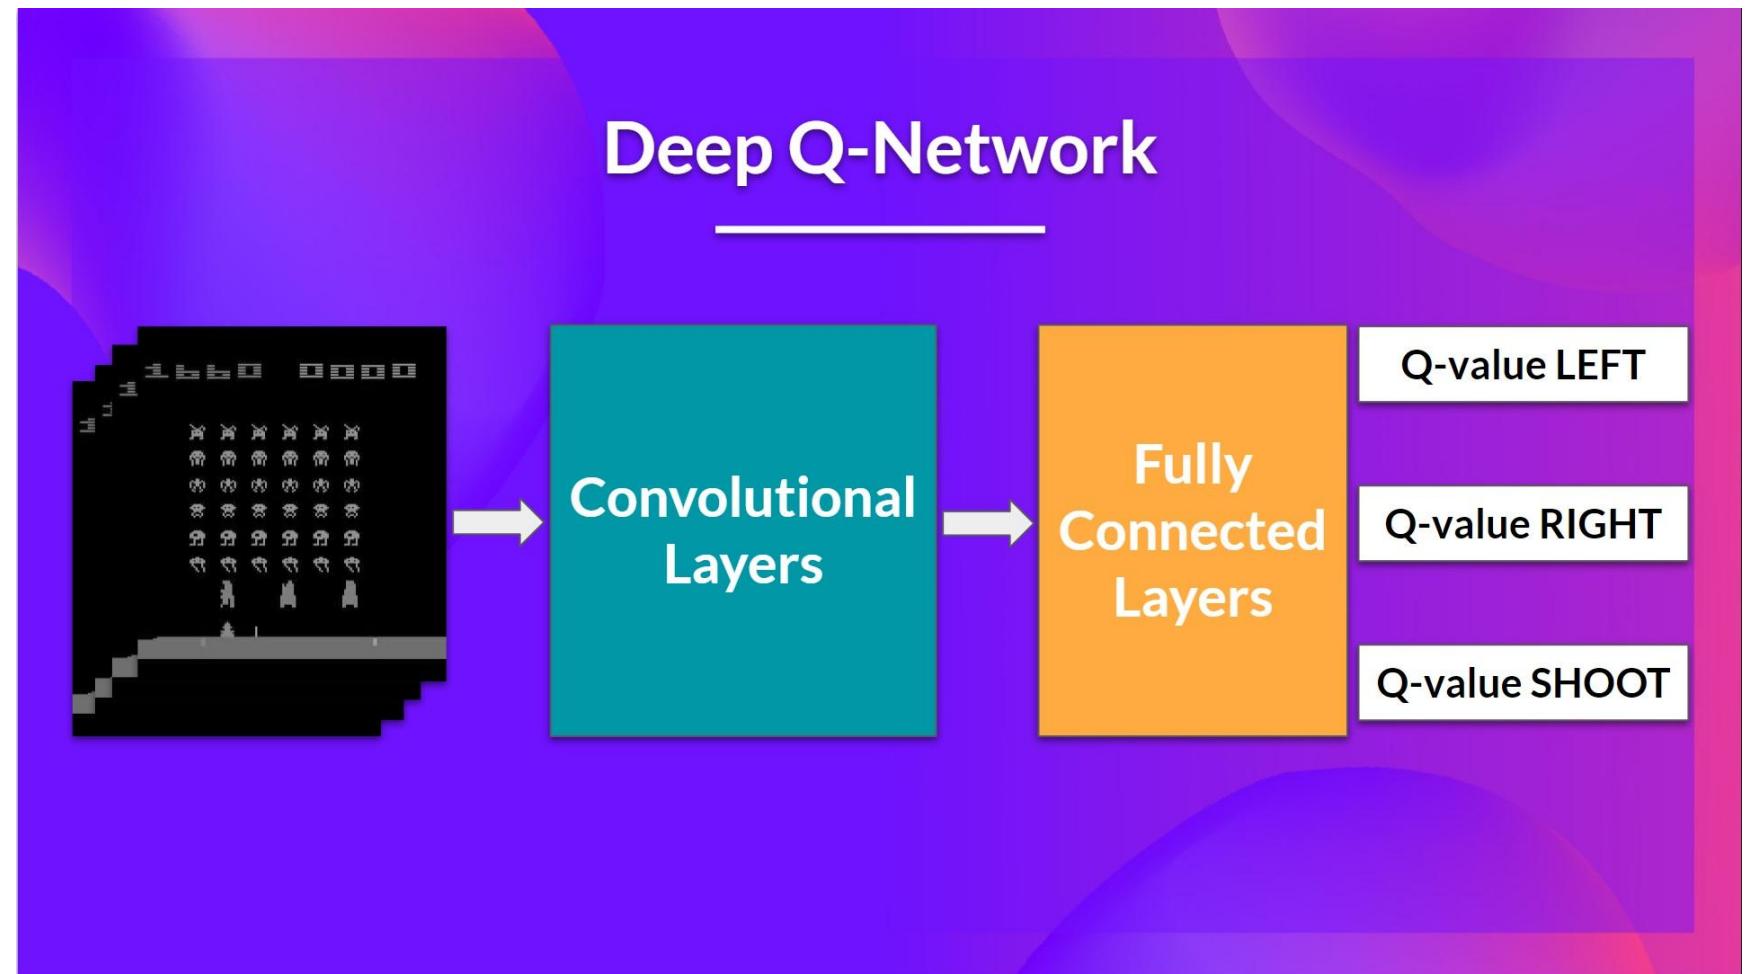
$$L_i(\theta_i) = \mathbb{E}_{s,a \sim \rho(\cdot)} \left[(y_i - Q(s, a; \theta_i))^2 \right]$$

Architecture

Input: Stack of 4 gray-scale frames

Q: Why?

Output: A vector of Q-values for each possible action at that state

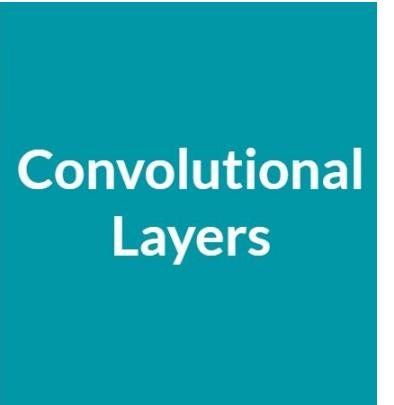


Playing Atari with Deep Reinforcement Learning ([Paper](#))

Architecture: What Is a Convolutional Layer?

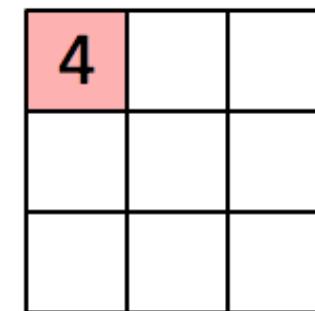
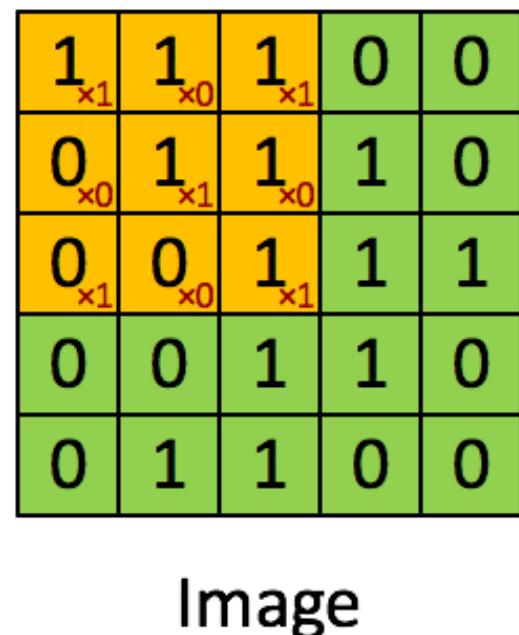
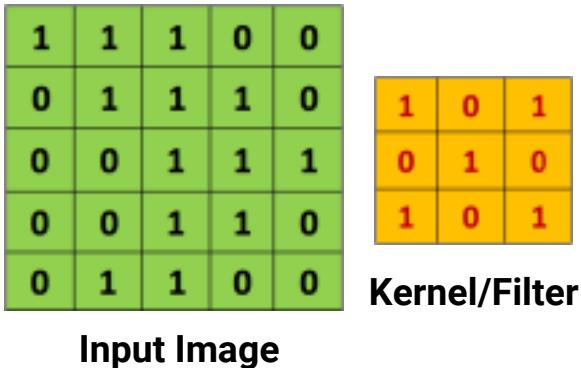


Input

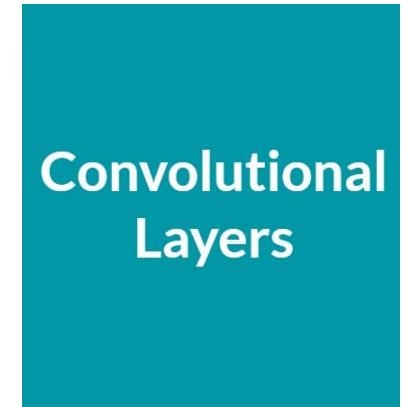


Playing Atari with Deep Reinforcement Learning ([Paper](#))

Architecture: What Is a Convolutional Layer?

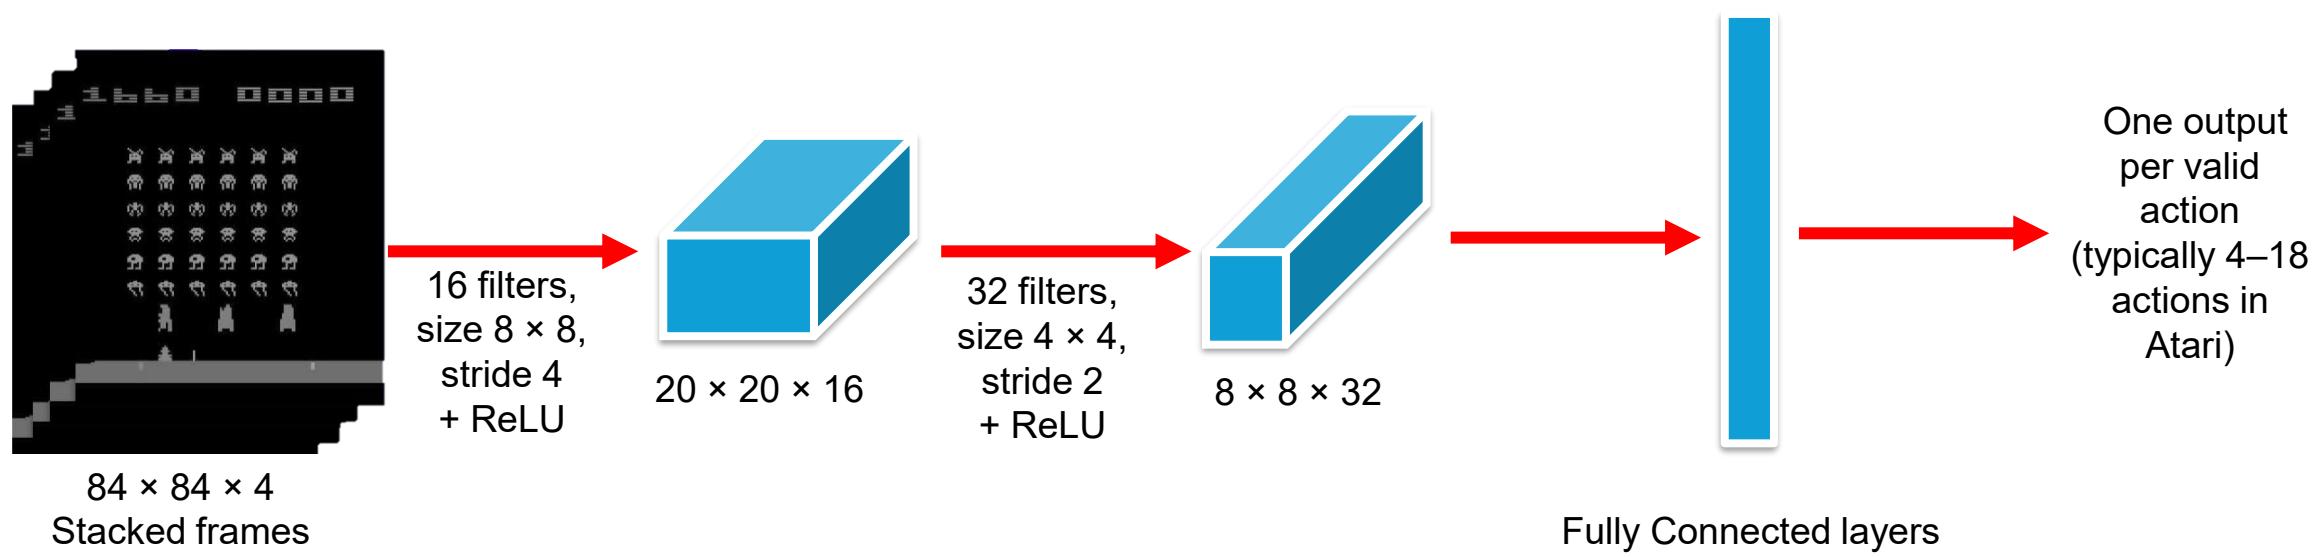


Convolved
Feature



Playing Atari with Deep Reinforcement Learning ([Paper](#))

Architecture

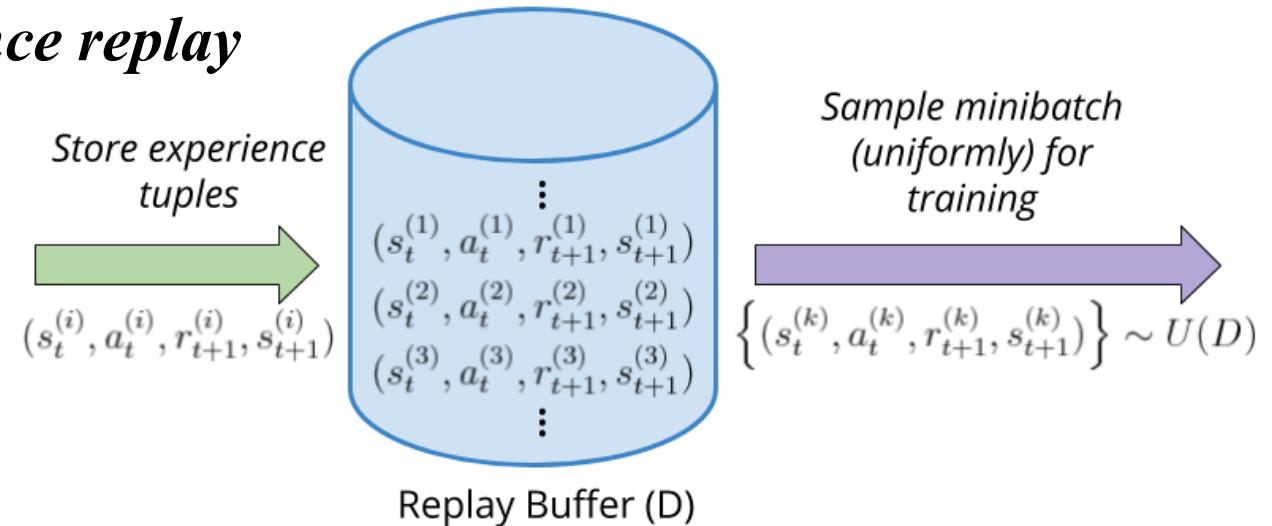


This architecture is known as a **Deep Q-Network (DQN)**

Playing Atari with Deep Reinforcement Learning ([Paper](#))

Algorithm: Replay Buffer

*“We utilize a technique known as **experience replay***

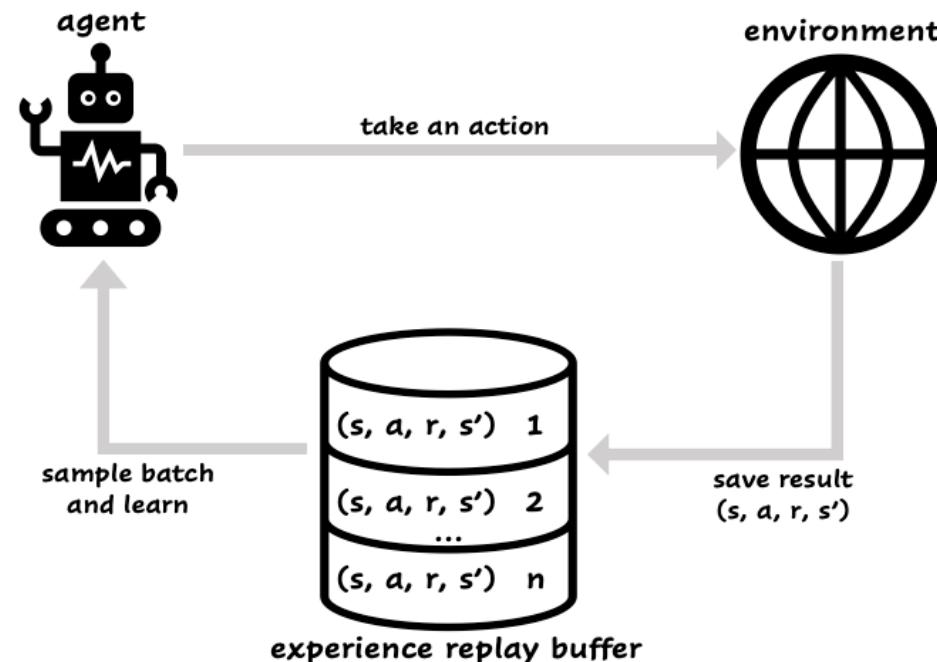


where we store the agent’s experiences at each time-step, $e_t = (s_t, a_t, r_t, s_{t+1})$ in a data-set $D = e_1, \dots, e_N$, pooled over many episodes into a replay memory. During the inner loop of the algorithm, we apply Q-learning updates, or minibatch updates, to samples of experience, $e \sim D$, drawn at random from the pool of stored samples.”

Playing Atari with Deep Reinforcement Learning ([Paper](#))

Algorithm: Replay Buffer

“In practice, our algorithm only stores the last N experience tuples in the replay memory, and samples uniformly at random from D when performing updates.”



Q: Why Replay Buffer?

Playing Atari with Deep Reinforcement Learning ([Paper](#))

Algorithm

Reminder Box

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

New
Q-value
estimation

Former
Q-value
estimation

Learning
Rate

Immediate
Reward

Discounted Estimate
optimal Q-value
of next state

Former
Q-value
estimation

TD Target

TD Error

Playing Atari with Deep Reinforcement Learning ([Paper](#))

Algorithm

Intuition

Q-Target

$$y_j = r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-)$$

$$R_{t+1} + \gamma \max_a Q(S_{t+1}, a)$$

Immediate Reward Discounted Estimate optimal Q-value of next state

TD Target

Q-Loss

$$y_j - Q(\phi_j, a_j; \theta)$$

$$[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

Immediate Reward Discounted Estimate optimal Q-value of next state

TD Target

Former Q-value estimation

TD Error

Playing Atari with Deep Reinforcement Learning ([Paper](#))

Algorithm

Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory \mathcal{D} to capacity N

Initialize action-value function Q with random weights

for episode = 1, M **do**

 Initialise sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in \mathcal{D}

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from \mathcal{D}

 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

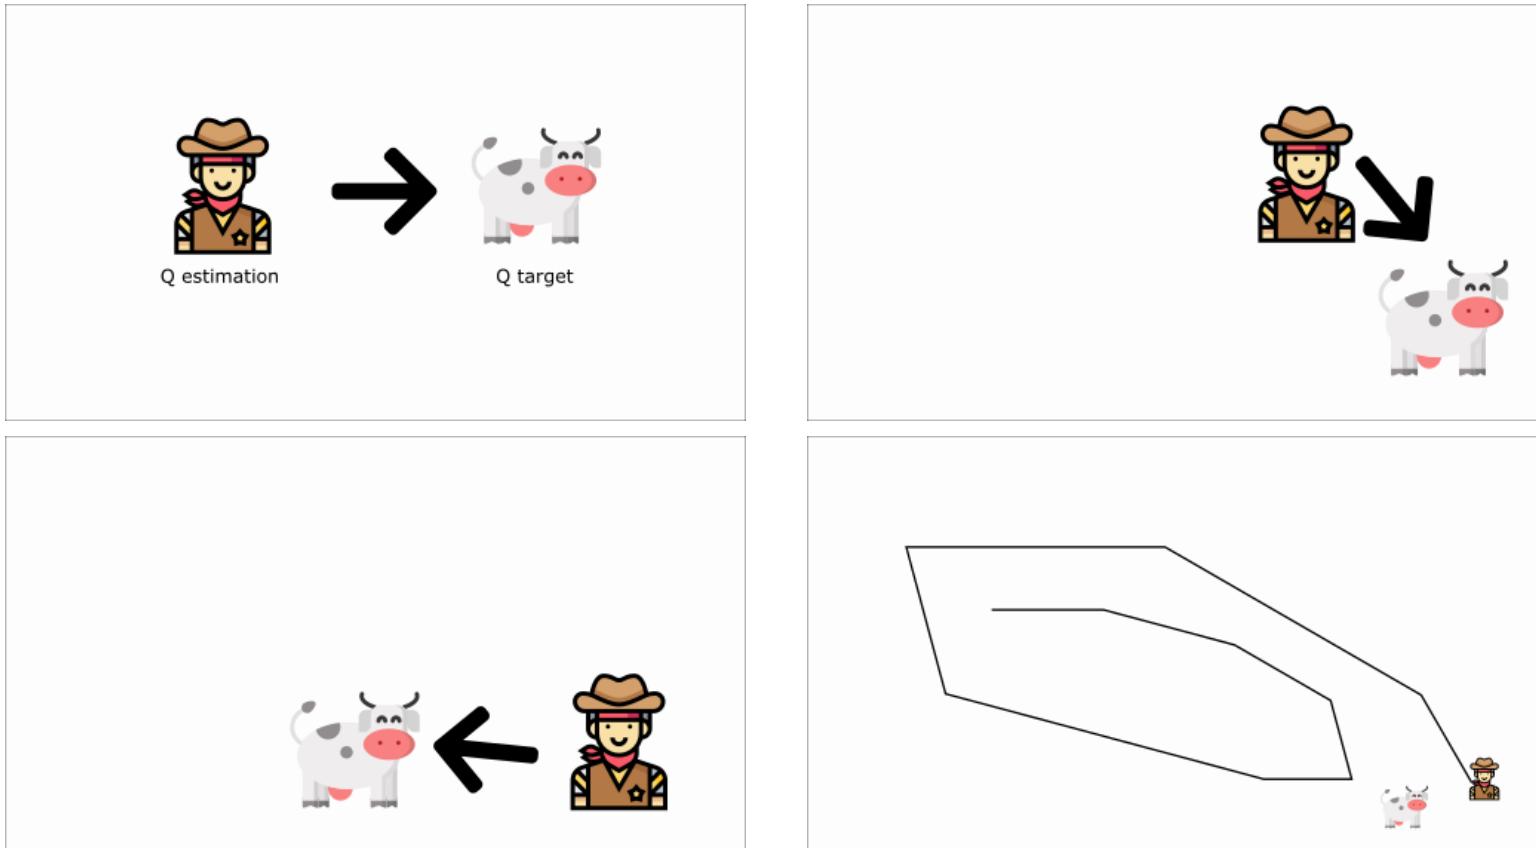
 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3

end for

end for

Playing Atari with Deep Reinforcement Learning ([Paper](#))

Algorithm: A Challenge!



Playing Atari with Deep Reinforcement Learning ([Paper](#))

Algorithm: A Challenge!

Solution:

- Use a **separate network** with fixed parameters for estimating the TD Target
- Copy the parameters from our Deep Q-Network **every C steps** to update the target network.

Algorithm

Algorithm 1: deep Q-learning with experience replay.

```

Initialize replay memory  $D$  to capacity  $N$ 
Initialize action-value function  $Q$  with random weights  $\theta$ 
Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- = \theta$ 
For episode = 1,  $M$  do
    Initialize sequence  $s_1 = \{x_1\}$  and preprocessed sequence  $\phi_1 = \phi(s_1)$ 
    For  $t = 1, T$  do
        With probability  $\varepsilon$  select a random action  $a_t$ 
        otherwise select  $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$ 
        Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$ 
        Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$ 
        Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $D$ 
        Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $D$ 
        Set  $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$ 
        Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  with respect to the
        network parameters  $\theta$ 
        Every  $C$  steps reset  $\hat{Q} = Q$ 
    End For
End For

```

Playing Atari with Deep Reinforcement Learning ([Paper](#), [Paper](#))

Conclusion

“We apply our method to seven Atari 2600 games from the Arcade Learning Environment, with no adjustment of the architecture or learning algorithm. We find that it outperforms all previous approaches on six of the games and surpasses a human expert on three of them, with no adjustment of the architecture or hyperparameters.”

Playing Atari with Deep Reinforcement Learning ([Paper](#), [Paper](#))