

Reinforcement Learning in Control

Dr. Saeed Shamaghdari

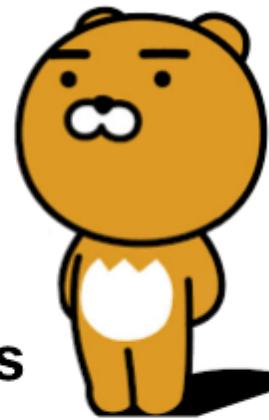
**Electrical Engineering Department
Control Group**

Fall 2025 | 4041

Policy Gradient Methods

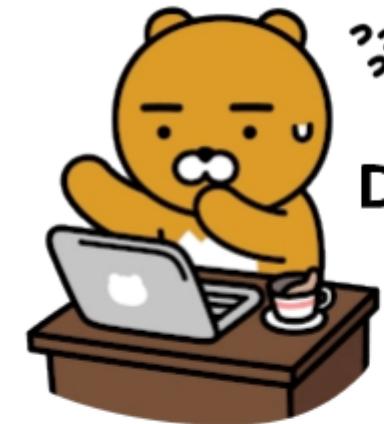
Quick Review

Policy Gradients



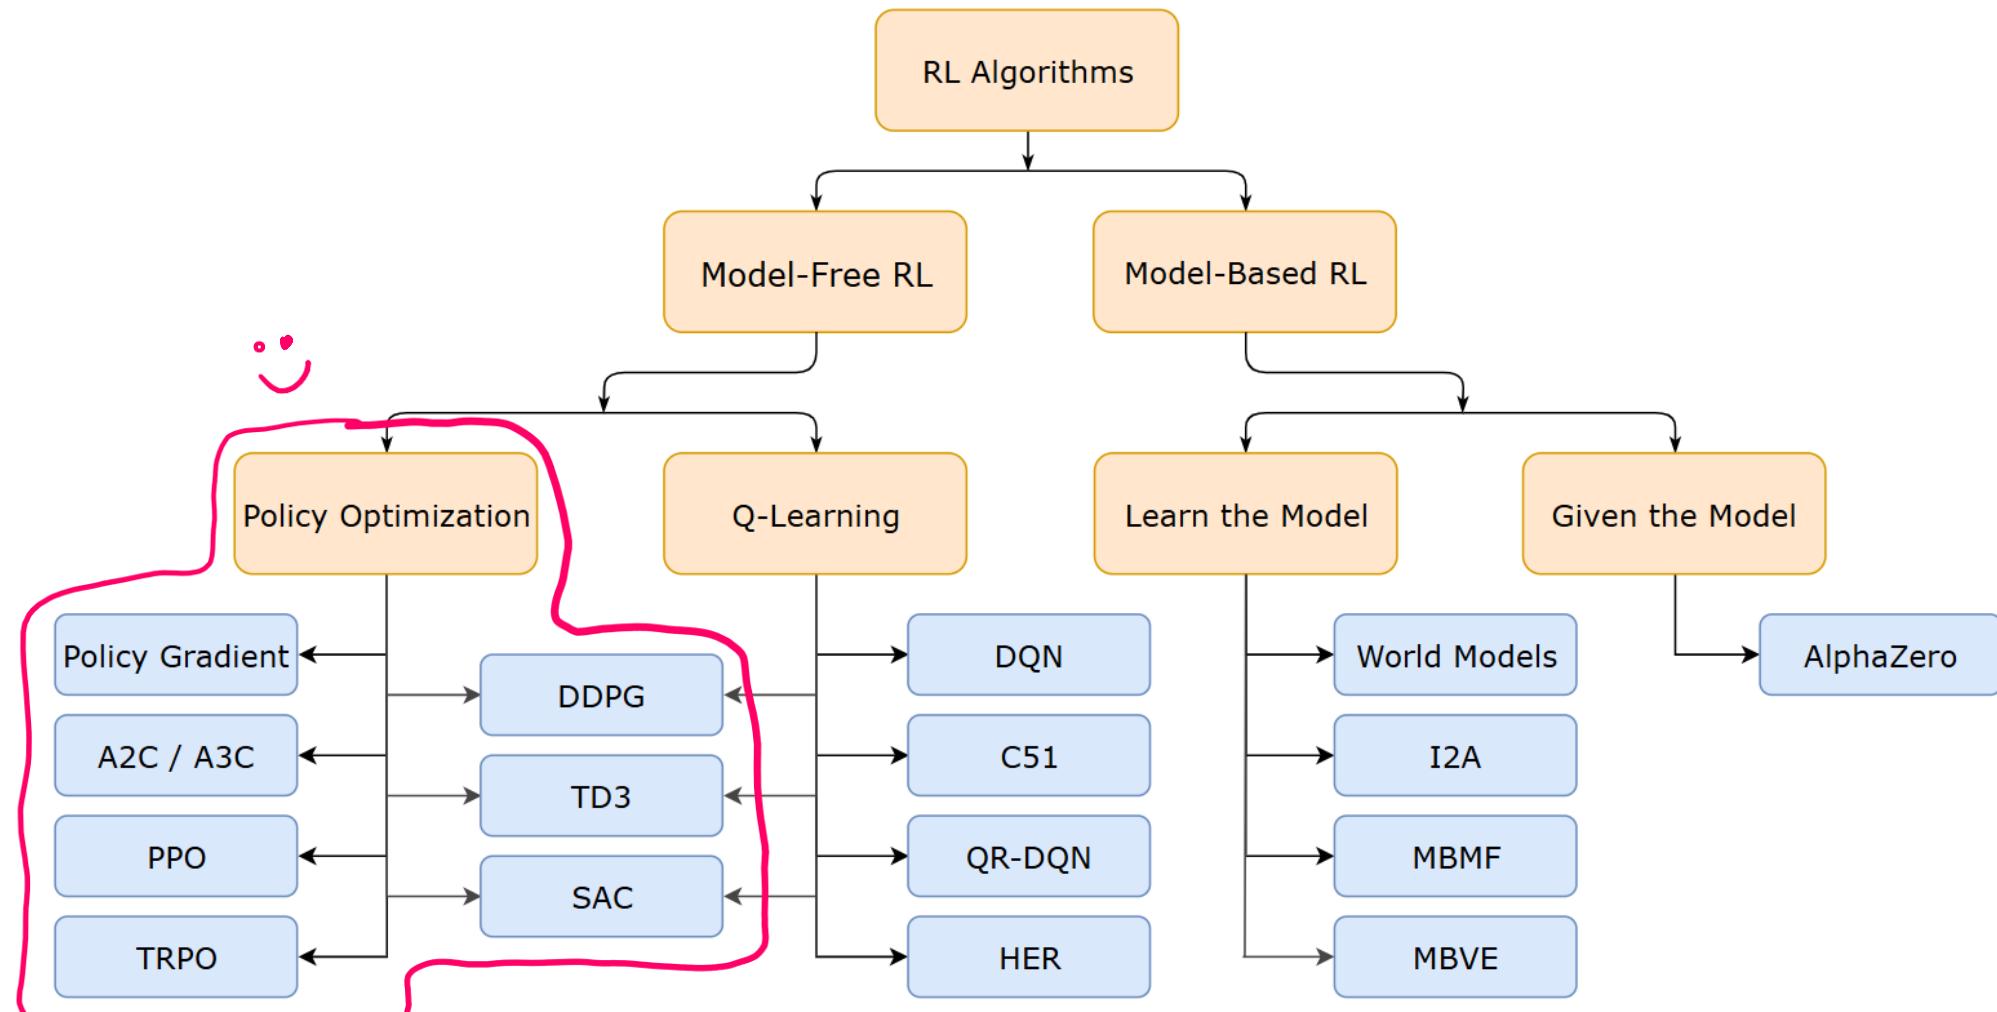
Go Right

Deep Q-Learning

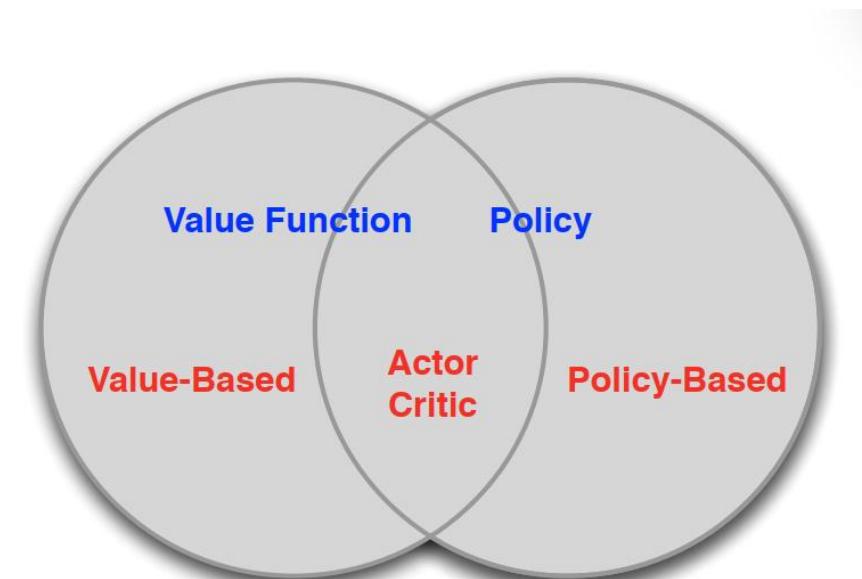
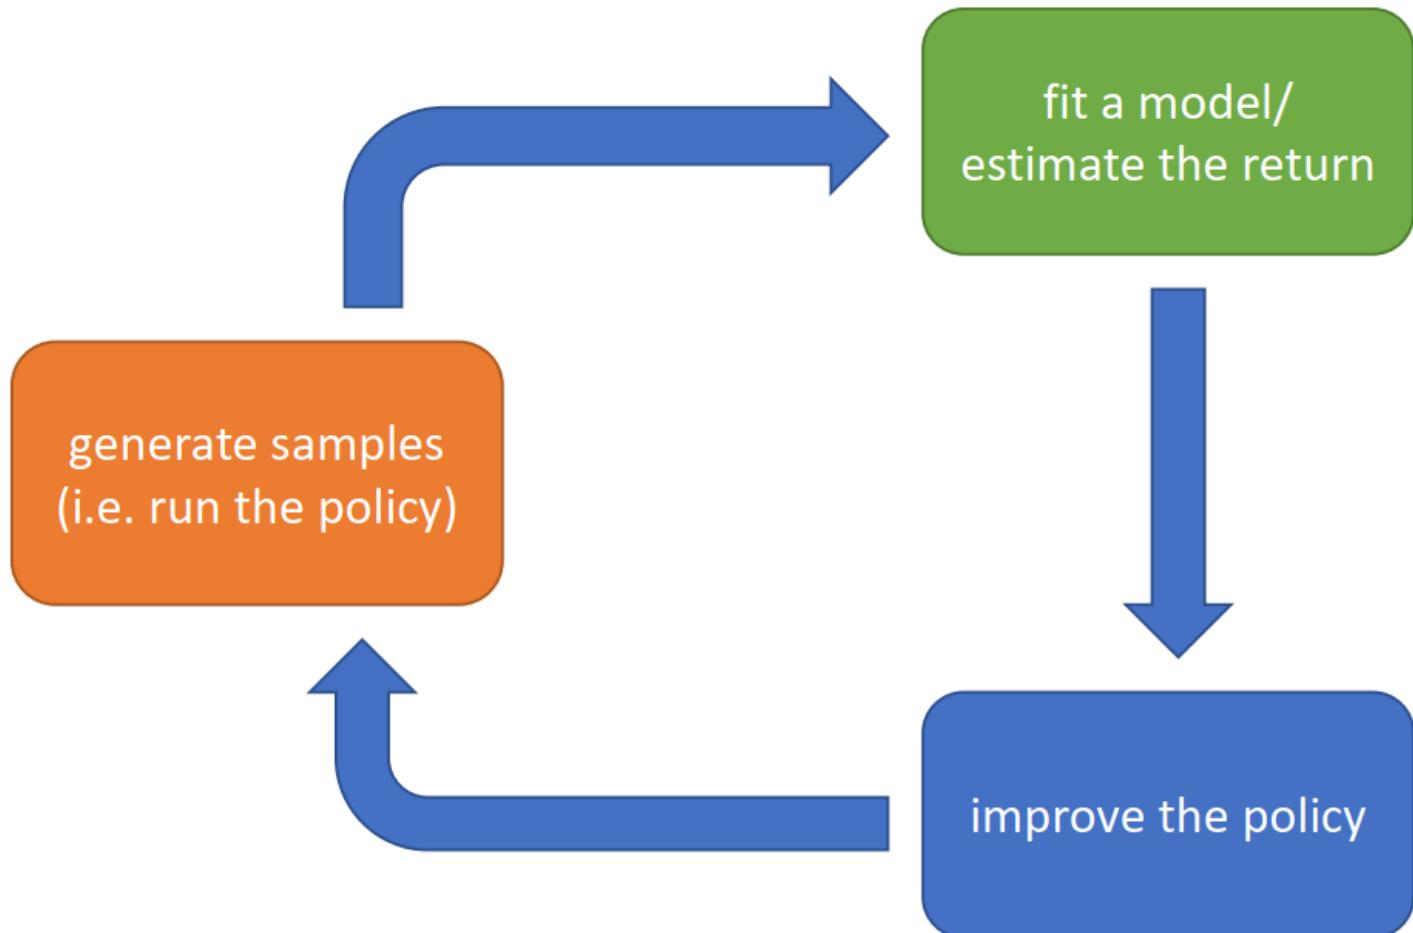


Please wait, I am still calculating Q value, only 41891 actions left...

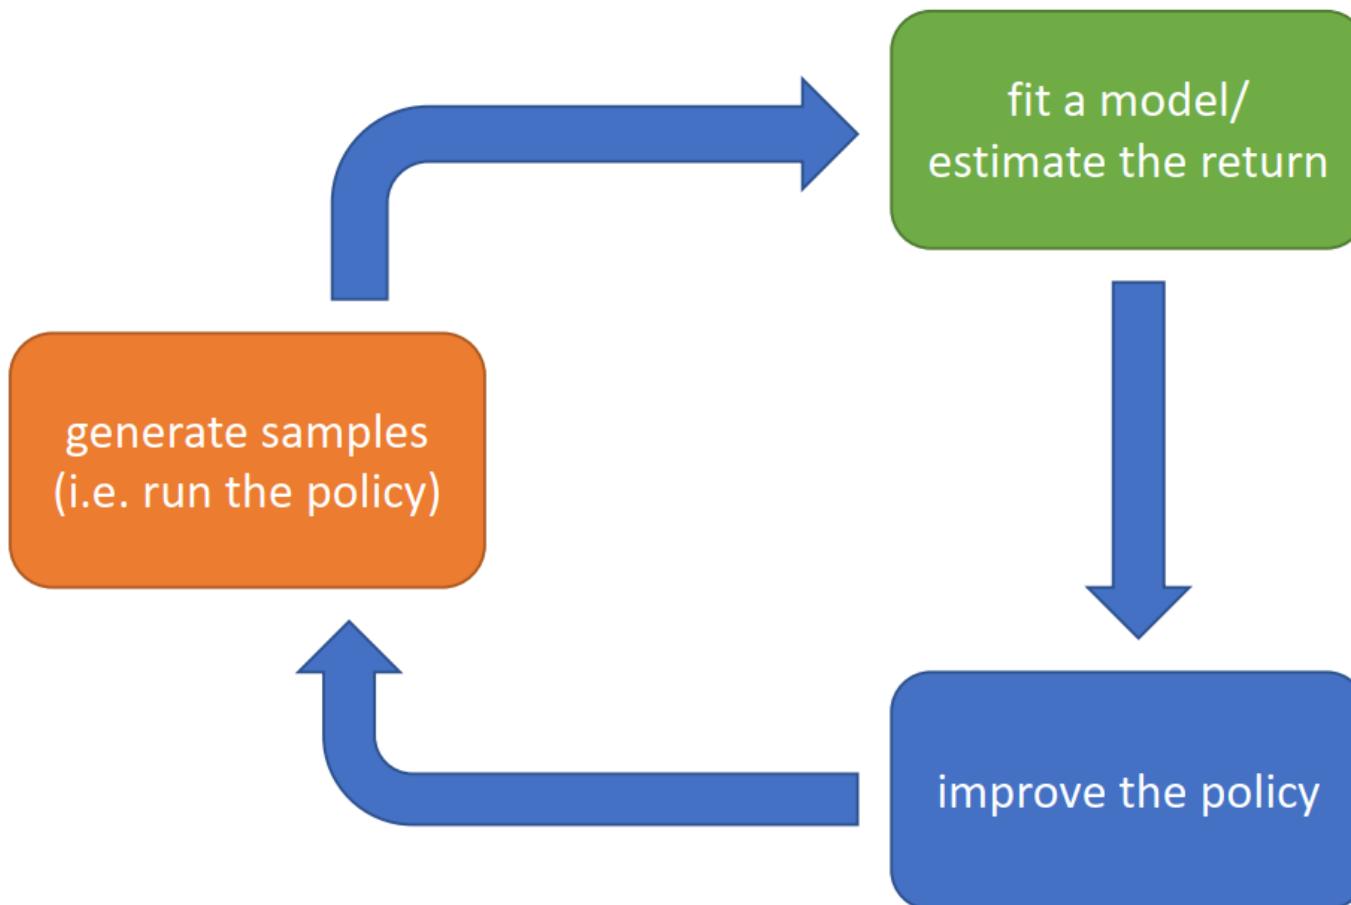
RL Algorithms



Anatomy of RL Algorithms



Anatomy of RL Algorithms: Value Function Based Methods



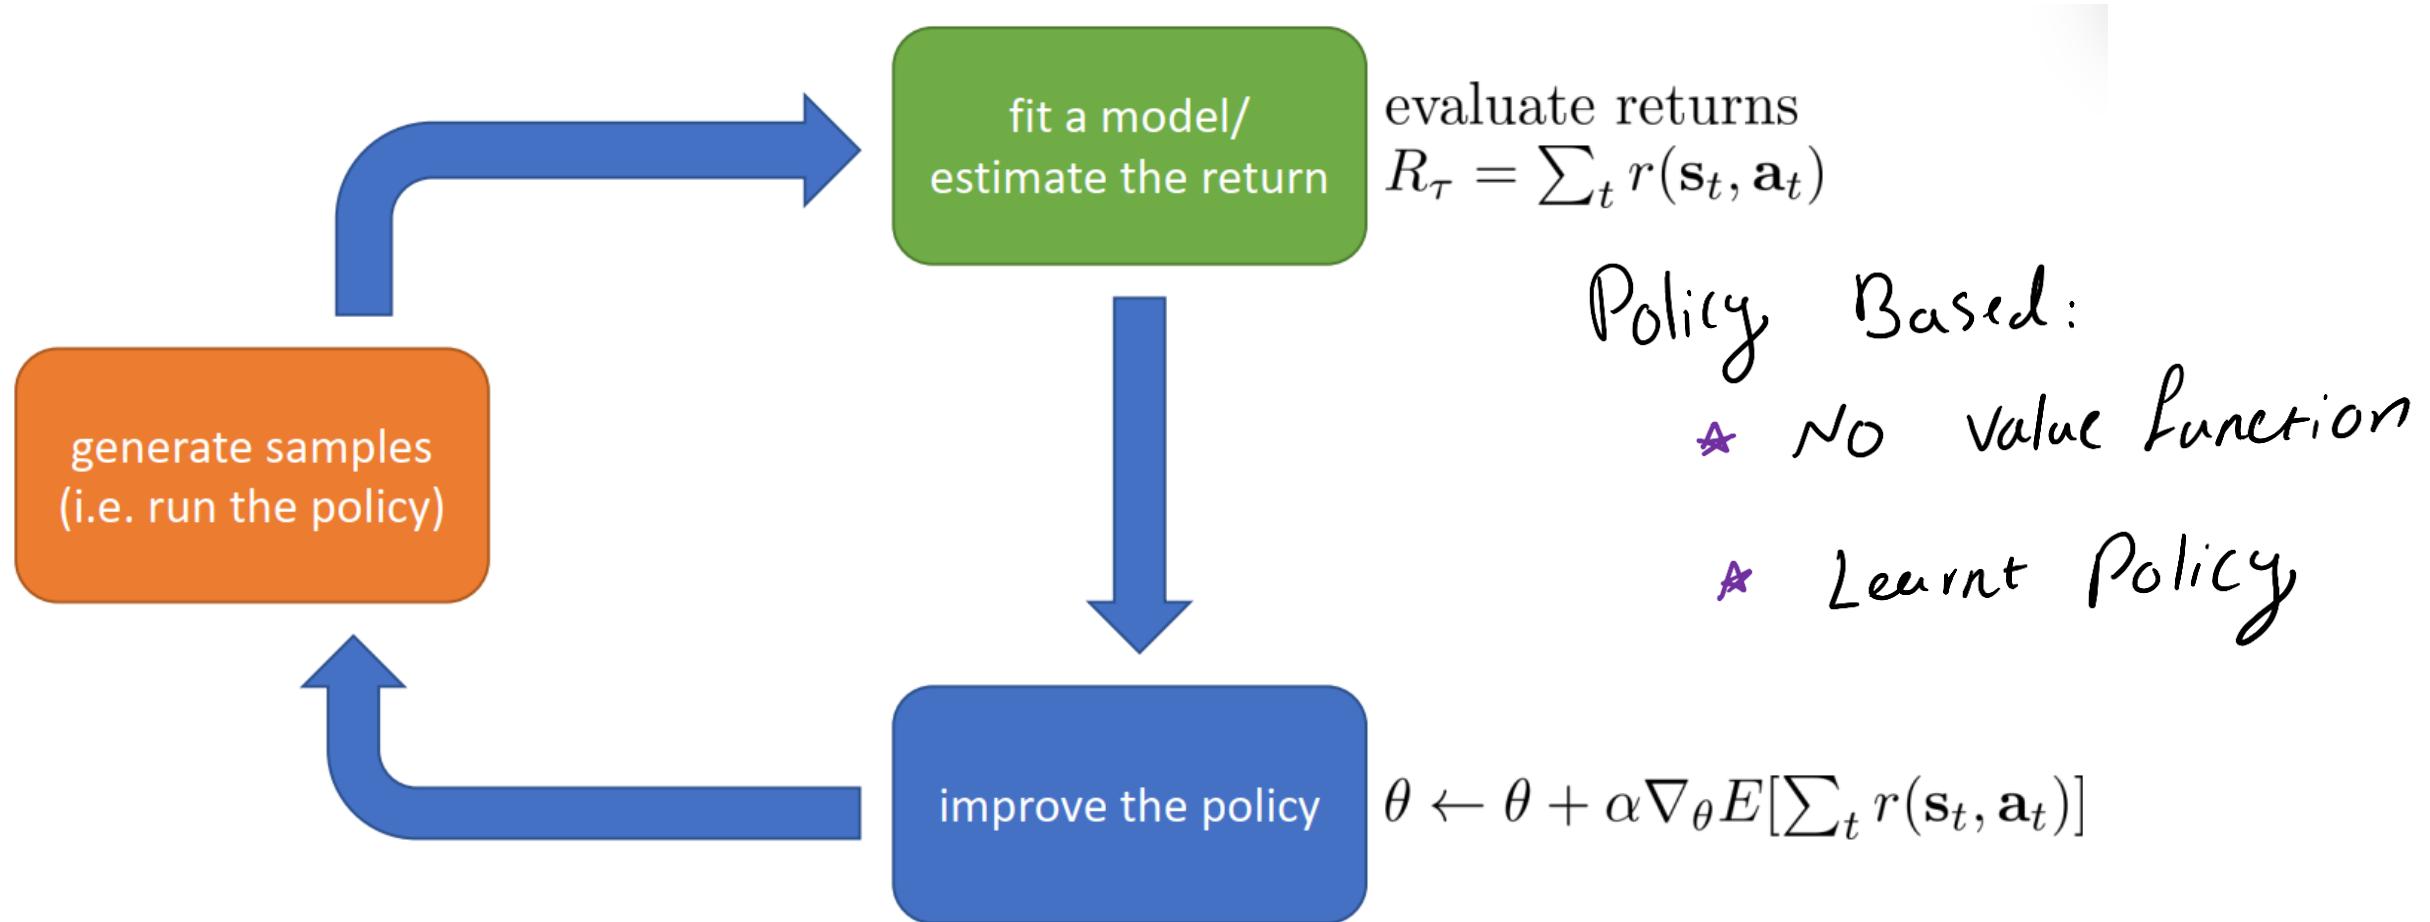
fit $V(\mathbf{s})$ or $Q(\mathbf{s}, \mathbf{a})$

Value Based:

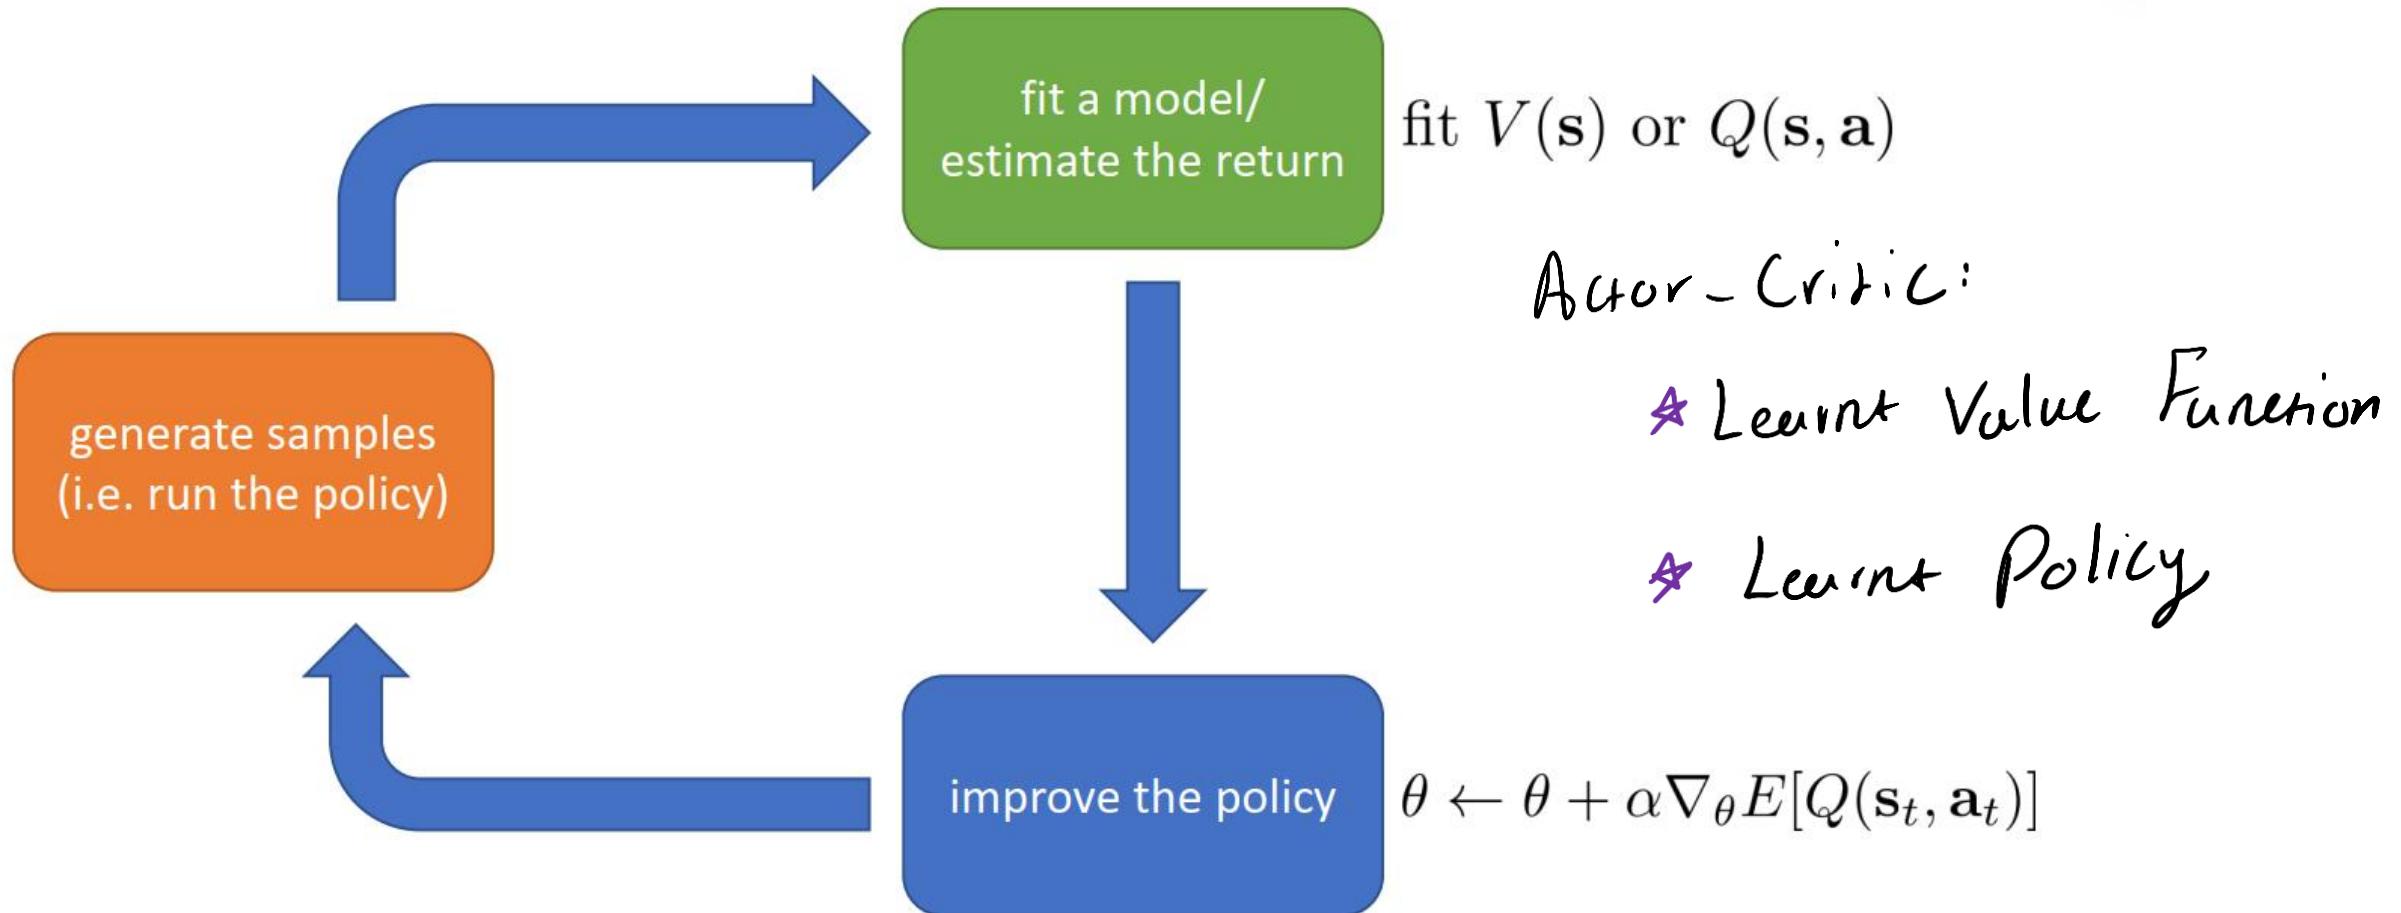
- ★ Learnt Value Function
- ★ Implicit Policy
(e.g. ϵ -greedy)

set $\pi(\mathbf{s}) = \arg \max_{\mathbf{a}} Q(\mathbf{s}, \mathbf{a})$

Anatomy of RL Algorithms: Policy Gradients



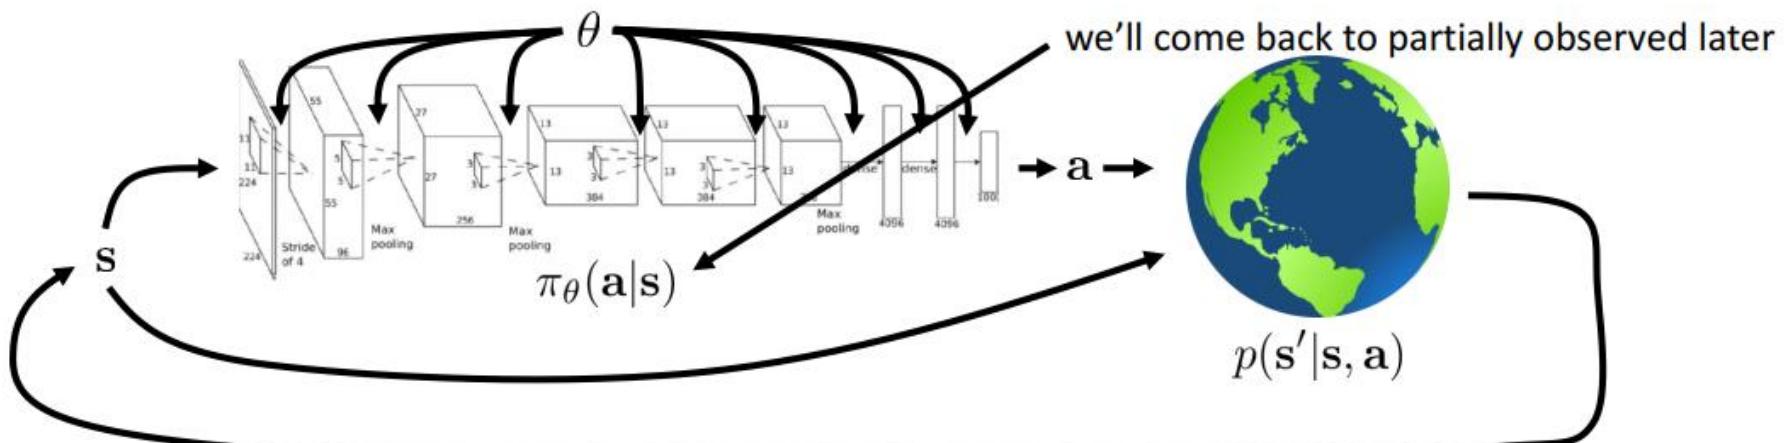
Anatomy of RL Algorithms: Actor-Critic



Policy Gradients

In Action-Value methods $\pi^*(s) = \arg \max_a Q^*(s, a)$

This Chapter: Learn a *parameterized policy*



Policy Gradients

Definition

We write

$$\pi(a|s, \theta) = \Pr\{A_t = a \mid S_t = s, \theta_t = \theta\}$$

for the probability that action a is taken at time t given that the environment is in state s at time t with parameter θ .

$\theta \in \mathbb{R}^{d'}$  policy's parameter vector

Policy Gradients

Policy gradients are methods for learning the policy parameter based on the gradient of some scalar performance measure

$J(\theta)$  Any Policy Objective Function

Policy gradients 

Search for maximum in $J(\theta)$
using Gradient Ascent

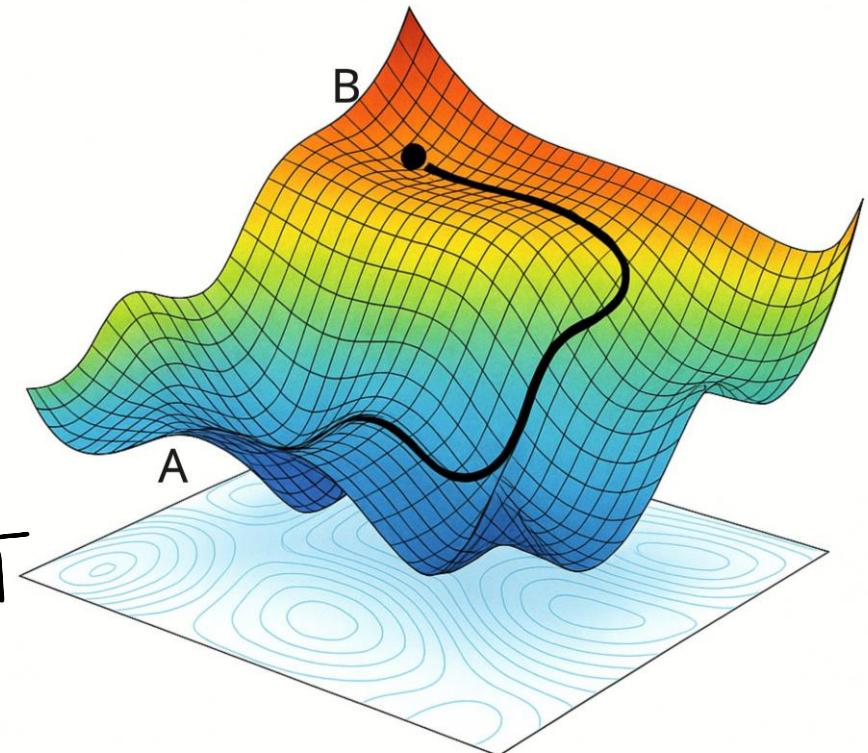
Policy Gradients

Maximize performance using *gradient ascent*

$$\theta_{t+1} = \theta_t + \alpha \widehat{\nabla J}(\theta_t)$$

Where $\widehat{\nabla J}(\theta_t) \in \mathbb{R}^{d'}$ is a stochastic estimate whose expectation approximates the gradient of the performance measure with respect to its argument θ_t .

$$\Delta\theta, \alpha \widehat{\nabla_\theta J}(\theta) \rightarrow \nabla_\theta J(\theta), \left[\frac{\partial J}{\partial \theta_1}, \dots, \frac{\partial J}{\partial \theta_n} \right]^T$$



Policy Gradients: Recap

Objective	Policy Gradient Methods	Actor-Critic Methods
Learn Policy Function	✓	✓
Learn Value Function	✗	✓

Policy Approximation: Soft-max in Action Preferences

Most common parameterization for discrete action spaces (and not too large):

$h(s, a, \theta) \in \mathbb{R}$ (for each action state pairs)

Parametrized numerical preferences

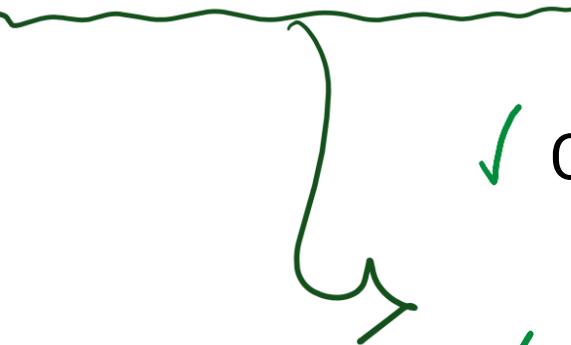
Actions with the $\uparrow h \rightarrow \uparrow$ probabilities of being selected

$$\pi(a|s, \theta) \doteq \frac{e^{h(s, a, \theta)}}{\sum_b e^{h(s, b, \theta)}}$$

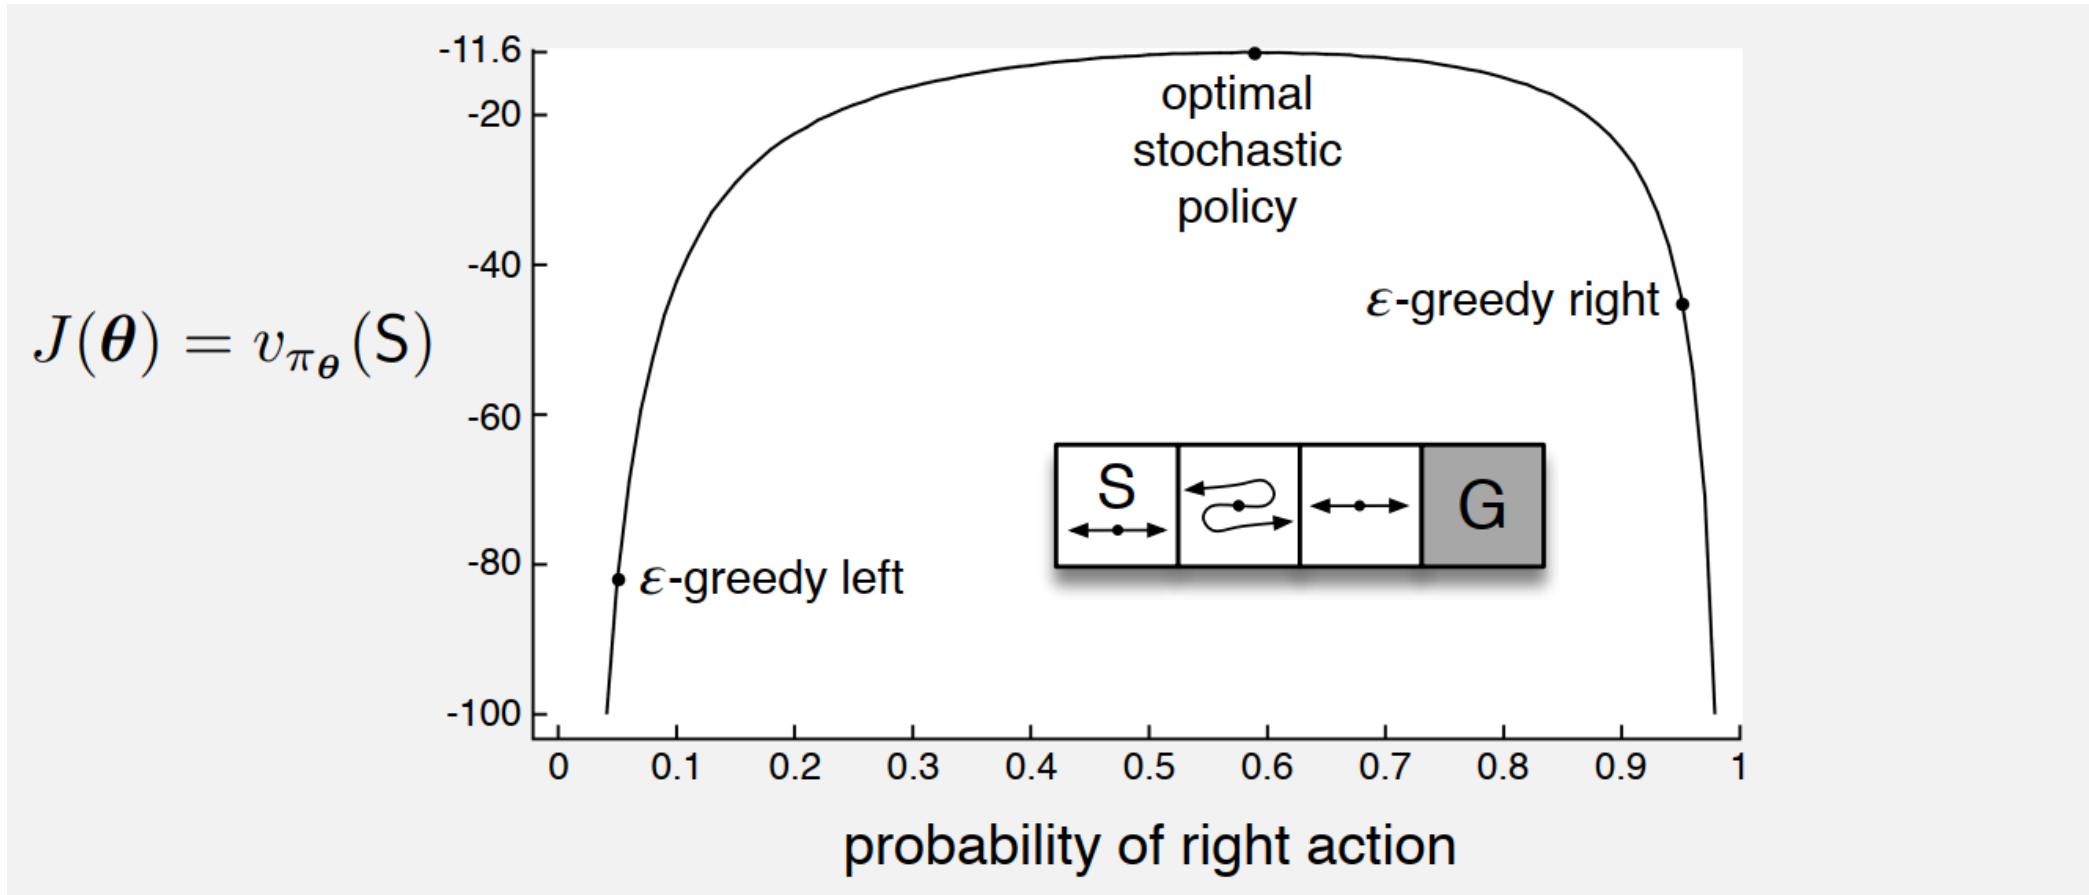
ANN ✓
Linear ✓

h with linear feature vectors $x(s, a) \in \mathbb{R}^{d'}$ $\Rightarrow h(s, a, \theta) = \theta^\top x(s, a)$

Policy Approximation: Advantages

- 
- ✓ Can become nearly deterministic
 - ✓ Can represent arbitrary action probabilities, enabling truly **stochastic** policies.
 - ✓ Simpler function to approximate, learning faster with better final performance.

Example: Short Corridor with Switched Actions



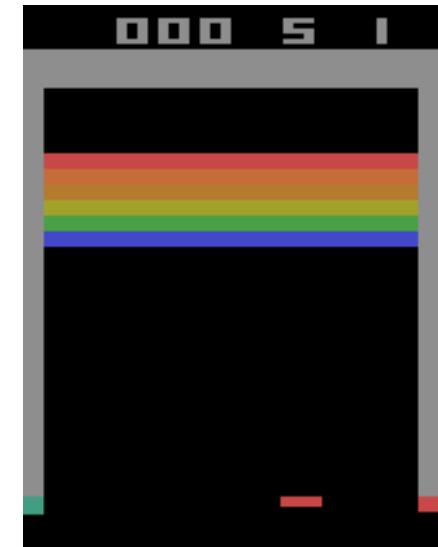
Another Example: Rock-Paper-Scissors

Deterministic Policy vs Stochastic Policy



Advantage of Policy-Based RL: Summary

- **Advantages:**
 - Better convergence properties
 - Effective in **high-dimensional** or **continuous** action spaces
 - Can learn **stochastic** policies
- **Disadvantages:**
 - Typically converge to a local rather than global optimum
 - Evaluating a policy is typically inefficient and **high variance**



The Policy Gradient Theorem

In the episodic case we define performance as ($\gamma=1$)

$$J(\theta) \doteq v_{\pi_\theta}(s_0)$$

true value function

Policy changes affect both actions and the induced state distribution.

Q: How to compute the gradient when policy changes alter an unknown state distribution?

The Policy Gradient Theorem

Proof (Episodic Case)

To keep the notation simple, we leave it implicit in all cases that π is a function of θ , and all gradients are also implicitly with respect to θ . First note that the gradient of the state-value function can be written in terms of the action-value function as

$$\nabla v_\pi(s) = \nabla \left[\sum_a \pi(a|s) q_\pi(s, a) \right], \quad \text{for all } s \in \mathcal{S}$$

The Policy Gradient Theorem

Proof (Episodic Case)

$$\begin{aligned}\nabla v_\pi(s) &= \nabla \left[\sum_a \pi(a|s) q_\pi(s, a) \right], \quad \text{for all } s \in \mathcal{S} \\ &= \sum_a \left[\nabla \pi(a|s) q_\pi(s, a) + \pi(a|s) \nabla q_\pi(s, a) \right] \\ &= \sum_a \left[\nabla \pi(a|s) q_\pi(s, a) + \pi(a|s) \nabla \sum_{s', r} p(s', r | s, a) (r + v_\pi(s')) \right] \\ &= \sum_a \left[\nabla \pi(a|s) q_\pi(s, a) + \pi(a|s) \sum_{s'} p(s' | s, a) \nabla v_\pi(s') \right]\end{aligned}$$

The Policy Gradient Theorem

Proof (Episodic Case)

$$\begin{aligned} &= \sum_a \left[\nabla \pi(a|s) q_\pi(s, a) + \pi(a|s) \sum_{s'} p(s'|s, a) \nabla v_\pi(s') \right] \\ &= \sum_a \left[\nabla \pi(a|s) q_\pi(s, a) + \pi(a|s) \sum_{s'} p(s'|s, a) \right. \\ &\quad \left. \sum_{a'} \left[\nabla \pi(a'|s') q_\pi(s', a') + \pi(a'|s') \sum_{s''} p(s''|s', a') \nabla v_\pi(s'') \right] \right] \\ &= \sum_{x \in \mathcal{S}} \sum_{k=0}^{\infty} \Pr(s \rightarrow x, k, \pi) \sum_a \nabla \pi(a|x) q_\pi(x, a), \end{aligned}$$

The Policy Gradient Theorem

Proof (Episodic Case)

$$= \sum_{x \in \mathcal{S}} \sum_{k=0}^{\infty} \Pr(s \rightarrow x, k, \pi) \sum_a \nabla \pi(a|x) q_{\pi}(x, a),$$

after repeated unrolling, where $\Pr(s \rightarrow x, k, \pi)$ is the probability of transitioning from state s to state x in k steps under policy π . It is then immediate that

$$\nabla J(\theta) = \nabla v_{\pi}(s_0)$$

The Policy Gradient Theorem

Proof (Episodic Case)

$$\begin{aligned}\nabla J(\boldsymbol{\theta}) &= \nabla v_{\pi}(s_0) \\ &= \sum_s \left(\sum_{k=0}^{\infty} \Pr(s_0 \rightarrow s, k, \pi) \right) \sum_a \nabla \pi(a|s) q_{\pi}(s, a) \\ &= \sum_s \eta(s) \sum_a \nabla \pi(a|s) q_{\pi}(s, a) \\ &= \sum_{s'} \eta(s') \sum_s \frac{\eta(s)}{\sum_{s'} \eta(s')} \sum_a \nabla \pi(a|s) q_{\pi}(s, a)\end{aligned}$$

The Policy Gradient Theorem

Proof (Episodic Case)

$$\begin{aligned} &= \sum_{s'} \eta(s') \sum_s \frac{\eta(s)}{\sum_{s'} \eta(s')} \sum_a \nabla \pi(a|s) q_\pi(s, a) \\ &= \sum_{s'} \eta(s') \sum_s \mu(s) \sum_a \nabla \pi(a|s) q_\pi(s, a) \\ &\propto \sum_s \mu(s) \sum_a \nabla \pi(a|s) q_\pi(s, a) \end{aligned}$$

that does *not* involve the derivative of the state distribution.

The Policy Gradient Theorem

The policy gradient theorem for the episodic case establishes that

$$\begin{aligned}\nabla J(\theta) &\propto \sum_s \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a|s, \theta) \\ &= \mathbb{E}_\pi \left[\sum_a q_\pi(S_t, a) \nabla \pi(a|S_t, \theta) \right].\end{aligned}$$

Stochastic gradient-ascent algorithm

$$\theta_{t+1} \doteq \theta_t + \alpha \sum_a \hat{q}(S_t, a, \mathbf{w}) \nabla \pi(a|S_t, \theta)$$

Our current interest is the classical REINFORCE algorithm (Williams, 1992) whose update at time t involves just A_t , the one action actually taken at time t .

REINFORCE: Monte Carlo Policy Gradient

$$\nabla J(\boldsymbol{\theta}) \propto \mathbb{E}_\pi \left[\sum_a q_\pi(S_t, a) \nabla \pi(a|S_t, \boldsymbol{\theta}) \right]$$

The REINFORCE Update:

$$\boldsymbol{\theta}_{t+1} \doteq \boldsymbol{\theta}_t + \alpha G_t \frac{\nabla \pi(A_t|S_t, \boldsymbol{\theta}_t)}{\pi(A_t|S_t, \boldsymbol{\theta}_t)}$$

$$\propto \mathbb{E}_\pi \left[\sum_a \pi(a|S_t, \boldsymbol{\theta}) q_\pi(S_t, a) \frac{\nabla \pi(a|S_t, \boldsymbol{\theta})}{\pi(a|S_t, \boldsymbol{\theta})} \right]$$

$$= \mathbb{E}_\pi \left[q_\pi(S_t, A_t) \frac{\nabla \pi(A_t|S_t, \boldsymbol{\theta})}{\pi(A_t|S_t, \boldsymbol{\theta})} \right]$$

(replacing a by the sample $A_t \sim \pi$)

$$= \mathbb{E}_\pi \left[G_t \frac{\nabla \pi(A_t|S_t, \boldsymbol{\theta})}{\pi(A_t|S_t, \boldsymbol{\theta})} \right],$$

(because $\mathbb{E}_\pi[G_t|S_t, A_t] = q_\pi(S_t, A_t)$)

REINFORCE: Monte Carlo Policy Gradient

The REINFORCE Update:

$$\theta_{t+1} \doteq \theta_t + \alpha G_t \frac{\nabla \pi(A_t | S_t, \theta_t)}{\pi(A_t | S_t, \theta_t)}$$

$$\nabla \ln x = \frac{\nabla x}{x}$$

$$\nabla \ln \pi(A_t | S_t, \theta_t) = \frac{\nabla \pi(A_t | S_t, \theta_t)}{\pi(A_t | S_t, \theta_t)}$$

$$\theta_{t+1} \doteq \theta_t + \alpha G_t \nabla \ln \pi(A_t | S_t, \theta_t)$$

REINFORCE: Algorithm (Discounted)

REINFORCE: Monte-Carlo Policy-Gradient Control (episodic) for π_*

Input: a differentiable policy parameterization $\pi(a|s, \theta)$

Algorithm parameter: step size $\alpha > 0$

Initialize policy parameter $\theta \in \mathbb{R}^{d'}$ (e.g., to $\mathbf{0}$)

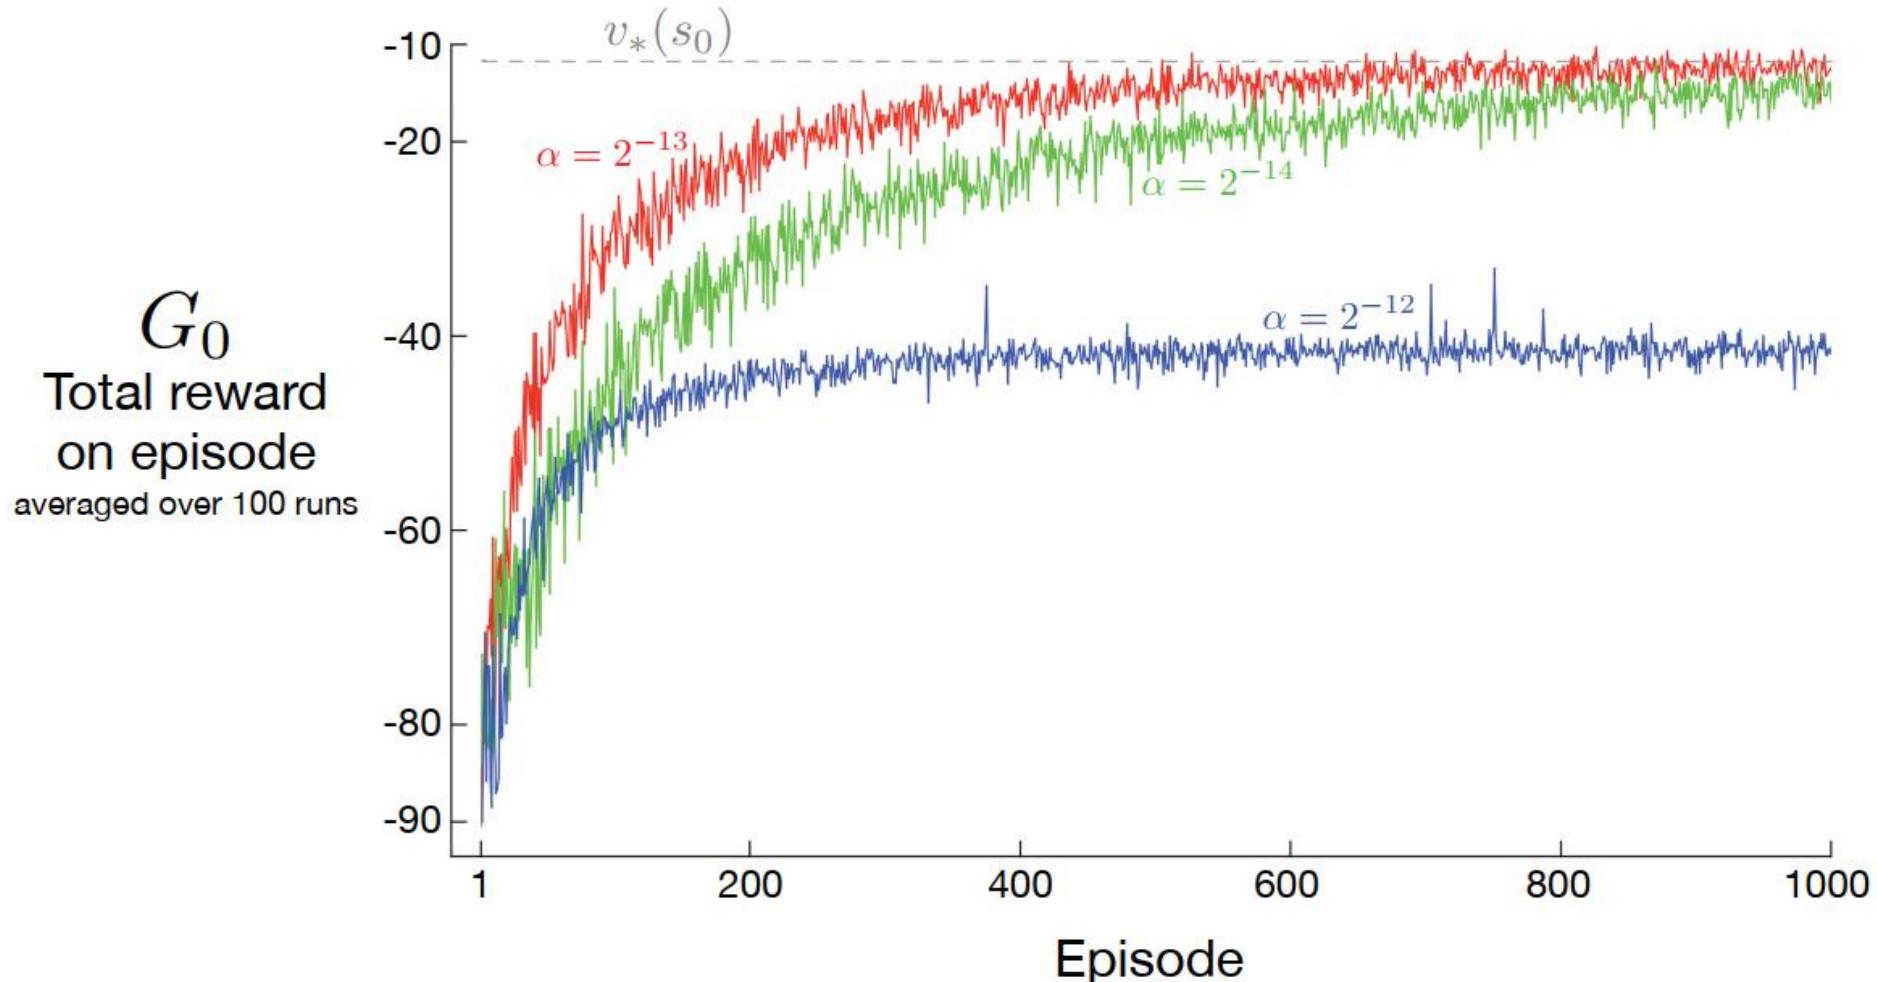
Loop forever (for each episode):

Generate an episode $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$, following $\pi(\cdot| \cdot, \theta)$

Loop for each step of the episode $t = 0, 1, \dots, T - 1$:

$$\begin{aligned} G &\leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} R_k & (G_t) \\ \theta &\leftarrow \theta + \alpha \gamma^t G \nabla \ln \pi(A_t | S_t, \theta) \end{aligned}$$

REINFORCE: Short Corridor Example



REINFORCE: Problem!

Good convergence properties, However ...

Q: What problem does the REINFORCE method have as a Monte Carlo approach?

REINFORCE with Baseline

The policy gradient theorem can be generalized to include a comparison of the action value to an arbitrary *baseline* $b(s)$:

$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a \left(q_\pi(s, a) - b(s) \right) \nabla \pi(a|s, \theta)$$

The baseline can be any function, even a random variable, as long as it does not vary with a ; the equation remains valid because the subtracted quantity is zero:

$$\sum_a b(s) \nabla \pi(a|s, \theta) = b(s) \nabla \sum_a \pi(a|s, \theta) = b(s) \nabla 1 = 0$$

REINFORCE with Baseline: Update

$$\boldsymbol{\theta}_{t+1} \doteq \boldsymbol{\theta}_t + \alpha \left(G_t - b(S_t) \right) \frac{\nabla \pi(A_t | S_t, \boldsymbol{\theta}_t)}{\pi(A_t | S_t, \boldsymbol{\theta}_t)}$$

One natural choice for the baseline is an estimate of the state value

$$\hat{v}(S_t, \mathbf{w})$$

$\mathbf{w} \in \mathbb{R}^d$ is a weight vector

REINFORCE with Baseline: Algorithm

REINFORCE with Baseline (episodic), for estimating $\pi_\theta \approx \pi_*$

Input: a differentiable policy parameterization $\pi(a|s, \theta)$

Input: a differentiable state-value function parameterization $\hat{v}(s, w)$

Algorithm parameters: step sizes $\alpha^\theta > 0$, $\alpha^w > 0$

Initialize policy parameter $\theta \in \mathbb{R}^{d'}$ and state-value weights $w \in \mathbb{R}^d$ (e.g., to $\mathbf{0}$)

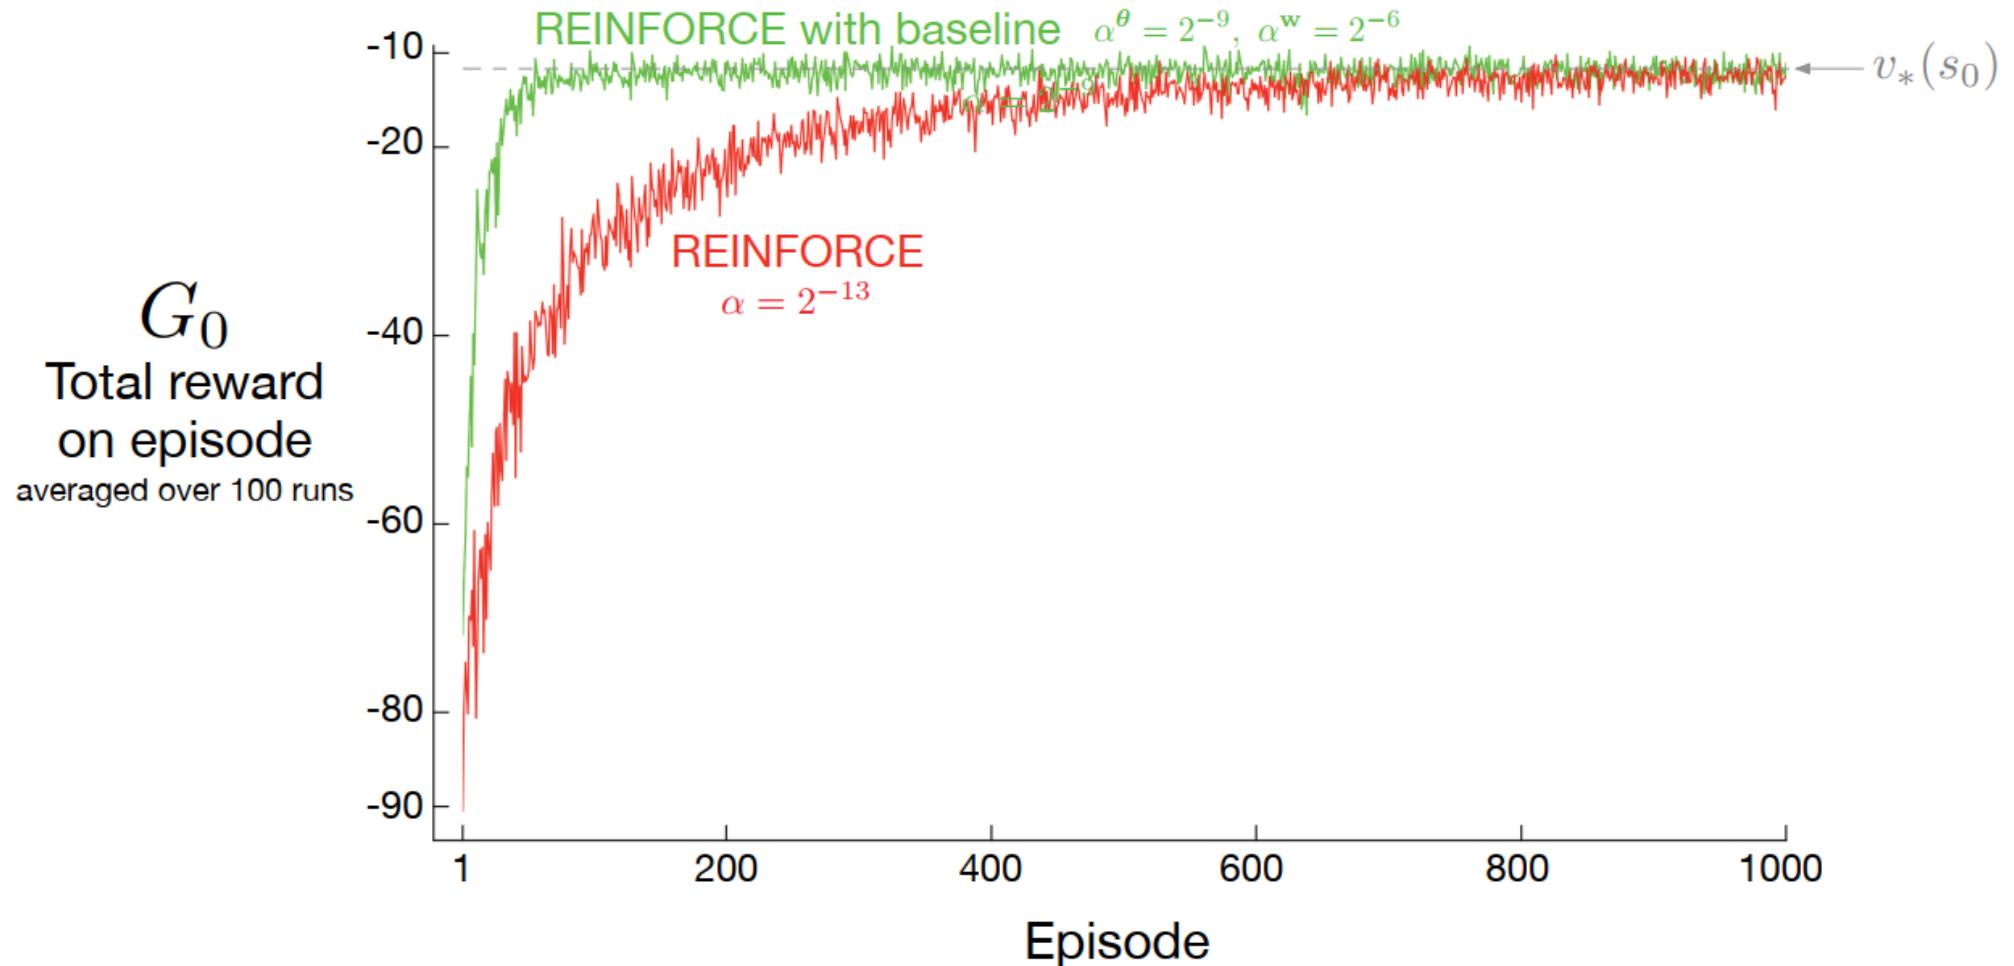
Loop forever (for each episode):

Generate an episode $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$, following $\pi(\cdot|\cdot, \theta)$

Loop for each step of the episode $t = 0, 1, \dots, T - 1$:

$$\begin{aligned} G &\leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} R_k & (G_t) \\ \delta &\leftarrow G - \hat{v}(S_t, w) \\ w &\leftarrow w + \alpha^w \delta \nabla \hat{v}(S_t, w) \\ \theta &\leftarrow \theta + \alpha^\theta \gamma^t \delta \nabla \ln \pi(A_t | S_t, \theta) \end{aligned}$$

REINFORCE with Baseline: Short Corridor Example



Actor-Critic Methods

REINFORCE with baseline:

$$\theta_{t+1} \doteq \theta_t + \alpha \left(G_t - \hat{v}(s, \omega) \right) \frac{\nabla \pi(A_t | S_t, \theta_t)}{\pi(A_t | S_t, \theta_t)}$$

$$\Rightarrow \theta_{t+1} \doteq \theta_t + \alpha (G_t - \hat{v}(s_t, \omega)) \frac{\nabla \pi(A_t | S_t, \theta_t)}{\pi(A_t | S_t, \theta_t)}$$

Replace full return G_t with one-step return $G_{t:t+1}$

Actor-Critic Methods

REINFORCE with baseline:

$$\Rightarrow \theta_{t+1} = \theta_t + \alpha (G_t - \hat{V}(s_t, \omega)) \frac{\nabla \ln \pi(A_t | s_t, \theta_t)}{\pi(A_t | s_t, \theta_t)}$$

Replace full return G_t with one-step return $G_{t:t+1}$

$$G_{t:t+1} = R_{t+1} + \gamma \hat{V}(s_{t+1}, \omega)$$

$$\theta_{t+1} = \theta_t + \alpha \underbrace{[R_{t+1} + \gamma \hat{V}(s_{t+1}, \omega) - \hat{V}(s_t, \omega)]}_{\text{TD Error } \delta_t} \nabla \ln \pi(A_t | s_t, \theta_t)$$

Actor-Critic Methods

Actor Critic Update

$$\theta_{t+1} = \theta_t + \alpha \underbrace{[R_{t+1} + \gamma \hat{v}(s_{t+1}, \omega) - \hat{v}(s_t, \omega)]}_{\text{TD Error } \delta_t} \nabla \ln \pi(A_t | s_t, \theta_t)$$

$$\Rightarrow \theta_{t+1} = \theta_t + \alpha \delta_t \nabla \ln \pi(A_t | s_t, \theta_t)$$

Actor-Critic Methods: Algorithm

One-step Actor–Critic (episodic), for estimating $\pi_\theta \approx \pi_*$

Input: a differentiable policy parameterization $\pi(a|s, \theta)$

Input: a differentiable state-value function parameterization $\hat{v}(s, w)$

Parameters: step sizes $\alpha^\theta > 0$, $\alpha^w > 0$

Initialize policy parameter $\theta \in \mathbb{R}^{d'}$ and state-value weights $w \in \mathbb{R}^d$ (e.g., to $\mathbf{0}$)

Loop forever (for each episode):

 Initialize S (first state of episode)

$I \leftarrow 1$

 Loop while S is not terminal (for each time step):

$A \sim \pi(\cdot|S, \theta)$

 Take action A , observe S', R

$\delta \leftarrow R + \gamma \hat{v}(S', w) - \hat{v}(S, w)$ (if S' is terminal, then $\hat{v}(S', w) \doteq 0$)

$w \leftarrow w + \alpha^w \delta \nabla \hat{v}(S, w)$

$\theta \leftarrow \theta + \alpha^\theta I \delta \nabla \ln \pi(A|S, \theta)$

$I \leftarrow \gamma I$

$S \leftarrow S'$