

Reinforcement Learning in Control

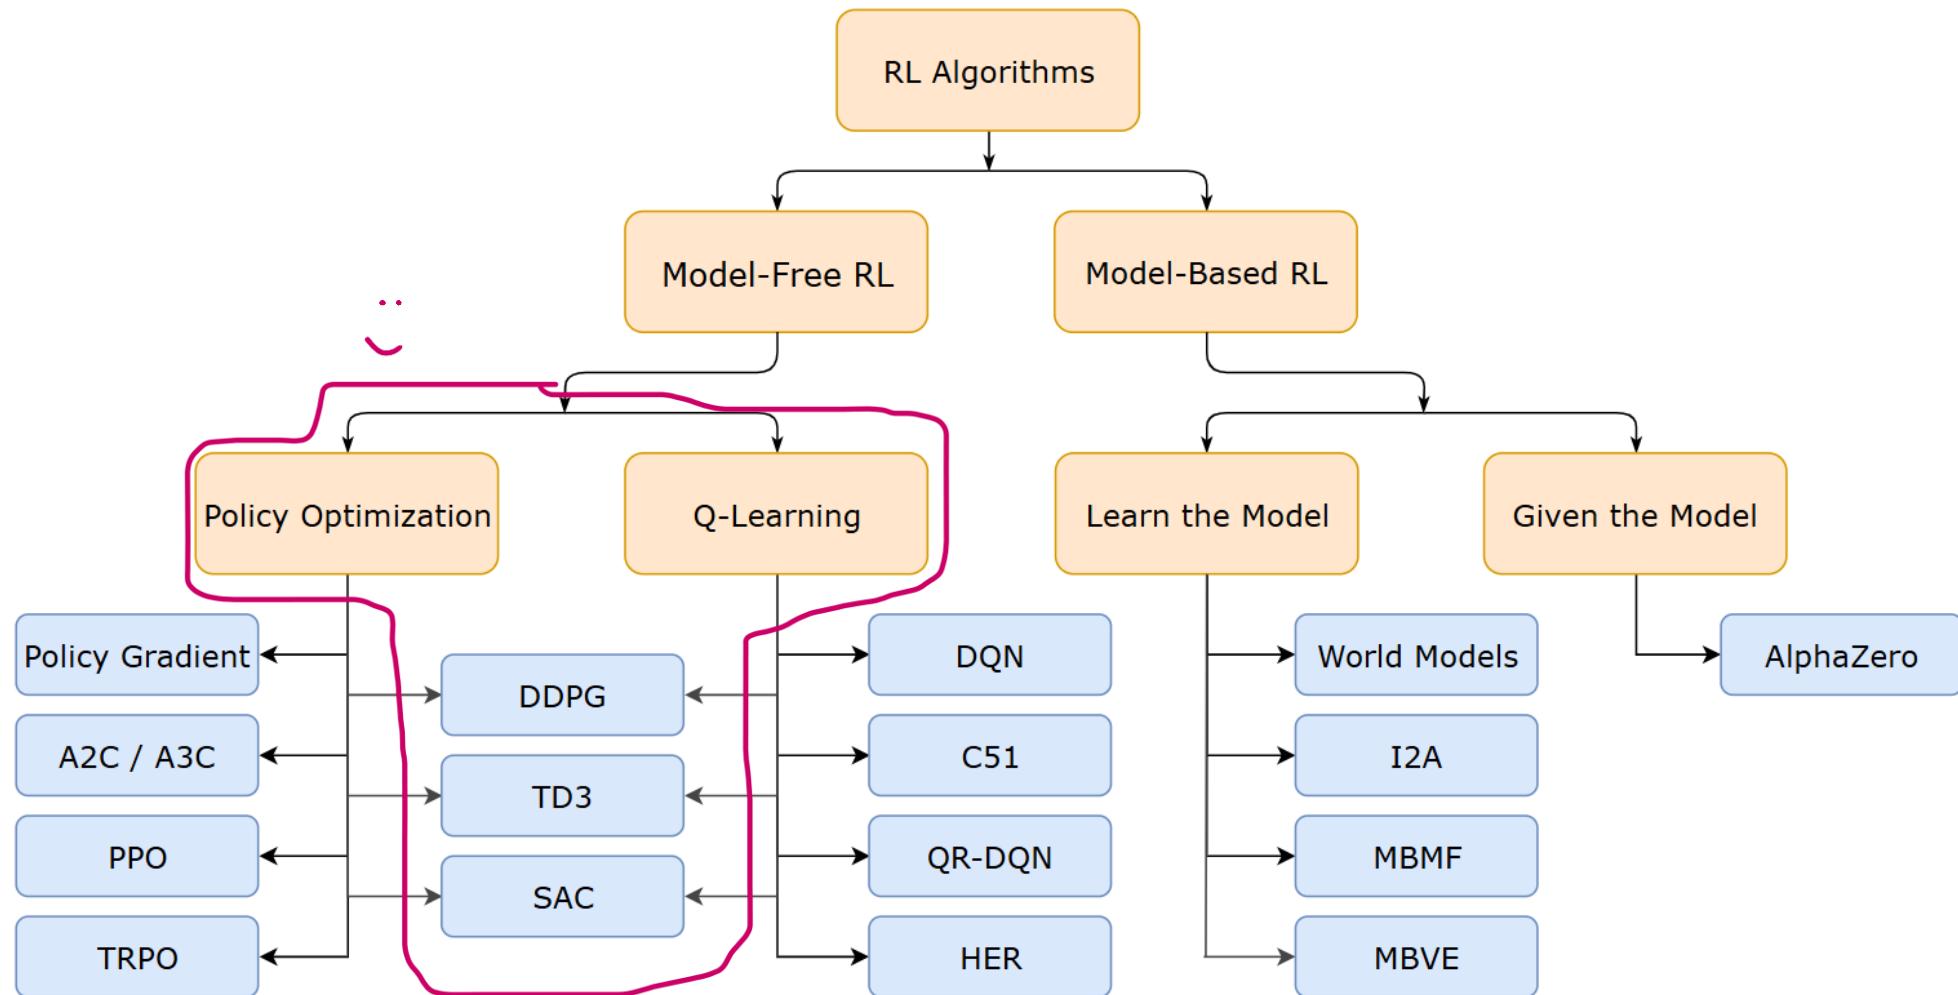
Dr. Saeed Shamaghdari

**Electrical Engineering Department
Control Group**

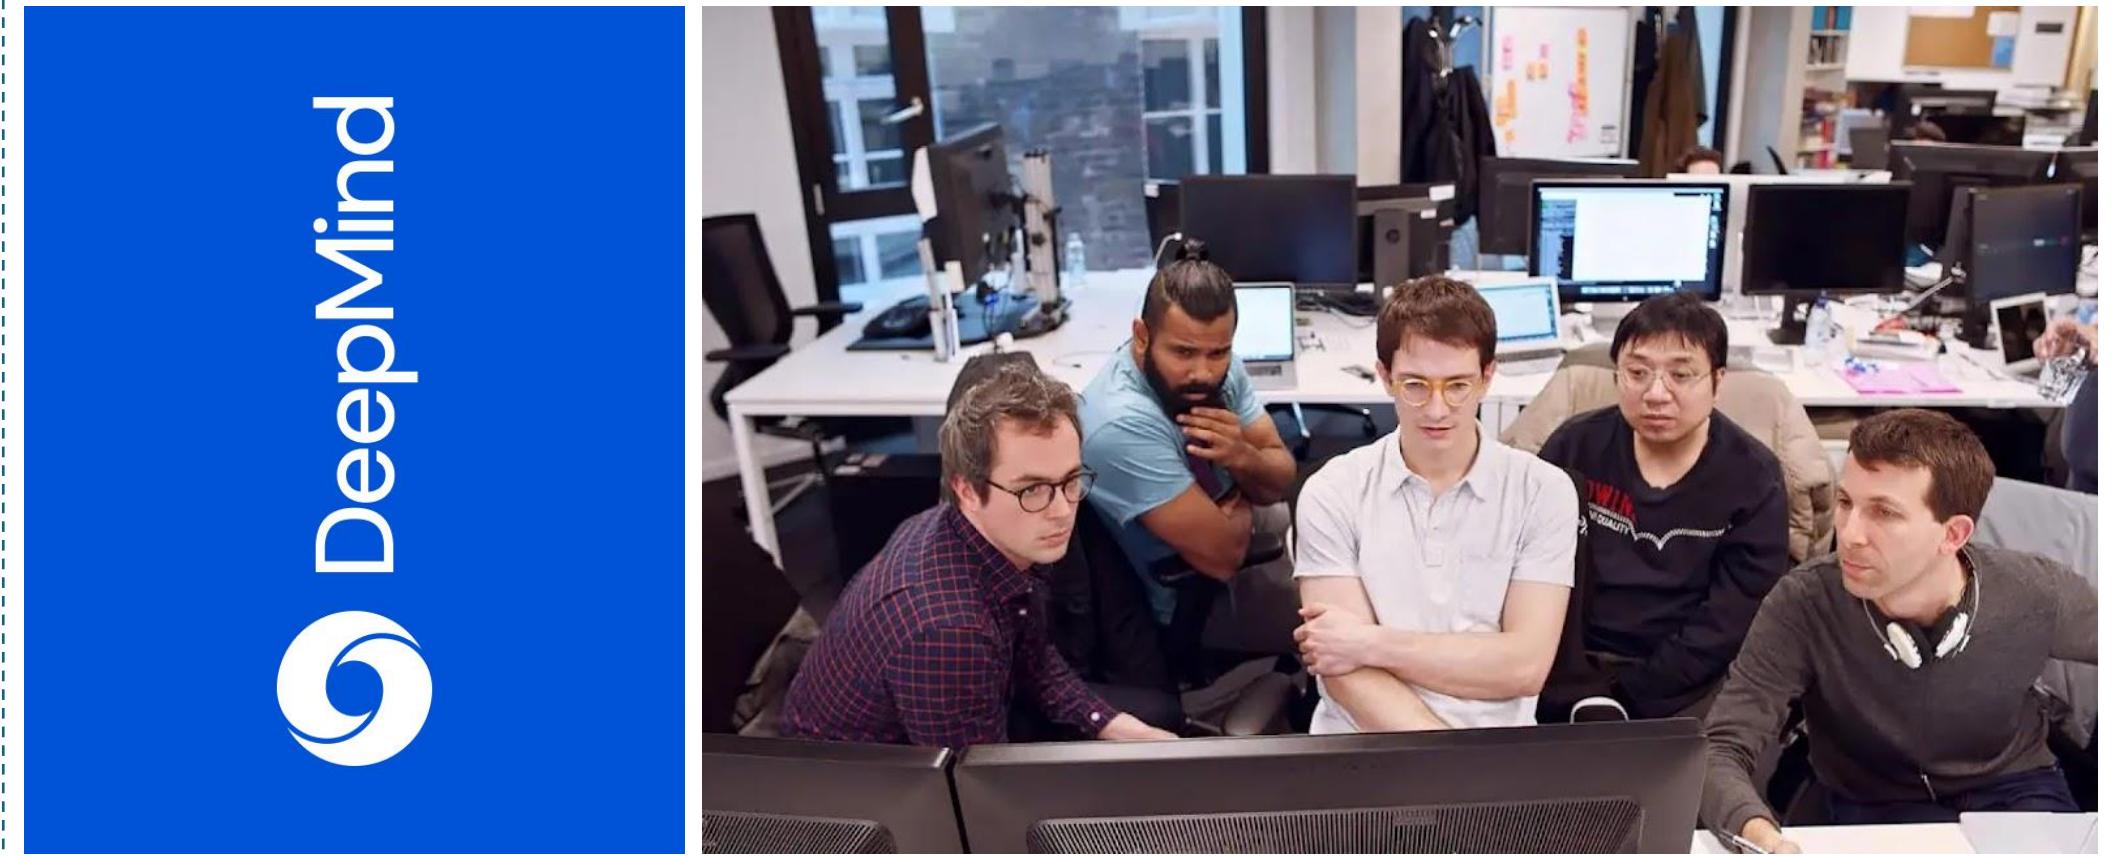
Fall 2025 | 4041

(Deep) Deterministic Policy Gradient

RL Algorithms



Once Again, Google DeepMind!



Deterministic Policy Gradient

Deterministic Policy Gradient Algorithms

David Silver

DeepMind Technologies, London, UK

DAVID@DEEPMIND.COM

Guy Lever

University College London, UK

GUY.LEVER@UCL.AC.UK

Nicolas Heess, Thomas Degris, Daan Wierstra, Martin Riedmiller

DeepMind Technologies, London, UK

*@DEEPMIND.COM

Deterministic Policy Gradient Algorithms ([Paper](#))

Deterministic Policy Gradient

exploration?

Deterministic
Policy Gradients



Go Right

- ✓ Continuous Action Space
- ✓ Off-Policy
- ✓ Stochastic Behaviour Policy

Deterministic Policy Gradient

$$\pi_{\theta}(a|s) = \mathbb{P}[a|s; \theta]$$

Reminder Box

Stochastic Policy

Adjusting policy parameters using
gradient ascent

Policy Gradients



Go Right

Deterministic Policy Gradient

“In this paper we instead consider *deterministic* policies”

$$a = \mu_{\theta}(s)$$

Q: Does a deterministic policy gradient actually exist?

“yes! *Deterministic policy gradient is the limiting case, as policy variance tends to zero, of the stochastic policy gradient.*”



Stochastic PG vs Deterministic PG

Stochastic PG integrates over **states + actions** 

Deterministic PG integrates only over **states** 

Needs more samples!

More sample efficient!

Q: Does this mean stochastic policies are no longer useful?
What happens to **exploration**?

Off-Policy DPG!

Quick Review of RL

A Markov Decision Process with: An action Space \mathcal{A} and a state space \mathcal{S}

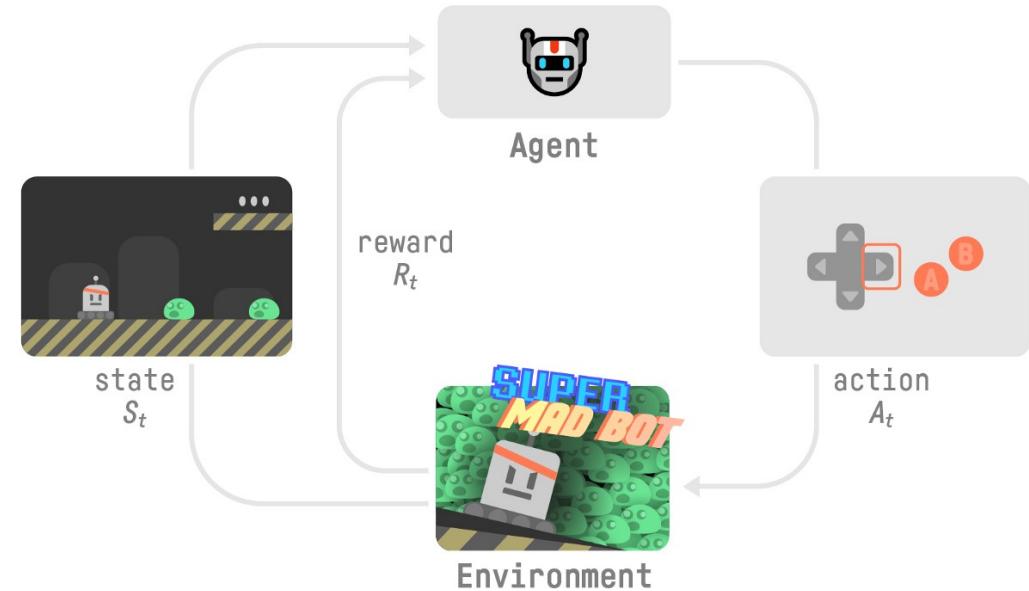
$$p(s_{t+1}|s_1, a_1, \dots, s_t, a_t) = p(s_{t+1}|s_t, a_t)$$

A reward function $r: \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$

A policy

$$\pi_\theta : \mathcal{S} \rightarrow \mathcal{P}(\mathcal{A})$$

is the set of probability measures on \mathcal{A}



Quick Review of RL

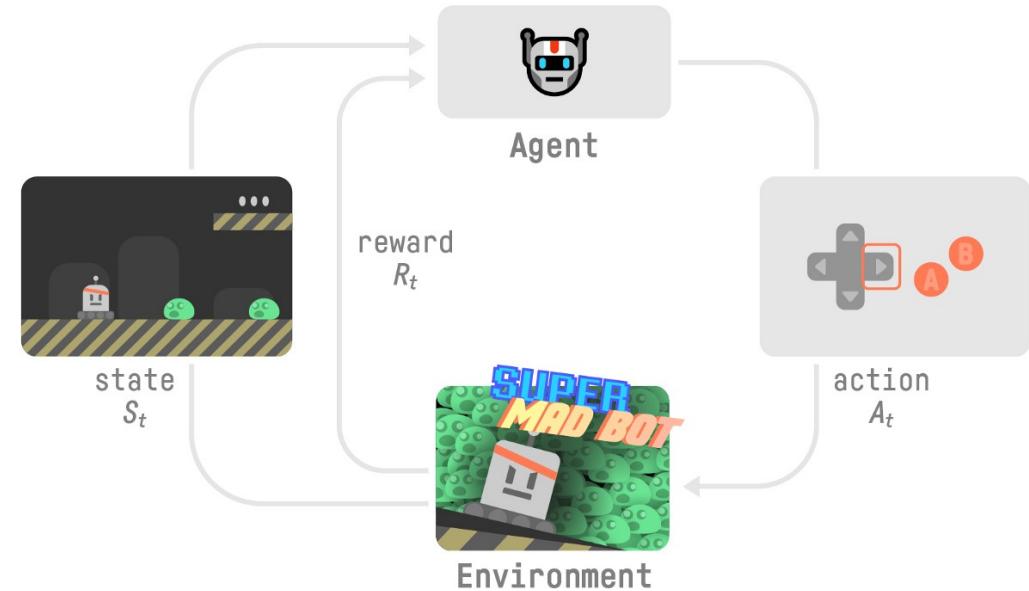
$h_{1:T} = s_1, a_1, r_1, \dots, s_T, a_T, r_T$ over $\mathcal{S} \times \mathcal{A} \times \mathbb{R}$ *→ Trajectory*

$$r_t^\gamma = \sum_{k=t}^{\infty} \gamma^{k-t} r(s_k, a_k) \text{ return}$$

$$V^\pi(s) = \mathbb{E}[r_1^\gamma | S_1 = s; \pi]$$

$$Q^\pi(s, a) = \mathbb{E}[r_1^\gamma | S_1 = s, A_1 = a; \pi]$$

Value functions



Quick Review of RL

Performance Objective:

$$J(\pi) = \mathbb{E} [r_1^\gamma | \pi]$$

(Improper) Discounted state distribution:

$$\rho^\pi(s') := \int_{\mathcal{S}} \sum_{t=1}^{\infty} \gamma^{t-1} p_1(s) \underbrace{p(s \rightarrow s', t, \pi)}_{\text{(the density at state } s' \text{ after transitioning for } t \text{ time steps from state } s)} ds$$

Then we can write:

$$\begin{aligned} J(\pi_\theta) &= \int_{\mathcal{S}} \rho^\pi(s) \int_{\mathcal{A}} \pi_\theta(s, a) r(s, a) da ds \\ &= \mathbb{E}_{s \sim \rho^\pi, a \sim \pi_\theta} [r(s, a)] \end{aligned}$$

Quick Review of RL

Stochastic Policy Gradient:

$$\begin{aligned}\nabla_{\theta} J(\pi_{\theta}) &= \int_{\mathcal{S}} \rho^{\pi}(s) \int_{\mathcal{A}} \nabla_{\theta} \pi_{\theta}(a|s) Q^{\pi}(s, a) da ds \\ &= \mathbb{E}_{s \sim \rho^{\pi}, a \sim \pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a|s) Q^{\pi}(s, a)]\end{aligned}$$

Stochastic Actor-Critic:

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{s \sim \rho^{\pi}, a \sim \pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a|s) Q^{\textcolor{brown}{w}}(s, a)]$$

Off-Policy Actor-Critic

It is often useful to estimate the policy gradient *off-policy* from trajectories sampled from a **distinct behaviour policy** $\beta(a|s) \neq \pi_\theta(a|s)$

$$\begin{aligned} J_\beta(\pi_\theta) &= \int_{\mathcal{S}} \rho^\beta(s) V^\pi(s) ds \\ &= \int_{\mathcal{S}} \int_{\mathcal{A}} \rho^\beta(s) \pi_\theta(a|s) Q^\pi(s, a) da ds \end{aligned}$$

$$\begin{aligned} \nabla_\theta J_\beta(\pi_\theta) &\approx \int_{\mathcal{S}} \int_{\mathcal{A}} \rho^\beta(s) \nabla_\theta \pi_\theta(a|s) Q^\pi(s, a) da ds \\ &= \mathbb{E}_{s \sim \rho^\beta, a \sim \beta} \left[\frac{\pi_\theta(a|s)}{\beta_\theta(a|s)} \nabla_\theta \log \pi_\theta(a|s) Q^\pi(s, a) \right] \end{aligned}$$

Off-Policy Actor-Critic

It is often useful to estimate the policy gradient *off-policy* from trajectories sampled from a **distinct behaviour policy** $\beta(a|s) \neq \pi_\theta(a|s)$

$$\begin{aligned} J_\beta(\pi_\theta) &= \int_{\mathcal{S}} \rho^\beta(s) V^\pi(s) ds \\ &= \int_{\mathcal{S}} \int_{\mathcal{A}} \rho^\beta(s) \pi_\theta(a|s) Q^\pi(s, a) da ds \end{aligned}$$

$$\begin{aligned} \nabla_\theta J_\beta(\pi_\theta) &\approx \int_{\mathcal{S}} \int_{\mathcal{A}} \rho^\beta(s) \nabla_\theta \pi_\theta(a|s) Q^\pi(s, a) da ds \\ &= \mathbb{E}_{s \sim \rho^\beta, a \sim \beta} \left[\frac{\pi_\theta(a|s)}{\beta_\theta(a|s)} \nabla_\theta \log \pi_\theta(a|s) Q^\pi(s, a) \right] \end{aligned}$$

Off-Policy Actor-Critic

$$\begin{aligned}
 \nabla_{\theta} J_{\beta}(\pi_{\theta}) &\approx \int_{\mathcal{S}} \int_{\mathcal{A}} \rho^{\beta}(s) \nabla_{\theta} \pi_{\theta}(a|s) Q^{\pi}(s, a) da ds \\
 &= E \left[\int_{\mathcal{A}} \nabla_{\theta} \pi_{\theta}(a|s) Q^{\pi}(s, a) \mid s \sim \rho^{\beta} \right] \\
 &= E \left[\underbrace{\int_{\mathcal{A}} \beta_{\theta}(a|s)}_{E[\dots | a \sim \beta]} \frac{\pi_{\theta}(a|s)}{\beta_{\theta}(a|s)} \underbrace{\frac{\nabla_{\theta} \pi_{\theta}(a|s)}{\pi_{\theta}(a|s)}}_{\nabla_{\theta} \log \pi_{\theta}(a|s)} Q^{\pi}(s, a) \mid s \sim \rho^{\beta} \right] \\
 &\quad \xrightarrow{\text{Green Arrow}} E \left[\frac{\pi_{\theta}(a|s)}{\beta_{\theta}(a|s)} \nabla_{\theta} \log \pi_{\theta}(a|s) Q^{\pi}(s, a) \mid s \sim \rho^{\beta}, a \sim \beta \right] \\
 &= \mathbb{E}_{s \sim \rho^{\beta}, a \sim \beta} \left[\frac{\pi_{\theta}(a|s)}{\beta_{\theta}(a|s)} \nabla_{\theta} \log \pi_{\theta}(a|s) Q^{\pi}(s, a) \right]
 \end{aligned}$$

Gradients of Deterministic Policies: Action-Value Gradients

Policy Improvement:

Reminder Box

$$\mu^{k+1}(s) = \operatorname{argmax}_a Q^{\mu^k}(s, a)$$

Q: What about continuous action space?

*“Requiring a global maximisation at every step
Instead, move the policy in the direction of the gradient of Q ”*

Gradients of Deterministic Policies: Action-Value Gradients

For each visited state s , the policy parameters θ^{k+1} , are updated in proportion to the gradient $\nabla_{\theta} Q^{\mu^k}(s, \mu_{\theta}(s))$.

Each state suggests a different direction of policy improvement; these may be averaged together by taking an expectation with respect to the state distribution $\rho^{\mu}(s)$.

$$\theta^{k+1} = \theta^k + \alpha \mathbb{E}_{s \sim \rho^{\mu^k}} \left[\nabla_{\theta} Q^{\mu^k}(s, \mu_{\theta}(s)) \right]$$

chain rule

$$\theta^{k+1} = \theta^k + \alpha \mathbb{E}_{s \sim \rho^{\mu^k}} \left[\nabla_{\theta} \mu_{\theta}(s) \left. \nabla_a Q^{\mu^k}(s, a) \right|_{a=\mu_{\theta}(s)} \right]$$

Deterministic Policy Gradient Theorem

A deterministic policy

$$\mu_{\theta}: s \rightarrow A, \theta \in \mathcal{R}'$$

Performance Objective

$$J(\mu_{\theta}), E[r_t^{\gamma} | \mu]$$

Probability Distribution

$$\rho(s \rightarrow s', t, \mu)$$

Discounted state distribution

$$\rho^{\mu}(s)$$

Then:

$$\begin{aligned} J(\mu_{\theta}) &= \int_{\mathcal{S}} \rho^{\mu}(s) r(s, \mu_{\theta}(s)) ds \\ &= \mathbb{E}_{s \sim \rho^{\mu}} [r(s, \mu_{\theta}(s))] \end{aligned}$$

Deterministic Policy Gradient Theorem

Theorem 1 (Deterministic Policy Gradient Theorem)

Suppose that the MDP satisfies condition below:

$p(s'|s, a)$, $\nabla_a p(s'|s, a)$, $\mu_\theta(s)$, $\nabla_\theta \mu_\theta(s)$, $r(s, a)$, $\nabla_a r(s, a)$, $p_1(s)$ are continuous in all parameters and variables s, a, s', x .

These imply that $\nabla_\theta \mu_\theta(s)$ and $\nabla_a Q^\mu(s, a)$ exist and that the deterministic policy gradient exists. Then

$$\begin{aligned}\nabla_\theta J(\mu_\theta) &= \int_{\mathcal{S}} \rho^\mu(s) \nabla_\theta \mu_\theta(s) \left. \nabla_a Q^\mu(s, a) \right|_{a=\mu_\theta(s)} ds \\ &= \mathbb{E}_{s \sim \rho^\mu} \left[\nabla_\theta \mu_\theta(s) \left. \nabla_a Q^\mu(s, a) \right|_{a=\mu_\theta(s)} \right]\end{aligned}$$

Limit of the Stochastic Policy Gradient

For a wide class of stochastic policies, including many bump functions, the deterministic policy gradient is indeed a special (limiting) case of the stochastic policy gradient.

Theorem 2

Consider a stochastic policy $\pi_{\mu_\theta, \sigma}$ where μ_θ is a deterministic policy and σ is a parameter controlling the variance. Then,

$$\lim_{\sigma \downarrow 0} \nabla_\theta J(\pi_{\mu_\theta, \sigma}) = \nabla_\theta J(\mu_\theta)$$

$$\begin{aligned}\sigma &\rightarrow 0 \\ \pi_{\mu_\theta, 0} &\equiv \mu_\theta\end{aligned}$$

Where on the l.h.s. the gradient is the standard stochastic policy gradient and on the r.h.s. the gradient is the deterministic policy gradient.

Deterministic Actor-Critic: On-Policy Q: Optimality?

Algorithm: On-Policy Actor-Critic

In this *deterministic actor-critic* algorithm, the critic uses Sarsa updates to estimate the action-value function.

$$\delta_t = r_t + \gamma Q^w(s_{t+1}, a_{t+1}) - Q^w(s_t, a_t)$$

$$w_{t+1} = w_t + \alpha_w \delta_t \nabla_w Q^w(s_t, a_t) \quad \text{Critic}$$

$$\theta_{t+1} = \theta_t + \alpha_\theta \nabla_\theta \mu_\theta(s_t) \left. \nabla_a Q^w(s_t, a_t) \right|_{a=\mu_\theta(s)} \quad \text{Actor}$$

Deterministic Actor-Critic: Off-Policy

Learn a deterministic target policy $\mu_\theta(s)$, from trajectories generated by an arbitrary stochastic behaviour policy $\pi(a|s) = \beta(a|s)$.

$$\begin{aligned} J_\beta(\mu_\theta) &= \int_{\mathcal{S}} \rho^\beta(s) V^\mu(s) ds \\ &= \int_{\mathcal{S}} \rho^\beta(s) Q^\mu(s, \mu_\theta(s)) ds \end{aligned}$$

$$\begin{aligned} \nabla_\theta J_\beta(\mu_\theta) &\approx \int_{\mathcal{S}} \rho^\beta(s) \nabla_\theta \mu_\theta(a|s) Q^\mu(s, a) ds \\ &= \mathbb{E}_{s \sim \rho^\beta} \left[\nabla_\theta \mu_\theta(s) \left. \nabla_a Q^\mu(s, a) \right|_{a=\mu_\theta(s)} \right] \end{aligned}$$

Deterministic Actor-Critic: Off-Policy Q: Importance Sampling?

“We now develop an actor-critic algorithm that updates the policy in the direction of the off-policy deterministic policy gradient.”

Algorithm: On-Policy Actor-Critic

Substitute a differentiable action-value function $Q^w(s, a)$ in place of the true action-value function $Q^\mu(s, a)$. A critic estimates the action-value function $Q^w(s, a) \approx Q^\mu(s, a)$ off-policy from trajectories generated by $\beta(a|s)$, and uses Q-learning updates to estimate the action-value function.

$$\delta_t = r_t + \gamma Q^w(s_{t+1}, \mu_\theta(s_{t+1})) - Q^w(s_t, a_t)$$

$$w_{t+1} = w_t + \alpha_w \delta_t \nabla_w Q^w(s_t, a_t) \quad \text{Critic}$$

$$\theta_{t+1} = \theta_t + \alpha_\theta \nabla_\theta \mu_\theta(s_t) \left. \nabla_a Q^w(s_t, a_t) \right|_{a=\mu_\theta(s)} \quad \text{Actor}$$

I Deterministic Policy Gradient

For further details on the algorithms, including compatible function approximation and related methods, please refer to the [original paper](#). We will not go into more depth here.

Deep Deterministic Policy Gradient

Published as a conference paper at ICLR 2016

CONTINUOUS CONTROL WITH DEEP REINFORCEMENT LEARNING

Timothy P. Lillicrap,* Jonathan J. Hunt,* Alexander Pritzel, Nicolas Heess,

Tom Erez, Yuval Tassa, David Silver & Daan Wierstra

Google Deepmind

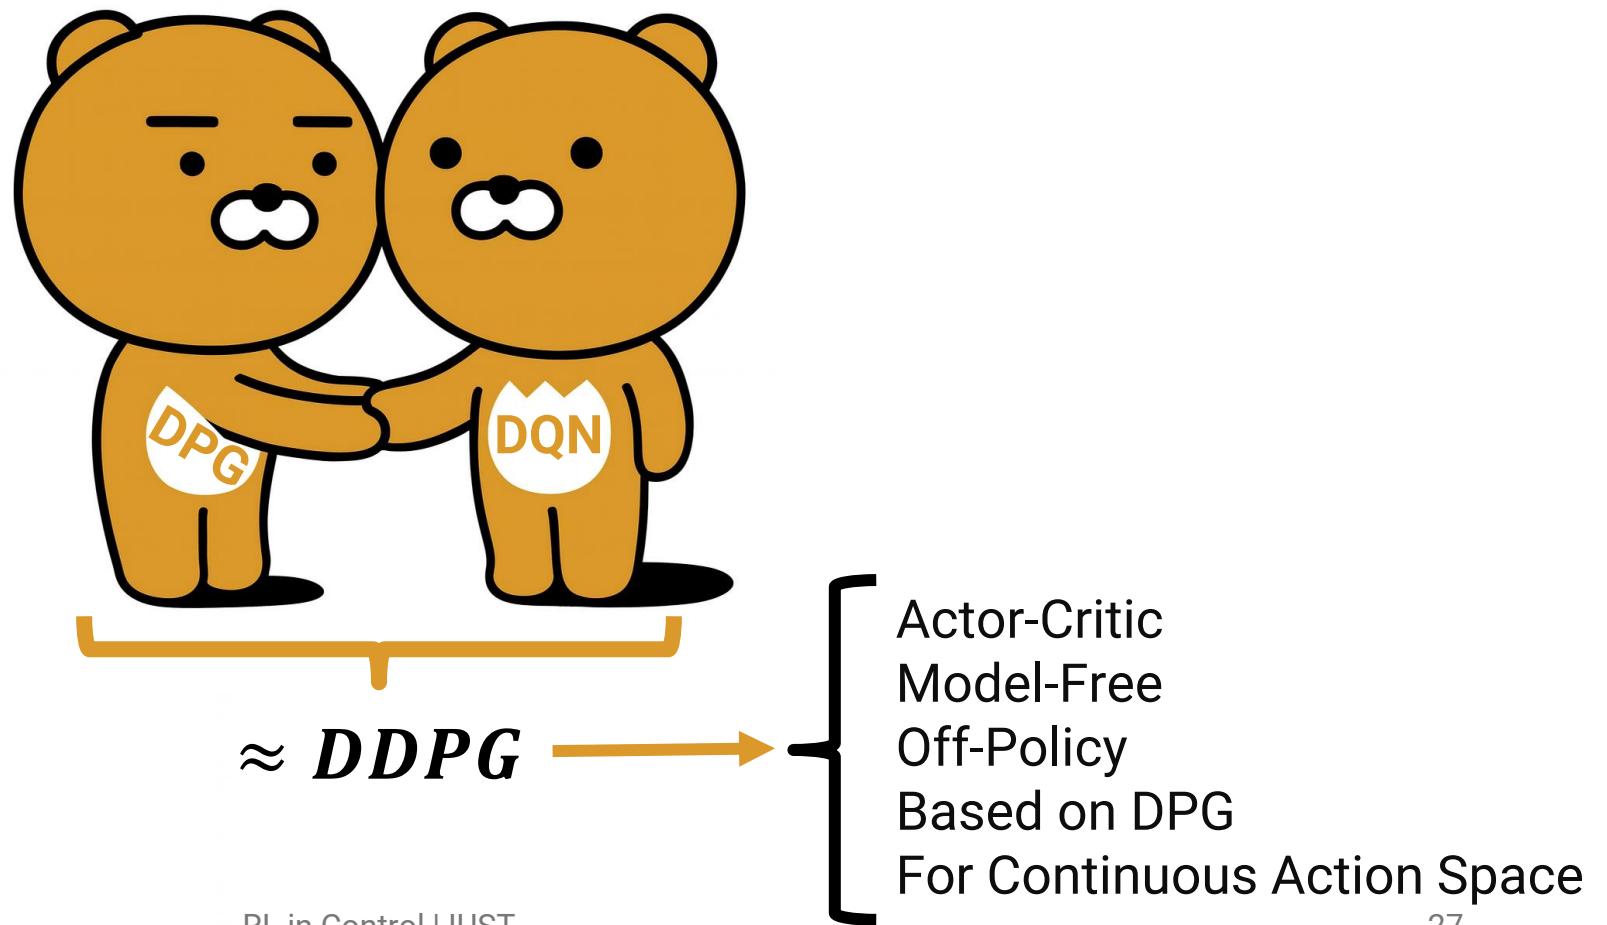
London, UK

{countzero, jjhunt, apritzel, heess,
erez, tassa, davidsilver, wierstra} @ google.com

Continuous Control with Deep Reinforcement Learning ([Paper](#))

Deep Deterministic Policy Gradient

learn policies “end-to-end”: directly from raw pixel inputs.



DQN: Discussion Questions

Q: What is the main limitation of standard DQN?

A: Discrete / low dim Action Space!

Q: Why can't DQN be directly applied to continuous action spaces?

A: Maximization!

Well, discretize the action space! 😊



Q: What problems arise when discretizing continuous action spaces, and why is it inefficient?

A: Curse of dimensionality

Exploration, ...

Example: 7 DOF \times 3 values each:
 $3^7 = 2187$ actions.

DQN: Discussion Questions

Q: Is DQN only limited and problematic, or does it also have important advantages?

A: Replay buffer }
Target network } Stability in training 😎

DDPG

In the DPG paper, we saw

$$\begin{aligned}\nabla_{\theta} J_{\beta}(\mu_{\theta}) &\approx \int_{\mathcal{S}} \rho^{\beta}(s) \nabla_{\theta} \mu_{\theta}(a|s) Q^{\mu}(s, a) ds \\ &= \mathbb{E}_{s \sim \rho^{\beta}} \left[\nabla_{\theta} \mu_{\theta}(s) \left. \nabla_a Q^{\mu}(s, a) \right|_{a=\mu_{\theta}(s)} \right]\end{aligned}$$

In the DDPG paper, the notation is slightly **different**:

$$\begin{aligned}\nabla_{\theta^{\mu}} J &\approx \mathbb{E}_{s_t \sim \rho^{\beta}} \left[\nabla_{\theta^{\mu}} Q(s, a | \theta^Q) \Big|_{s=s_t, a=\mu(s_t | \theta^{\mu})} \right] \\ &= \mathbb{E}_{s_t \sim \rho^{\beta}} \left[\nabla_a Q(s, a | \theta^Q) \Big|_{s=s_t, a=\mu(s_t)} \nabla_{\theta^{\mu}} \mu(s | \theta^{\mu}) \Big|_{s=s_t} \right]\end{aligned}$$

Reminder Box

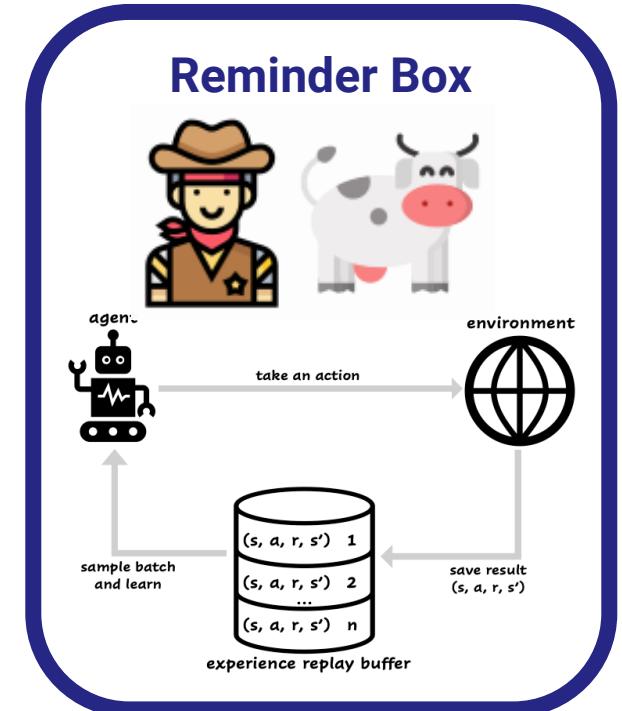
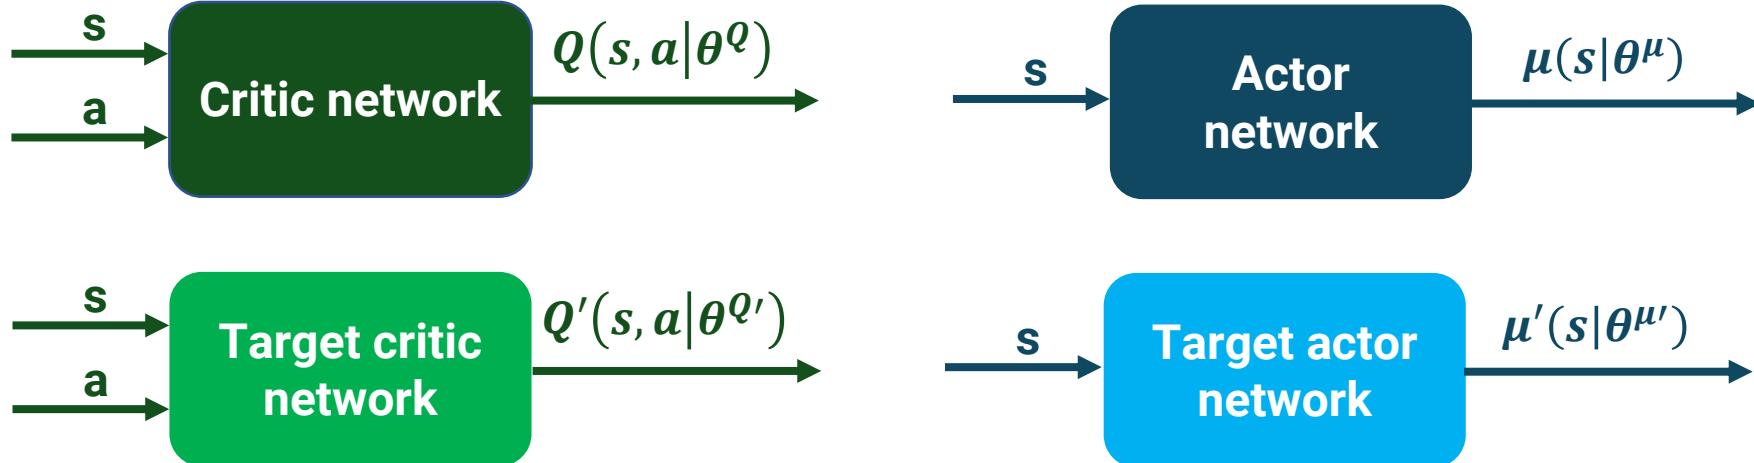
$$\begin{aligned}R_t &= \sum_{i=t}^T \gamma^{(i-t)} r(s_i, a_i) \\ J &= \mathbb{E}_{r_i, s_i \sim E, a_i \sim \pi} [R_1]\end{aligned}$$

I Deep Deterministic Policy Gradient

DDPG *→ 4 NNs + Batch Norm technique (?)*

Similar to DQN, we also use:

- separate target networks, **Q**: Why?
- and a replay buffer here



The weights of these target networks are then updated by having them slowly track the learned networks: $\theta' \leftarrow \tau\theta + (1 - \tau)\theta'$, $\tau \ll 1$

Exploration Challenge

“A major challenge of learning in continuous action spaces is exploration. An advantage of off-policies algorithms such as DDPG is that we can treat the problem of exploration independently from the learning algorithm. We constructed an exploration policy μ' by adding noise sampled from a noise process \mathcal{N} to our actor policy

$$\mu'(s_t) = \mu(s_t | \theta_t^\mu) + \mathcal{N}$$

\mathcal{N} can be chosen to suit the environment.”

DDPG: Algorithm

Algorithm 1 DDPG algorithm

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ .
 Initialize target network Q' and μ' with weights $\theta^{Q'} \leftarrow \theta^Q, \theta^{\mu'} \leftarrow \theta^\mu$
 Initialize replay buffer R

for episode = 1, M **do**

- Initialize a random process \mathcal{N} for action exploration
- Receive initial observation state s_1
- for** t = 1, T **do**

 - Select action $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$ according to the current policy and exploration noise
 - Execute action a_t and observe reward r_t and observe new state s_{t+1}
 - Store transition (s_t, a_t, r_t, s_{t+1}) in R
 - Sample a random minibatch of N transitions (s_i, a_i, r_i, s_{i+1}) from R
 - Set $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$
 - Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$
 - Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

Update the target networks:

$$\begin{aligned}\theta^{Q'} &\leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'} \\ \theta^{\mu'} &\leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}\end{aligned}$$

end for
end for

End of Part I



R L^E Control