

# یادگیری تقویتی در کنترل

دکتر سعید شمقدری

دانشکده مهندسی برق  
گروه کنترل

نیمسال اول ۱۴۰۵-۱۴۰۴

# Temporal Difference Learning

# Temporal Difference Learning

## Q Box

مشابهت با DP  
مشابهت با MC

## Reminder Box

تخمين تابع Value در DP و MC

$$\underline{V(S_t)} \leftarrow \underline{V(S_t)} + \alpha [\underline{G_t} - \underline{V(S_t)}]$$

New value of state t

Former estimation  
of value of state t  
(= Expected return  
starting at that state)

Learning  
Rate

Return at  
timestep  
t

Former estimation  
of value of state t  
(= Expected return  
starting at that state)

روش MC در Constant alpha (نیاز به اتمام اپیزود)  
استفاده از Return به عنوان تخمین امید ریاضی (Sampling)

## Temporal Difference Prediction

تخمین تابع Value در TD

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

New value  
of state t

Former  
estimation of  
value of state  
t

Learning Rate

Discounted value of next  
state

**Q Box:**

Target in MC and TD?

## Temporal Difference Prediction

تخمین تابع Value در TD

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

New value  
of state t

Former  
estimation of  
value of state  
t

Learning Rate

Discounted value of next  
state

TD Target

TD(0): one step TD

**Q Box:**  
Target in MC and TD?

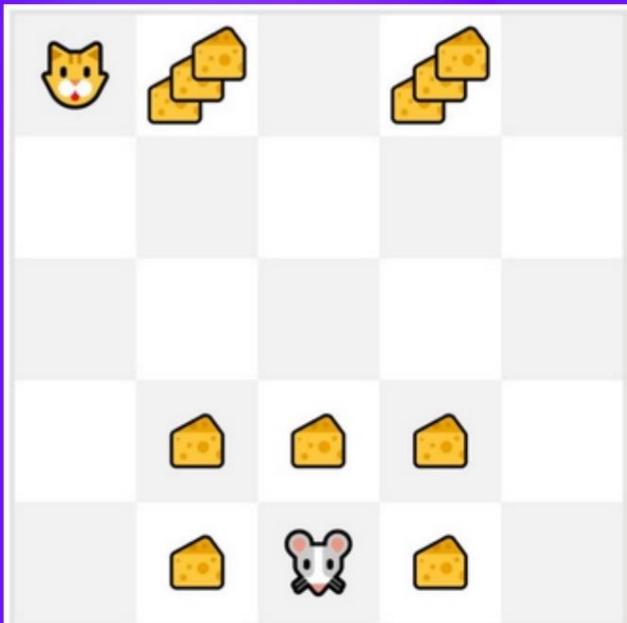
## TD Learning Approach:

*Temporal Difference Learning:* learning at each time step.

$$\underline{V(S_t)} \leftarrow \underline{V(S_t)} + \alpha [R_{t+1} + \gamma \underline{V(S_{t+1})} - \underline{V(S_t)}]$$

New value of state t      Former estimation of value of state t      Learning Rate      Reward      Discounted value of next state  
TD Target

## TD Approach:



At the end of one step (State, Action, Reward, Next State):

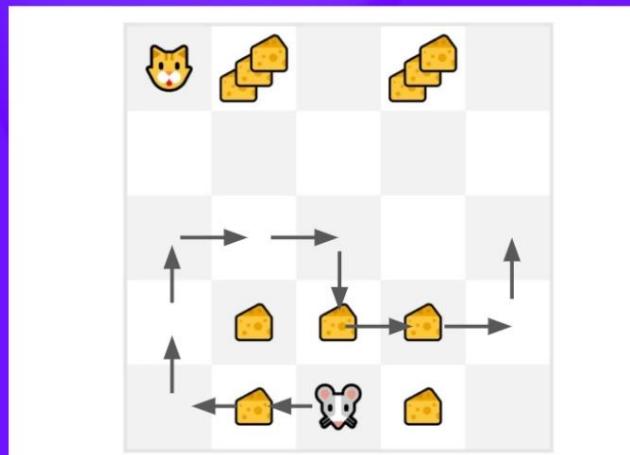
- We have  $R_{t+1}$  and  $S_{t+1}$
- We update  $V(S_t)$ :
  - We estimate  $G_t$  by adding  $R_{t+1}$  and the discounted value of next state.  
**TD target** :  $R_{t+1} + \gamma V(S_{t+1})$

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

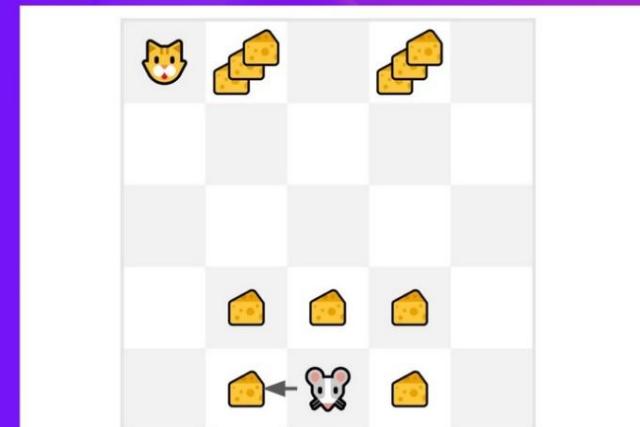
Now we continue to interact with this environment with our updated value function. By running more and more steps, the agent will learn to play better and better.

## Two Learning Approaches:

*Monte Carlo: learning at the end of the episode*



*Temporal Difference Learning: learning at each step.*



## Temporal Difference Prediction

تخمین تابع Value در TD

$$\begin{aligned}
 v_\pi(s) &\doteq \mathbb{E}_\pi[G_t \mid S_t = s] \\
 &= \mathbb{E}_\pi[R_{t+1} + \gamma G_{t+1} \mid S_t = s] \\
 &= \mathbb{E}_\pi[R_{t+1} + \gamma v_\pi(S_{t+1}) \mid S_t = s]
 \end{aligned}$$

- DP : محاسبه امید ریاضی با مدل، استفاده از تخمین قبلي تابع Value در تخمین جديد (bootstrapping)
- TD Sampling : استفاده از تخمین قبلي تابع Value در تخمین جديد (bootstrapping)

## Tabular TD(0) for estimating $v_\pi$

Input: the policy  $\pi$  to be evaluated

Algorithm parameter: step size  $\alpha \in (0, 1]$

Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$

Loop for each episode:

    Initialize  $S$

    Loop for each step of episode:

$A \leftarrow$  action given by  $\pi$  for  $S$

        Take action  $A$ , observe  $R, S'$

$V(S) \leftarrow V(S) + \alpha [R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

    until  $S$  is terminal

## TD Error

$$\delta_t \doteq R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$

## Monte Carlo Error

$$\begin{aligned}
G_t - V(S_t) &= R_{t+1} + \gamma G_{t+1} - V(S_t) + \gamma V(S_{t+1}) - \gamma V(S_{t+1}) \\
&= \delta_t + \gamma(G_{t+1} - V(S_{t+1})) \\
&= \delta_t + \gamma\delta_{t+1} + \gamma^2(G_{t+2} - V(S_{t+2})) \\
&= \delta_t + \gamma\delta_{t+1} + \gamma^2\delta_{t+2} + \cdots + \gamma^{T-t-1}\delta_{T-1} + \gamma^{T-t}(G_T - V(S_T)) \\
&= \delta_t + \gamma\delta_{t+1} + \gamma^2\delta_{t+2} + \cdots + \gamma^{T-t-1}\delta_{T-1} + \gamma^{T-t}(0 - 0) \\
&= \sum_{k=t}^{T-1} \gamma^{k-t} \delta_k.
\end{aligned}$$

**Driving Home****مثال**

Value function approximation TD

<i>State</i>	<i>Elapsed Time</i> (minutes)	<i>Predicted</i> <i>Time to Go</i>	<i>Predicted</i> <i>Total Time</i>
leaving office, friday at 6	0	30	30
reach car, raining	5	35	40
exiting highway	20	15	35
2ndary road, behind truck	30	10	40
entering home street	40	3	43
arrive home	43	0	43

Reward?

**Q Box**

Return? T\_Go

Value? E(T\_Go)

**Driving Home****مثال**

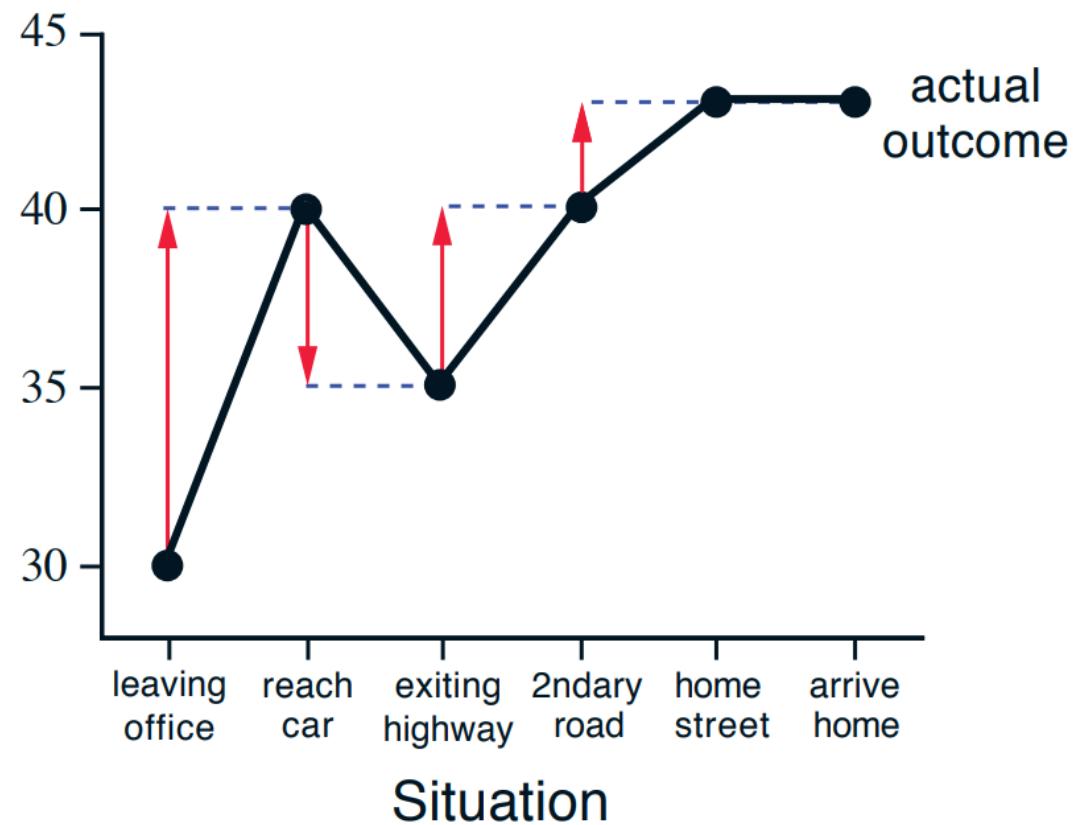
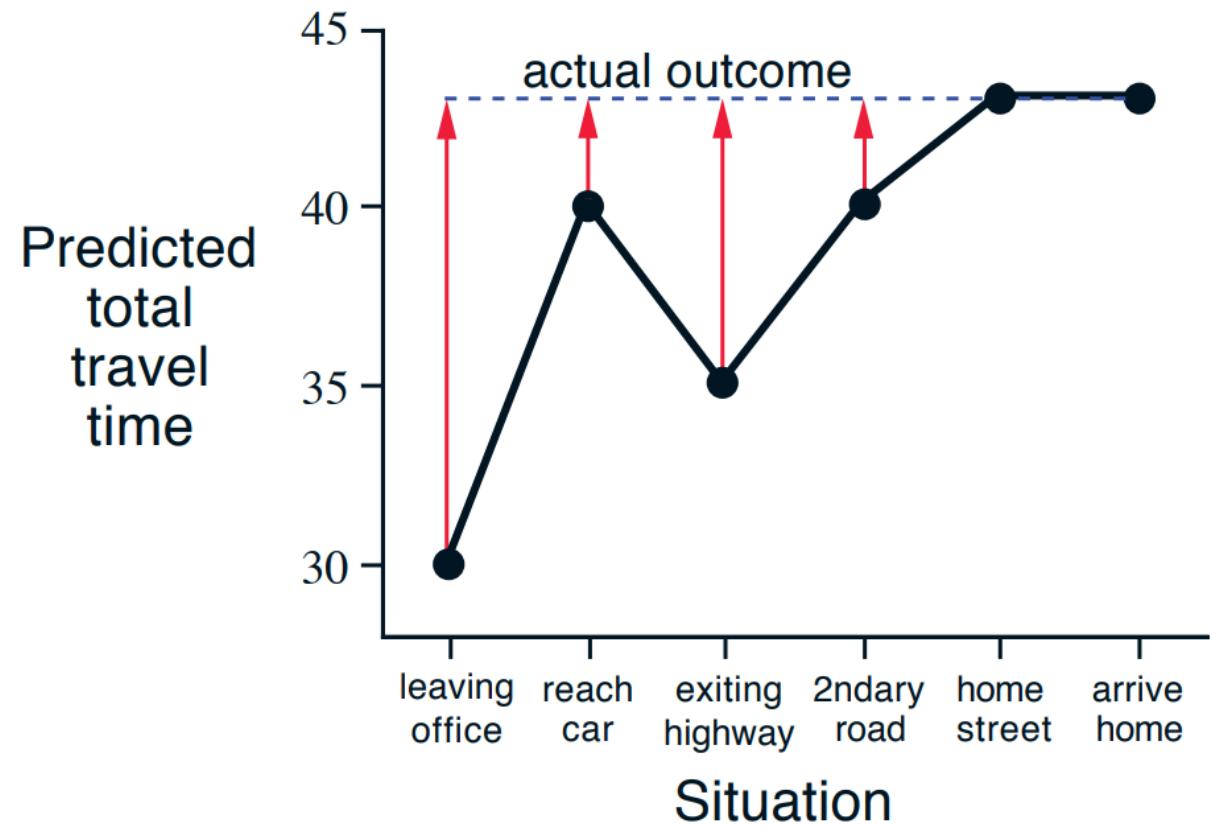
## Value function approximation TD

<i>State</i>	<i>Elapsed Time</i> (minutes)	<i>Predicted</i> <i>Time to Go</i>	<i>Predicted</i> <i>Total Time</i>
leaving office, friday at 6	0	30	30
reach car, raining	5	35	40
exiting highway	20	15	35
2ndary road, behind truck	30	10	40
entering home street	40	3	43
arrive home	43	0	43

**Q Box**

خطای تخمین در خروجی بزرگراه (بر اساس MC)  
اصلاح تخمین در خروجی بزرگراه (بر اساس MC) به ازای  $\alpha=0.5$  به ازای  
 $\alpha(G_t - V(s_t))$ ?

# یادگیری سریعتر: TD



## Temporal Difference Learning

همگرایی TD

- مقدار  $\alpha$  بسیار کوچک
- $\alpha$  با شرایط کاهشی

$$\sum_{n=1}^{\infty} \alpha_n(a) = \infty$$

and

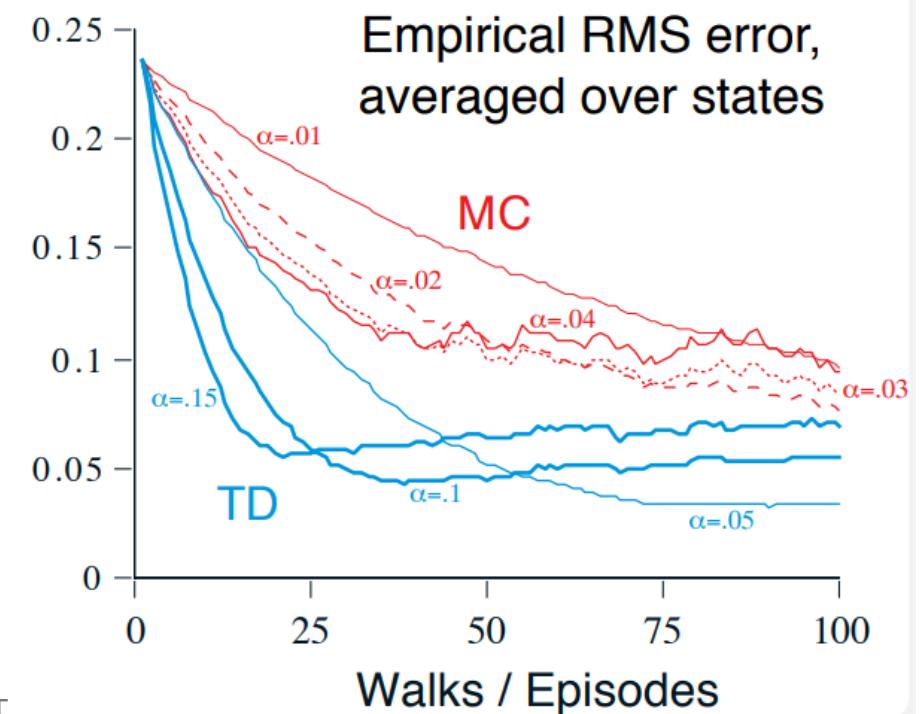
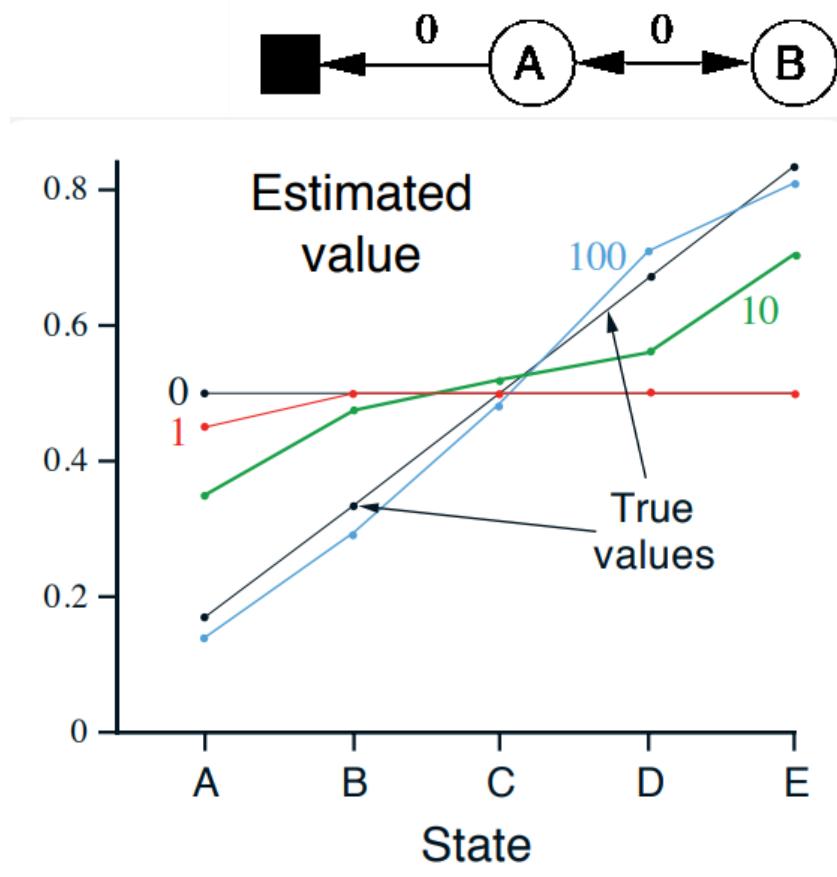
$$\sum_{n=1}^{\infty} \alpha_n^2(a) < \infty$$

سرعت همگرایی MC و TD

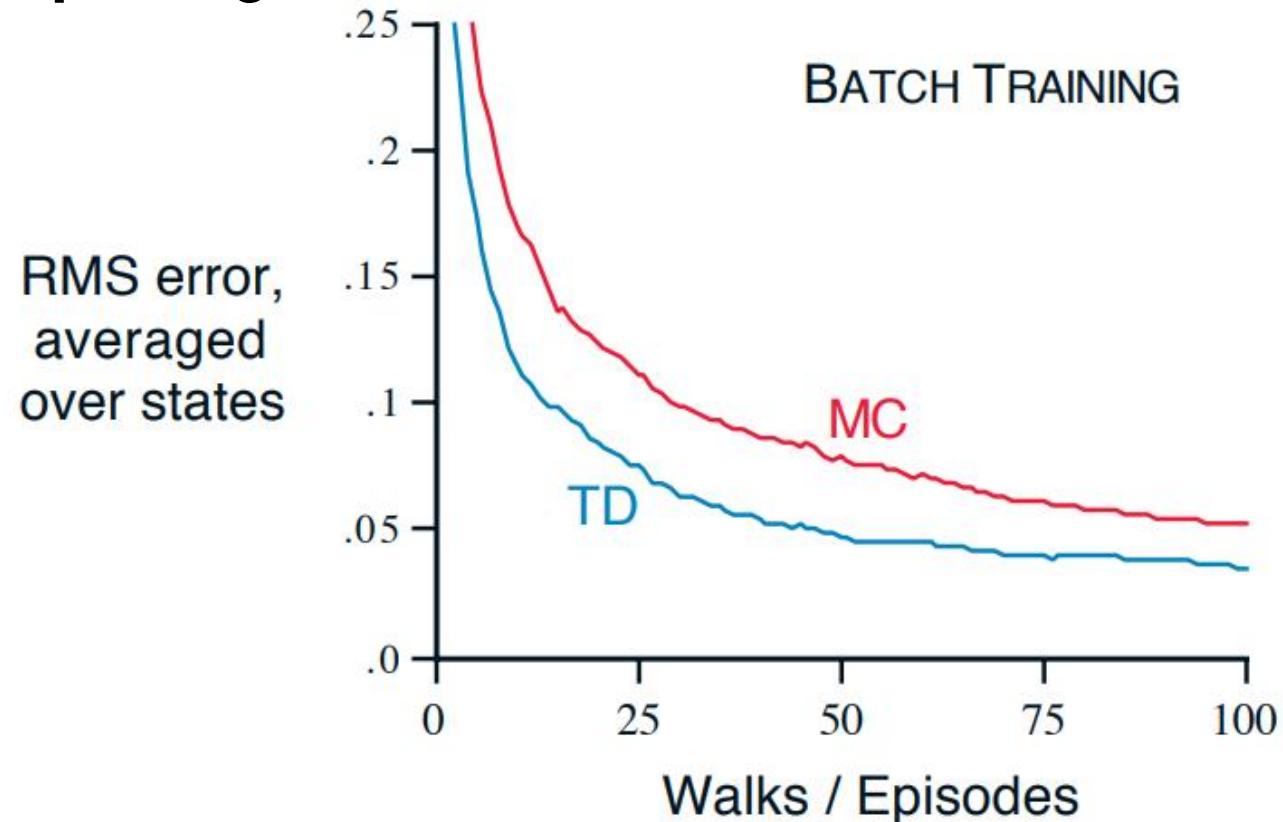
مثال

## Markov Reward Process: Random Walk

In this example we empirically compare the prediction abilities of TD(0) and constant- $\alpha$  MC when applied to the following Markov reward process:



## Random Walk Under Batch Updating



**Figure 6.2:** Performance of TD(0) and constant- $\alpha$  MC under batch training on the random walk task.

## You are the Predictor

مثال

A, 0, B, 0

B, 1

B, 1

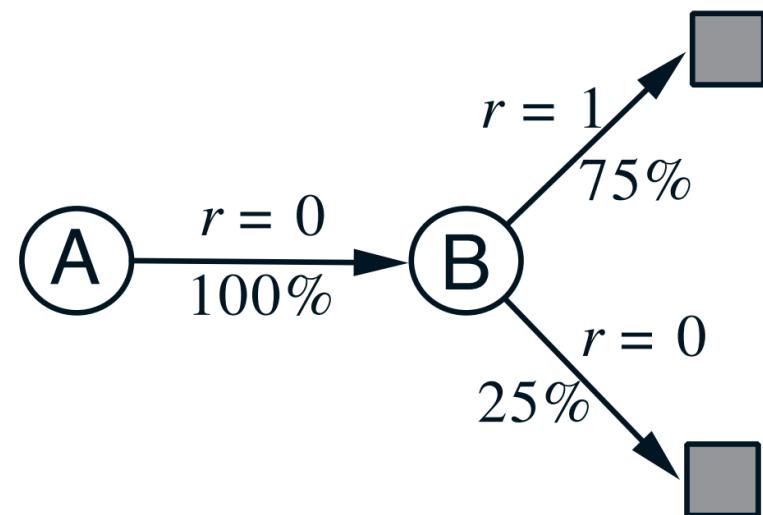
B, 1

B, 1

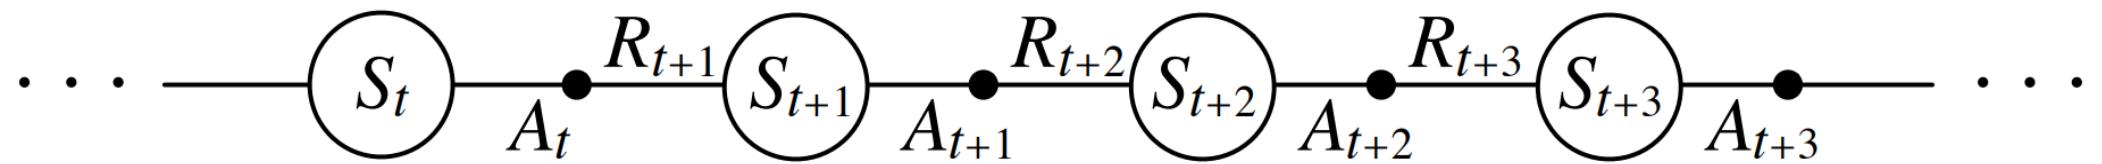
B, 1

B, 1

B, 0

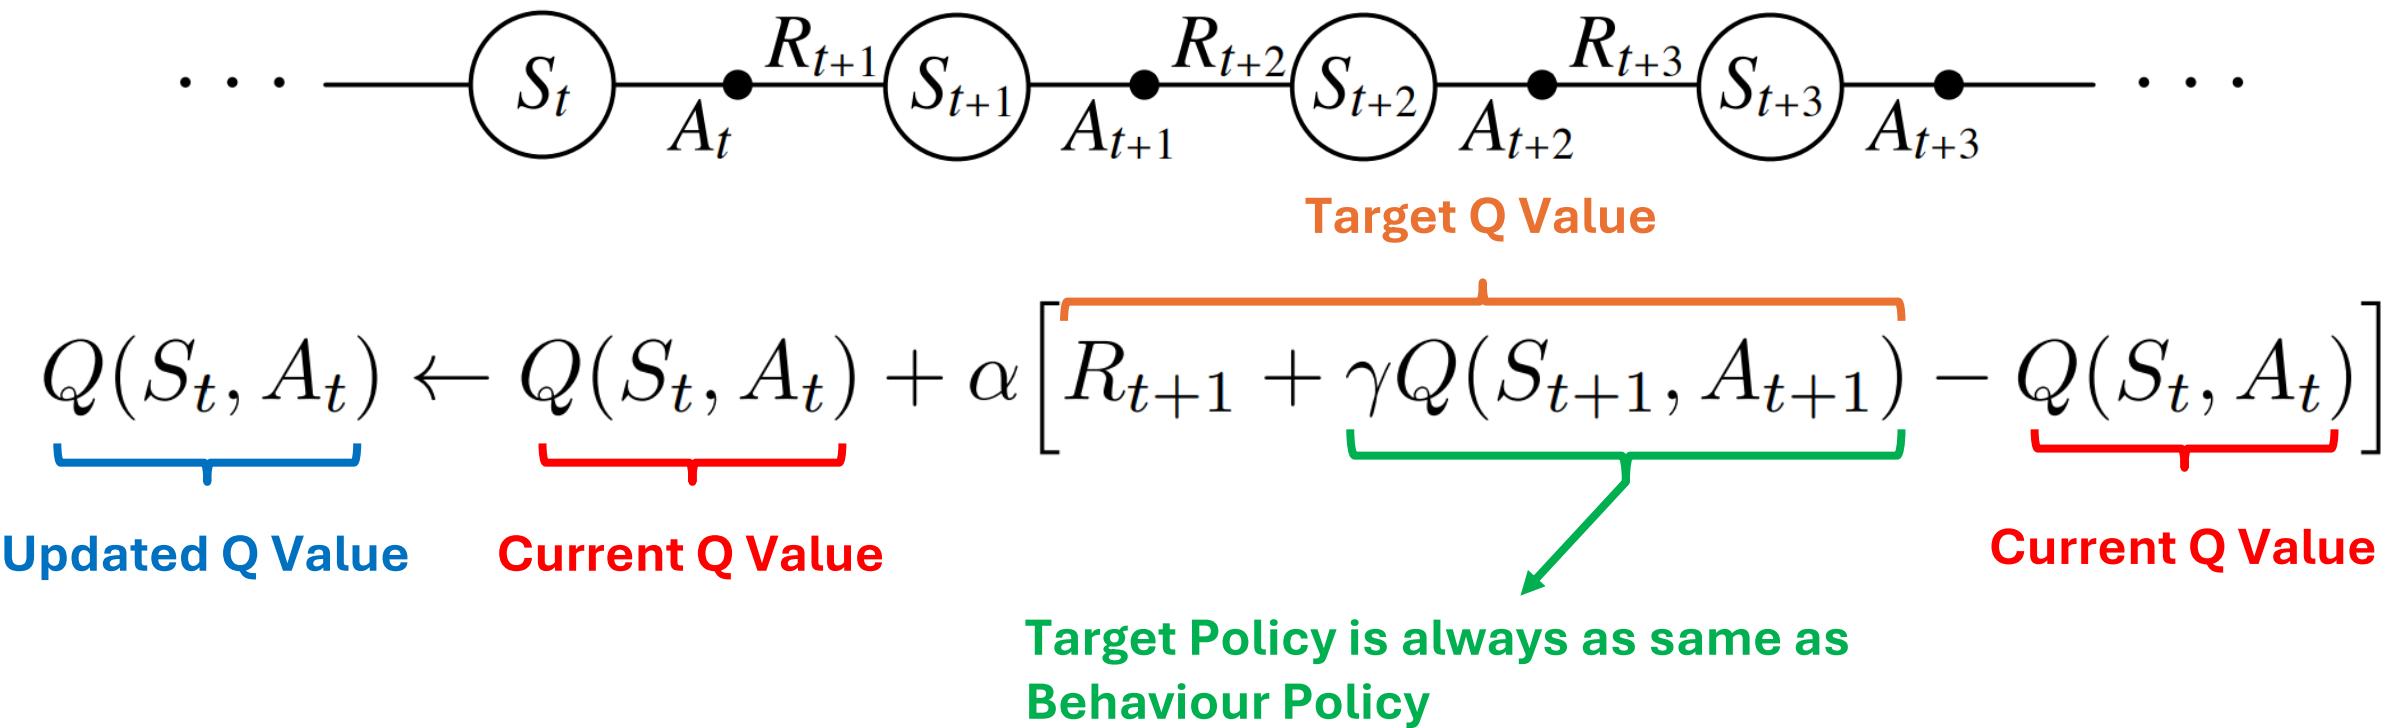


## SARSA: On-Policy TD Control



$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[ R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t) \right]$$

## SARSA: On-Policy TD Control



## Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+$ ,  $a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

    Loop for each step of episode:

        Take action  $A$ , observe  $R, S'$

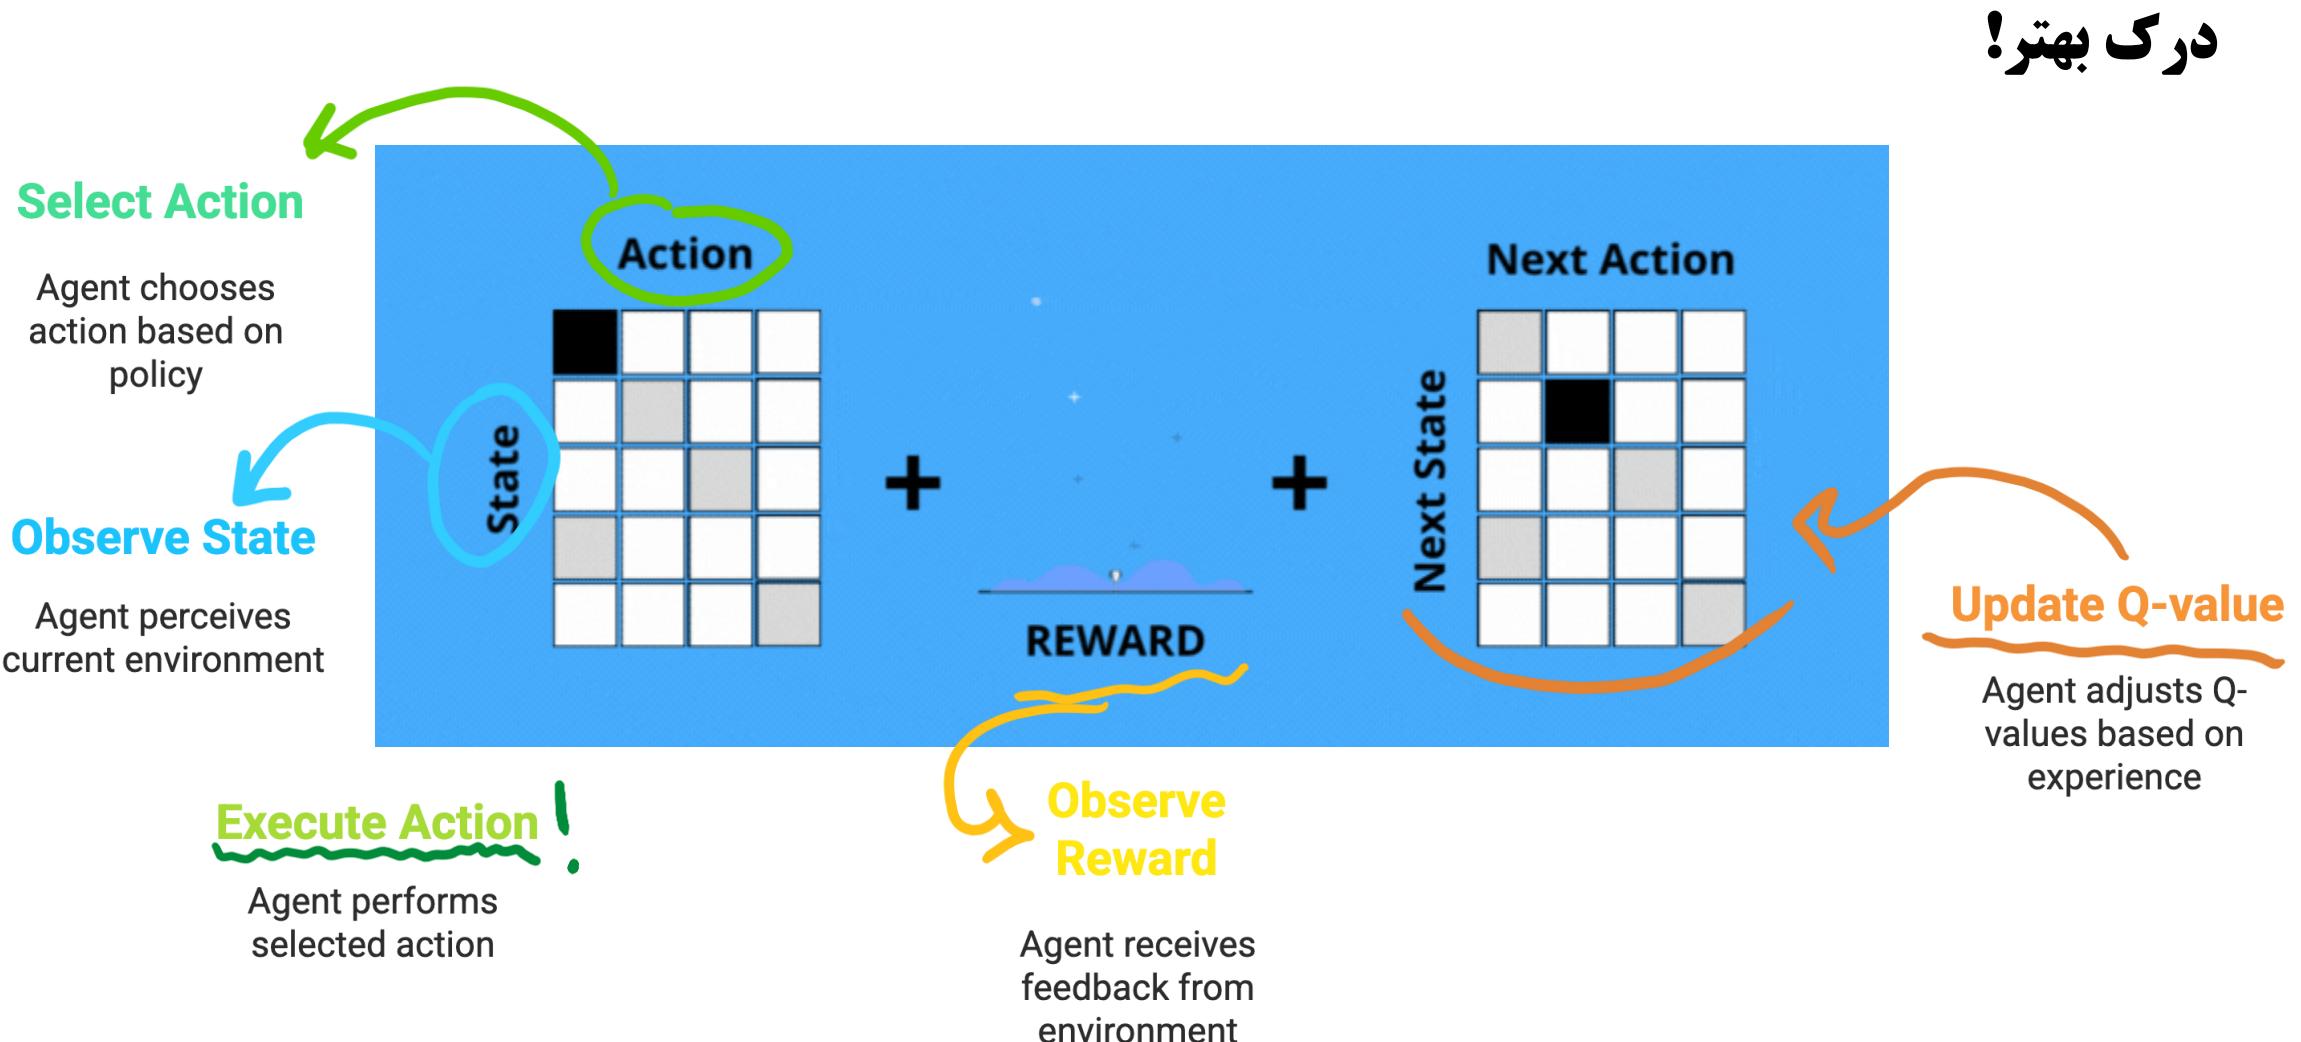
        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$$

$S \leftarrow S'; A \leftarrow A'$ ;

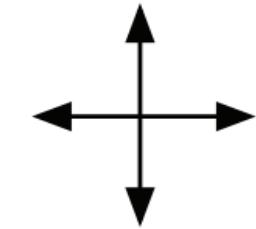
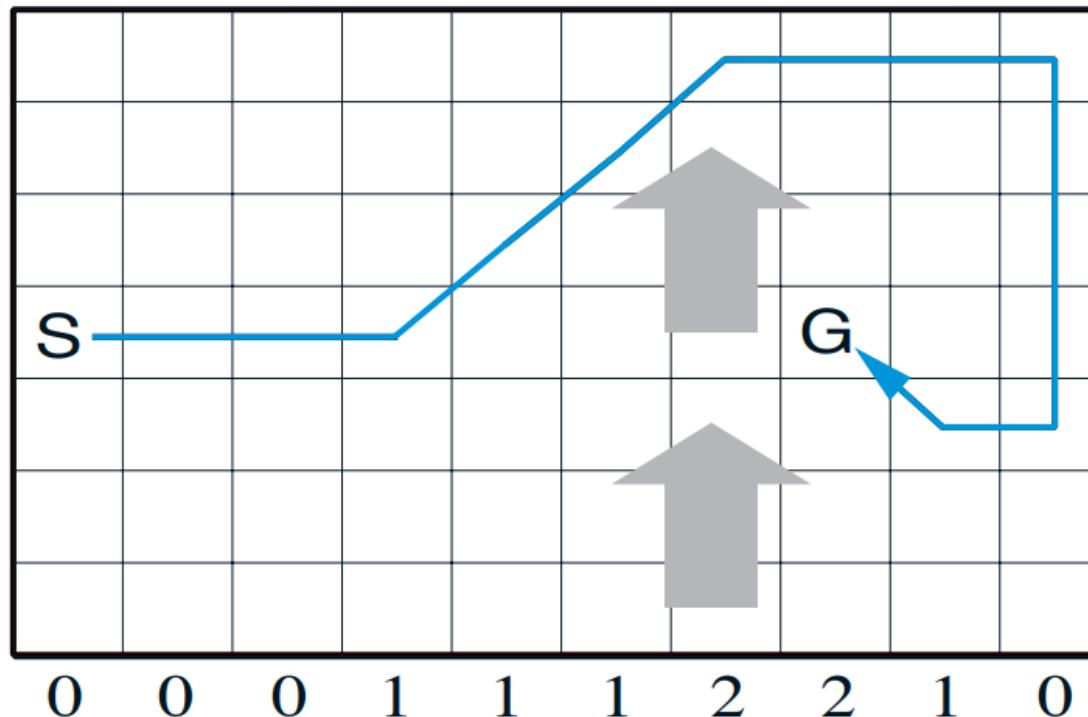
    until  $S$  is terminal

## I TD Learning



## Windy Grid

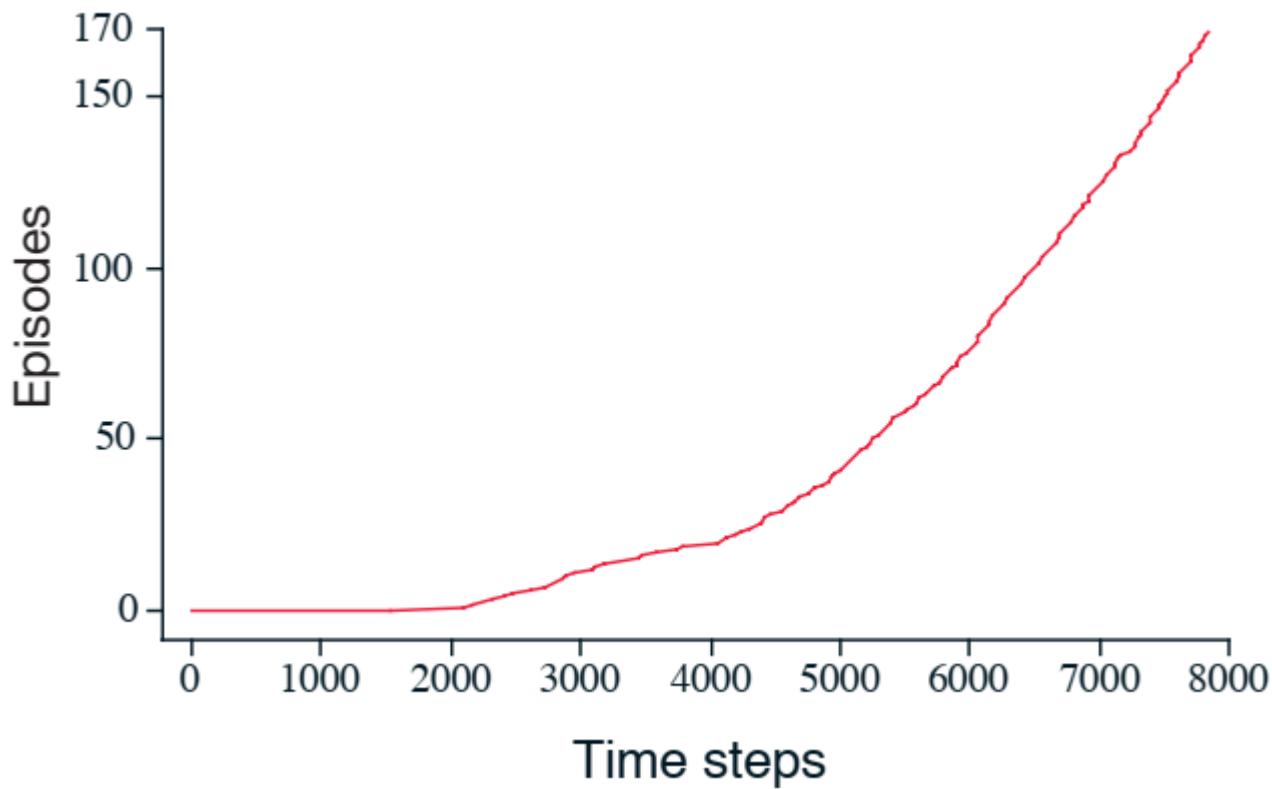
مثال



Actions

## Windy Grid

مثال



## Q-Learning: Off-Policy TD Control

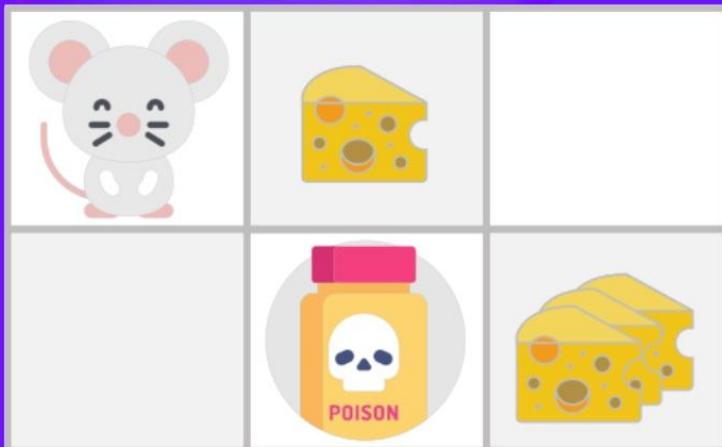
$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

The diagram illustrates the Q-Learning update rule with color-coded components:

- New Q-value estimation** (green bar)
- Former Q-value estimation** (blue bar)
- Learning Rate** (red bar)
- Immediate Reward** (orange bar)
- Discounted Estimate optimal Q-value of next state** (purple bar)
- TD Target** (dark blue bar)
- TD Error** (yellow bar)

The update rule is shown as:  $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$

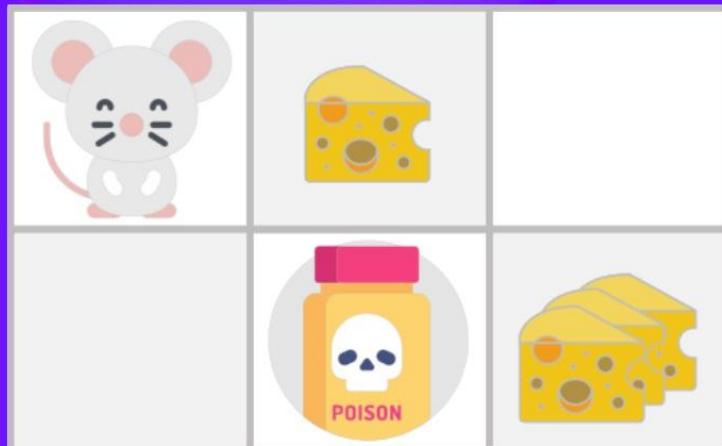
## Example



- You always start at the **same starting point**.
- The goal: eat the **big pile of cheese** (at the bottom right-hand corner) and **avoid the poison**.
- The episode ends if we eat the poison, eat the big pile of cheese or if we spent more than 5 steps.
- Learning rate = 0.1
- Gamma = 0.99

## Example

- The reward function:
  - 0: Going to a state **with no cheese in it.**
  - +1: Going to a state with a **small cheese in it.**
  - +10: Going to the state with **the big pile of cheese.**
  - -10: Going to the state **with the poison and thus die.**



## Example, Step 1

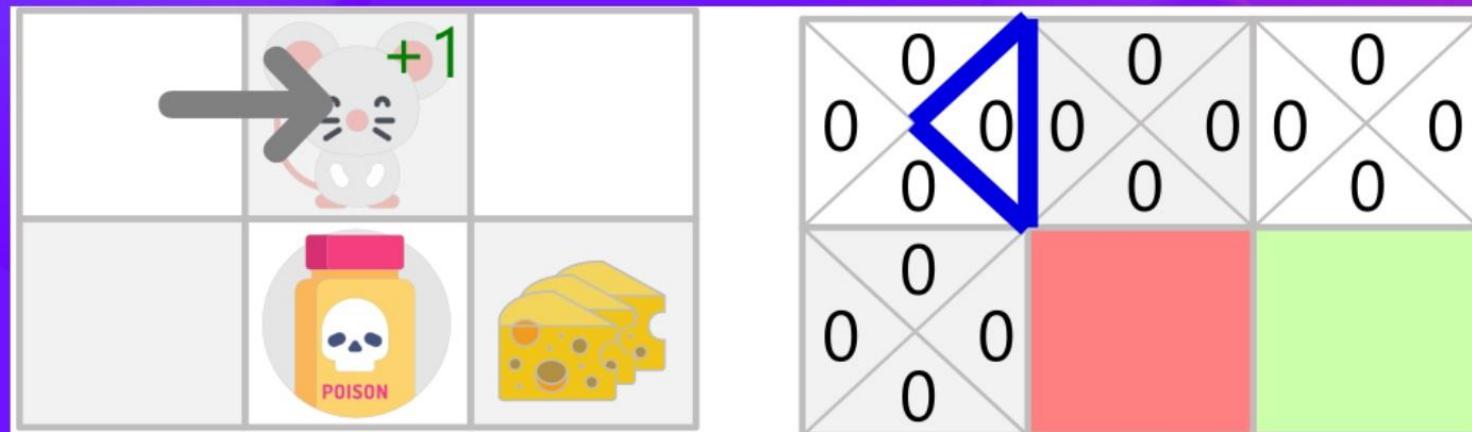
Initialize  $Q$  arbitrarily (e.g.,  $Q(s, a) = 0$  for all  $s \in \mathcal{S}$  and  $a \in \mathcal{A}(s)$ , and  $Q(\text{terminal-state}, \cdot) = 0$ )

	$\leftarrow$	$\rightarrow$	$\uparrow$	$\downarrow$
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0

We initialize the Q-Table

## Example, Step 2

Choose action  $A_t$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)



We took a random action (exploration)

## Example, Step 3

Take action  $A_t$  and observe  $R_{t+1}, S_{t+1}$



## Example, Step 4

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

New  
Q-value  
estimation

Former  
Q-value  
estimation

Learning  
Rate

Immediate  
Reward

Discounted Estimate  
optimal Q-value  
of next state

Former  
Q-value  
estimation

TD Target

TD Error

Update our Q-value estimation

## Example, Step 4

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t))$$

$$Q(\text{Initial state, Right}) = 0 + 0.1 * [1 + 0.99 * 0 - 0]$$

$$Q(\text{Initial state, Right}) = 0.1$$

	$\leftarrow$	$\rightarrow$	$\uparrow$	$\downarrow$
	0	0.1	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0

## Q-Learning Recap

Q-learning (off-policy TD control) for estimating  $\pi \approx \pi_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+$ ,  $a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Loop for each step of episode:

        Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

        Take action  $A$ , observe  $R, S'$

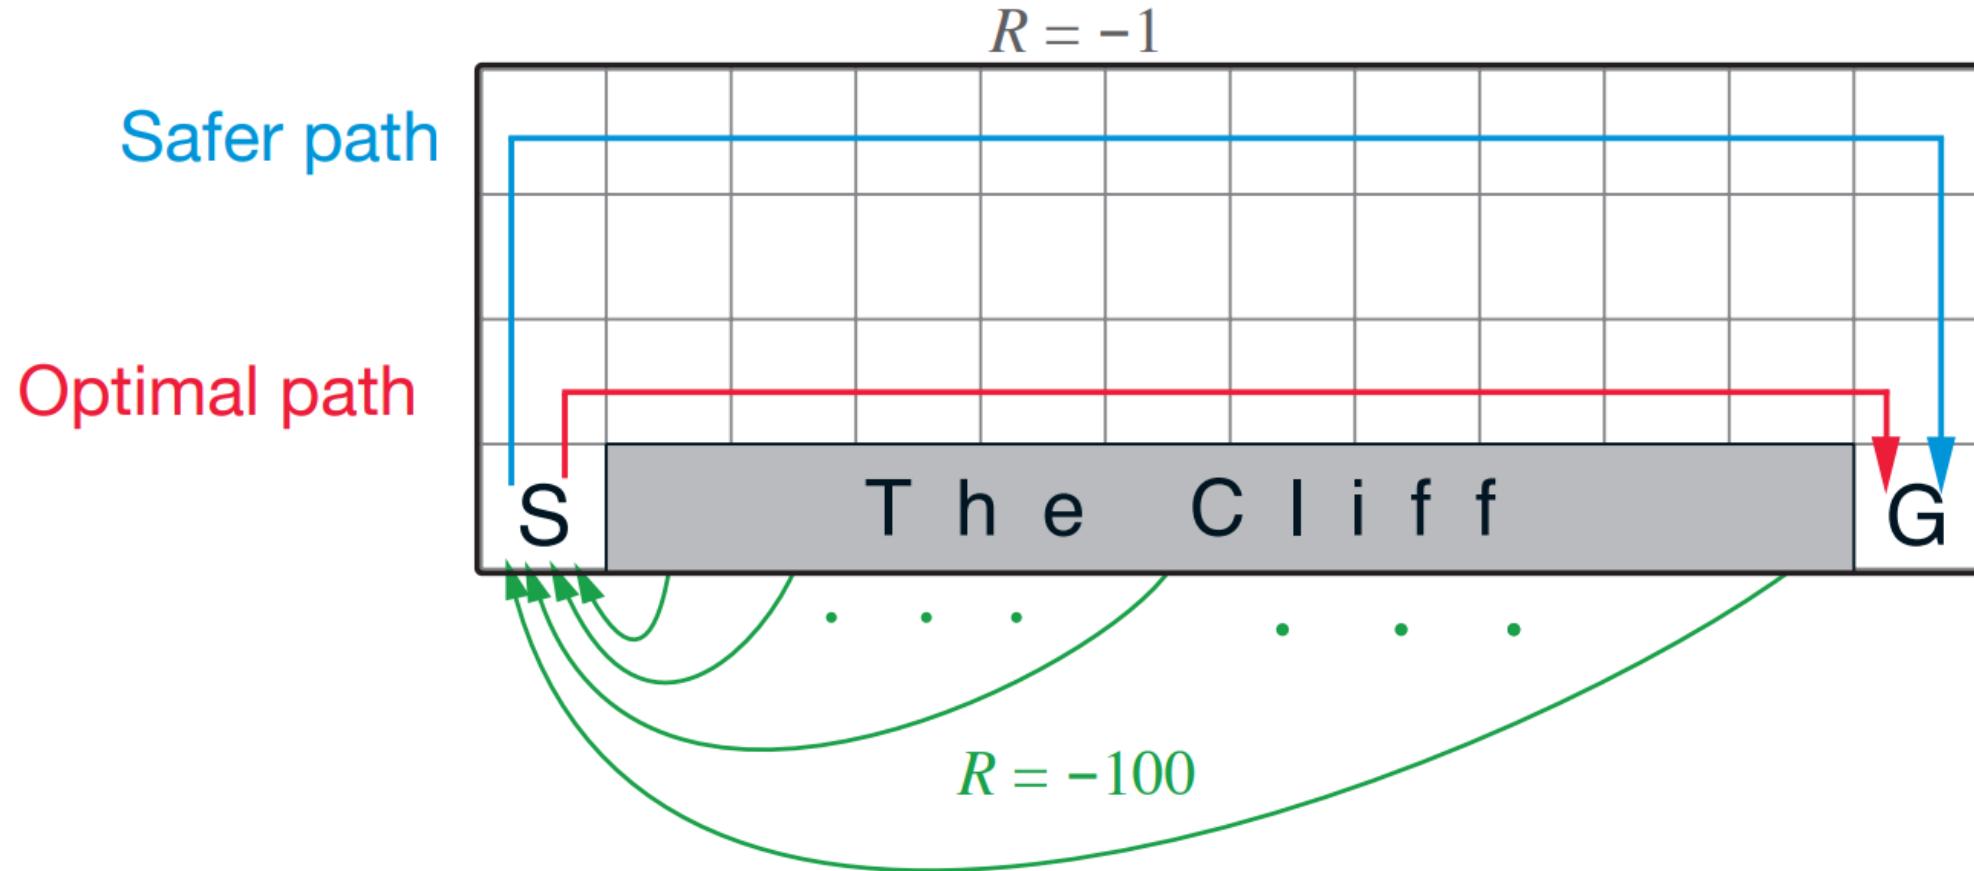
$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$$S \leftarrow S'$$

    until  $S$  is terminal

# Cliff Walking

مثال



**Cliff Walking**

مثال

